

Acquisition, Conversion, Service, Retention, and Expansion in the Age of AI

Jweokk · 2026-10-04 · v1.0.5

In 2024, Klarna handled the workload of 700 full-time customer-service agents with an AI assistant — and the global media went wild. By 2026, its CEO said "we went too far" and started hiring people again.

Not a story about whether AI works, but about how the full customer lifecycle should be managed in the age of AI — across the full lifecycle, from acquisition to expansion: successes and failures, both sides told.

Copyright © Jweokk. Free for reading and non-commercial sharing; commercial use requires prior written permission.

Contact: weokk2025@gmail.com

Contents

- 01** Preface: The Customer Is the Last Battlefield of the AI Era
- 02** Chapter 1: Year One — From Klarna to "We Went Too Far"
- 03** Chapter 2 The Customer Lifecycle and the AI Engine
- 04** Chapter 3 Acquisition: AI Turns Prospecting into Capacity
- 05** Chapter 4 Conversion: Personalization, Pricing, and Conversational Commerce
- 06** Chapter 5 Service: AI Support and Human Backup
- 07** Chapter 6 Retention and Expansion: AI-Driven Customer Success
- 08** Chapter 7 The Foundation: AI CRM and Customer Data
- 09** Chapter 8 Organization: People, Skills, and Change
- 10** Chapter 9 Industries: AI Customer Lifecycle Management in Eleven Industries
- 11** Chapter 10 The Dark Side and Boundaries: Backlash, Hallucination, Compliance, and Liability
- 12** Chapter 11 The Future: When the Customer Is an AI
- 13** Appendix A: Case Index and Sources
- 14** Appendix B: Common Metrics
- 15** Appendix C: Version History and Update Notes
- 16** Afterword

Preface: The Customer Is the Last Battlefield of the AI Era

In February 2024, Klarna, the Swedish fintech company, dropped a bombshell: in its first month, an AI assistant built on OpenAI technology had handled 2.3 million customer-service conversations, absorbing two-thirds of the workload—the equivalent of 700 full-time agents. Average resolution time fell from 11 minutes to 2 minutes, customer satisfaction matched human agents, and the company expected to improve profits by \$40 million a year.

The global media went into a frenzy. "AI replacing customer service" had moved from science fiction to earnings-report numbers.

Two years later, in 2026, the story flipped. In an interview, Klarna's CEO Sebastian Siemiatkowski admitted: "We went too far." Cost-driven evaluation had led to a decline in service quality; the company rebuilt its human customer-service channels and announced new hiring at its Stockholm headquarters.

The same company, two years apart, went from "AI doing the work of 700 agents" to "we went too far." What happened in between? Was AI not up to the job? Or is "customer operations" far more complex than "using AI to cut costs"?

That is the question this book sets out to answer.

What I Have Been Researching

Over the past two years, the most contested battlefield for AI inside the enterprise has not been writing code or creating designs—it has been the **customer**: how to find customers, how to get them to pay, how to serve them, and how to keep them. Salesforce's Agentforce reached \$1.5 billion in ARR within a year; when the AI sales agent Artisan raised \$25 million, its fundraising materials said it would "replace human salespeople with AI"; Gartner predicts that by 2028, AI agents will intermediate more than \$15 trillion in B2B purchases.

At the same time, the forces running the other way are just as powerful: AI cold email has become so pervasive that 95% draws zero replies, destroying the channel's credibility; hallucinations in AI customer service have cost airlines money and embarrassed delivery companies in public; and consumer trust in AI is far lower than vendors' slide decks claim.

On one side stands the great trend of AI reshaping customer operations; on the other, a trail of spectacular failures. Public discussion is either vendors hyping "AI can do anything" or victims complaining that "AI failed." Few people put the two side by side and carefully dissect them: **What exactly has AI changed about customer operations? What has not changed? Which parts can AI genuinely carry, and which parts must remain with people?**

This book attempts to do exactly that. It is not about "how to use AI tools"; it is a panorama of customer operations in the AI era—from AI-powered acquisition, conversion, service, and retention to organizational change. Every chapter speaks through real cases and data, and it covers successes and failures alike.

Where the Material Comes From

The material for this book comes from a continuously running collection pipeline: the knowledge base now holds 398 verified Chinese and English sources—official vendor data, third-party reports, research-institute studies, and founder interviews—each tagged [official claims] or [third-party verified] and traceable back to the original text. Every figure and case in the book is sourced in Appendix A.

This is not a book written all at once. It is a "living book"—the material keeps accumulating, and the edition keeps growing. If you find something incomplete in one edition, chances are the next update has already filled it in.

Who This Book Is For

It is written for three kinds of readers:

- **Practitioners in sales, customer success, and growth teams**—AI is changing how you work, and this book helps you see the full picture: which parts of the job AI will take over, and which skills will become more valuable;

- **Business operators**—you must decide how much AI to invest in customer operations, and how. The failure cases and ROI data in this book are closer to reality than any vendor sales demo;
- **Anyone curious about what people will do in the AI era**—customer operations is the best window for observing the division of labor between humans and machines, because every dollar here is tied to a real person.

The judgments in this book come from cross-validating the source material, and from my continued observation of this field. If you spot any errors or omissions, you are welcome to point them out (weokk2025@gmail.com).

Acknowledgments

The organization and public-release model of this book draw on the open-source practice behind *Frontline Deployment Engineer* by Fan Bing (Chinese: 范冰), better known as XDash — free public access to the full text, a GitHub repository plus website reading and PDF download, and sources cited in the appendix. Knowledge lives in movement, and this model deserves to be spread.

Chapter 1: Year One — From Klarna to "We Went Too Far"

1.1 From Klarna to "We Went Too Far"

On February 27, 2024, a Klarna press release swept across global tech media. The Swedish "buy now, pay later" company announced that in its first month, an AI assistant built on OpenAI technology had completed 2.3 million conversations, absorbing two-thirds of all customer-service chats—the workload of 700 full-time agents. Average time to resolve an issue dropped from 11 minutes to 2 minutes, customer satisfaction matched human agents, repeat inquiries fell by 25%, and profits were expected to improve by \$40 million a year^[5-28] [official claims].

This was the first time "AI customer service" had world-class benchmark data. Klarna's CEO Sebastian Siemiatkowski even said publicly that he wanted Klarna to be OpenAI's "favorite experimental subject."

If the story had stopped there, this book would need only one chapter: how impressive AI customer service is.

But the story did not stop. In 2025, Siemiatkowski admitted publicly: "Cost-driven evaluation led to a decline in service quality." The company rebuilt its human customer-service channels. By 2026, he said something even stronger to Bloomberg at the company's Stockholm headquarters: **"We went too far."** Klarna began hiring again—not agents to replace AI, but "around 100 highly skilled humans" to handle the situations AI cannot resolve, or resolves in ways customers do not accept^[12-06] [third-party verified].

Two years, one company, from "AI covering 700 agents" to "we went too far." International media called it Klarna's "sharp U-turn on AI customer service."

How should we understand this U-turn? On three levels:

First, Klarna's AI customer service did not fail. To this day it still handles about 80% of customer-service chats—31 million conversations over the whole of 2025—the workload of more than 700 full-time agents^[5-23] [third-party verified]. The cost savings are real.

Second, the "success" of AI customer service and "customer satisfaction" are not the same thing. AI can efficiently resolve procedural issues like "check my bill, change my address, ask about progress." But when customers hit complex, emotional situations that require judgment, the experience collapses in front of a bot that is always polite yet never truly understands you. It took Klarna two years to discover: **cutting costs and keeping customers are two separate ledgers.**

Third, and most important: Klarna's lesson is not that "AI doesn't work," but that "using AI as a cost-saving machine hurts customer operations itself." The right way to use AI is not to replace people, but to let humans and AI each do what they do best—and that is exactly what this book sets out to explain from beginning to end.

1.2 "Customer Management" Is Dying

Klarna's story is a microcosm. Behind it lies a larger paradigm shift: **"customer management" is becoming "customer operations."**

What is the difference between the two terms?

"Customer management" is a legacy of the industrial age. Its implicit assumption: the customer is an input to a process. Sales manages contracts, service manages complaints, marketing manages campaigns, finance manages collections—each department runs one small slice of the customer lifecycle, while customer information is scattered across the CRM, support tickets, and financial systems, and no one is truly accountable for "how much this customer is worth across their entire lifecycle." However well customer management is done, it only keeps customers "under control"—don't lose them, don't let them complain, don't let them default.

"Customer operations" rests on a different assumption: the customer is a source of assets. The operator's task is to make every customer keep generating value across the entire lifecycle—find the right customers, get them onboarded smoothly, serve them well, keep them around, get them to buy more, and get them to recommend others. This requires not just processes, but **continuous understanding** of customers and **proactive action** on their behalf.

AI makes "customer operations" scalable for the first time. Traditional customer management runs on people: one customer success manager can maintain dozens of customers, and one salesperson can follow up a few hundred leads—that is the ceiling of capacity. AI can analyze behavioral data from thousands of customers at once, spot patterns human teams cannot see, and automatically execute the right action at

every touchpoint^[1-17] [third-party verified]. ZoomInfo's customer lifecycle management framework calls this a "cross-functional operating system"—unifying data and processes across marketing, sales, customer success, and product teams to scale fulfillment around key buyer personas [third-party verified].

The blunter version comes from the B2C world: Blueshift's lifecycle marketing guide did the math—acquiring a new customer can cost 5 to 25 times as much as retaining one, and a 5% improvement in retention can lift profits by 25% to 95%^[1-06] [official claims]. Once AI can automate and scale the business of operating existing customers, customers are no longer merely managed—they are actively operated.

1.3 Data Can Lie: Three Truths About AI-Powered Customer Lifecycle Management

Before we get to the substance, let me put three easily misleading "truths" on the table. They come from cross-validating 398 pieces of source material, and from observing a great many failure cases.

Truth one: adoption of AI-powered customer lifecycle management is badly overestimated. Intercom surveyed 2,470 customer-service professionals worldwide: 82% of executives said they had invested in AI over the past 12 months, and 87% planned to keep investing in 2026—yet only 10% had reached "mature deployment" (AI deeply integrated into core operations and running at scale)^[1-09] [official claims]. Between buying AI and using it well lies an enormous gap.

Truth two: vendor-reported ROI runs 30-40 percentage points above independent data. The 53 AI customer-service ROI studies compiled by Digital Applied show an average return of \$3.50 for every \$1 invested by vendors' own accounting, while independent benchmarks are markedly lower^[5-04] [third-party verified]. It is not that all vendors are lying; rather, vendors tend to count "the volume of conversations AI handled," not "the improvement in customer satisfaction and retention."

Truth three: customers trust AI, but they trust people more. EY's Future Consumer Index survey of 21,000 respondents across 27 countries shows that consumers quickly adopt AI tools that save money and time while staying wary of AI—24% worry that AI will take their place^[1-16] [third-party verified]. Metrigy's research is

more direct: 85% of consumers prefer interacting with a real person in customer service^[11-15] [third-party verified]. The stronger AI gets, the scarcer—and the more valuable—the "human touch" becomes.

1.4 A Map of This Book

This book unfolds along the five stages of the customer lifecycle: **Acquisition → Conversion → Service → Retention → Expansion**. Each chapter answers one question:

- Chapter 2: the overall framework of AI-powered customer lifecycle management—the lifecycle model, the divide between AI-native and AI-assisted approaches, and when AI should not be used;
- Chapter 3: acquisition—how AI turns prospecting into capacity, and why AI outbound is killing cold email;
- Chapter 4: conversion—personalization, pricing, and conversational commerce: how AI is changing the business of "getting customers to pay";
- Chapter 5: service—AI customer-service platforms, human-AI collaboration, and the failure scenes;
- Chapter 6: retention and expansion—AI-powered customer success: churn prediction, health scores, and NDR;
- Chapter 7: the foundation—AI CRM and customer data: without data, everything else is empty talk;
- Chapter 8: organization—rebuilding the GTM, sales training, and the psychology of organizational change;
- Chapter 9: industries—a map of AI-powered customer lifecycle management across eleven industries;
- Chapter 10: the flip side—backlash, hallucination, compliance, and liability: the boundaries of AI-powered customer lifecycle management;
- Chapter 11: the future—when customers themselves become AI.

Every chapter follows the same writing discipline: **success stories and failure stories are told side by side**. In the field of AI-powered customer lifecycle management, failure cases often teach us more than success stories—and Klarna is the best example.

Chapter 2 The Customer Lifecycle and the AI Engine

2.1 The Five-Stage Model: The Master Framework of Customer Operations

The previous chapter said that "customer management is dying and customer operations are arriving." So what exactly does customer operations operate on? The answer is a five-stage framework:

Acquisition → Conversion → Service → Retention → Expansion

This framework is not a new invention. Blueshift's B2C lifecycle marketing guide frames it as five stages — "acquire, onboard, engage, retain, win back"^[1-06][official claims]; monday.com's 2026 lifecycle marketing guide argues that the industry is shifting from "campaign-centric" to "customer-centric" — traditional marketing consists of isolated campaigns with start and end dates, whereas lifecycle marketing continuously delivers personalized messages according to the stage of the customer relationship, and the metrics shift from clicks and opens to customer lifetime value and retention growth^[1-21][third-party verified].

Why does this framework suddenly become important in the AI era? Because AI turns every stage from "labor-intensive" into "scalable operations":

- **Acquisition:** an AI SDR can reach out to thousands of leads in a single day, something human salespeople cannot do;
- **Conversion:** AI-powered personalized recommendations can simultaneously serve millions of customers with different needs;
- **Service:** AI customer service can handle thousands of conversations at once;
- **Retention:** AI churn prediction can spot risk signals months before a customer leaves;
- **Expansion:** AI can identify upsell opportunities for each customer and push NDR (net dollar retention) higher.

Traditional customer management split these five stages across different departments, each looking only at its own segment. The first change AI-powered customer lifecycle management brings is **reconnecting the five stages into one continuous line** — because AI needs complete customer data to work, and once complete customer data starts flowing, the barriers between the five stages naturally break down. ZoomInfo calls this CLM (customer lifecycle management): systematically optimizing every touchpoint on the customer journey, like a cross-functional operating system that unifies marketing, sales, customer success, and product teams^[1-17][third-party verified].

2.2 AI-native vs. AI-assisted: A Spectrum

Any discussion of AI-powered customer lifecycle management must confront a conceptual divide: **AI-native and AI-assisted**.

Nutshell uses an "AI spectrum" to explain why so many companies invest in AI yet see no results: only a minority of companies achieve measurable enterprise-level gains; most CRMs are merely "AI-assisted" — layering an AI tier on top of legacy systems; whereas AI-native embeds AI into the core database, workflows, and automation — AI is not a feature attached to the database but the architecture itself^[1-05][third-party verified].

What is the difference? Take a concrete example. An AI-assisted CRM has salespeople manually entering customer information while AI prompts on the side, "this customer might be interested"; an AI-native CRM automatically captures every email, every call, every meeting, keeps records fresh on its own, and prepares for humans^[1-12] "next actions pending approval" — AI does the work, people make the decisions^[1-12][official claims].

This divide applies to the whole of customer operations:

Dimension	AI-assisted	AI-native
Where AI sits	Add-on feature	The architecture itself
Data	People enter it, AI analyzes	AI captures automatically, people review
Workflows	Existing processes + AI acceleration	Processes designed around AI

Dimension	AI-assisted	AI-native
Role of humans	Doers	Decision-makers/approvers
Typical outcome	15% adoption, "no results" at year-end review	Measurable metrics, running continuously

The "bought AI but got no results" cases that recur throughout this book are almost all products of the AI-assisted approach: tools purchased, training delivered, incentives given — yet the organizational structure, decision processes, and performance standards are all traditional, with AI merely pasted onto legacy systems.

There is a simple test for whether a customer operations system is AI-native: **if you were designing it from scratch tomorrow, would AI be written into the blueprint?** If the answer is no, it is just a band-aid.

2.3 Six Models of AI-native GTM

When customer operations takes AI as its architectural core, the way acquisition and revenue are done changes as well. A study of companies racing toward \$100M in revenue summarizes six models of AI-native GTM^[1-19][third-party verified]:

1. **AI-generated content for acquisition:** use AI to mass-produce high-quality content so search engines and AI search recommendations bring a steady stream of traffic. Jasper gained over 100,000 organic search visits in 18 months through programmatic content^[3-24]; Zapier's programmatic SEO quadrupled its organic traffic^[3-22][third-party verified].
2. **AI-powered personalization for conversion:** AI understands each visitor's intent in real time and delivers the right information at the right moment.
3. **Agentic sales:** AI agents autonomously handle outreach, follow-up, and booking, leaving only the final high-value step to humans.
4. **AI customer success:** AI proactively monitors customer health and intervenes before customers churn.
5. **AI pricing:** outcome-based pricing — customers pay for "work done" rather than "software seats."
6. **AI analytics for decisions:** AI synthesizes customer, pricing, competitive, and deal signals and directly recommends next actions.

What these six models share: **AI is not assisting humans in doing the same things as before — it changes the very way things are done.** EY-Parthenon's observation is sharper: agentic AI has dramatically increased the go-to-market complexity of SaaS companies, and "winners will be determined by the ability to reinvent business models, not by the size of sales teams"^[1-14][third-party verified].

2.4 The Engine and the Foundation of Customer Operations

Combine the five-stage model with AI-native, and you get this book's organizing framework:

- **The engine:** acquisition, conversion, service, retention, expansion — what AI transforms each of the five stages into (Chapters 3–6);
- **The foundation:** AI CRM and customer data — the fuel for all the engines (Chapter 7);
- **The organization:** who drives these engines — GTM teams, sales training, and change management (Chapter 8);
- **The industries:** different industries configure their engines differently (Chapter 9);
- **The boundaries:** the engines' braking system — backlash, compliance, liability (Chapter 10);
- **The future:** when customers themselves become AI, the engines must be redesigned (Chapter 11).

2.5 When Not to Use AI

This section sits in the overview chapter because it is one of the most important judgments in this book: **the first step of AI-powered customer lifecycle management is not deciding where to use AI, but deciding where not to use it.**

In 2026, an interesting counter-trend emerged: a group of brands turned "we don't use AI" into their strongest selling point. "No AI" is becoming a new marketing label after "all natural," triggered by consumers' rising aversion to AI content and their growing ability to recognize it^[1-01][third-party verified].

This is not anti-intellectualism — it is the market voting with its feet. Research draws clear boundary lines:

In high-trust scenarios, humans are irreplaceable. A survey Northwestern Mutual commissioned from Harris Poll shows that Americans trust human advisors significantly more than AI when making financial decisions^[1-08][official claims]. Psychologists at Teachers College, Columbia University warn that AI's most popular use cases — emotional support and companionship — are precisely where AI's architecture is the least suited, because AI cannot truly empathize^[1-07][third-party verified].

In luxury and experiential scenarios, AI should step back. McKinsey's luxury research (surveying 300 consumers and 90 merchants) finds that 85% of luxury consumers already use generative AI, but when it comes to purchase decisions, what they expect is "the human touch" — AI can assist, but cannot lead^[1-10][third-party verified].

In emotional and complex scenarios, AI fails. Klarna is the lesson here: AI handles process problems efficiently, but when customers arrive emotionally charged, the more polite the bot, the more infuriating it becomes.

So where should AI be used? The answer: **repetitive, standardized, high-volume, low-emotion scenarios** — first-touch outreach, answering common questions, billing inquiries, process guidance, behavior analysis, and risk alerts. In one sentence: **AI handles scale and efficiency; humans handle trust and judgment.** IBM's thought-leadership article puts it more elegantly: AI is making customer service "more human" — because it leaves the complex cases that require emotional intelligence to people^[1-18][third-party verified].

This boundary is not static. As AI capabilities improve, the boundary moves; but as customers' ability to recognize AI grows, the value of the "human touch" rises as well. The operator's task is not to draw the boundary once and for all, but to keep recalibrating it.

The next chapter begins the first stop of the five stages: acquisition.

Chapter 3 Acquisition: AI Turns Prospecting into Capacity

3.1 Artisan: The Company That Printed "Stop Hiring Humans" on a Billboard

In 2025, a series of billboards appeared on the streets of San Francisco — white background, black text, a single line: "Stop Hiring Humans." They were signed by a startup called Artisan, whose product is Ava, an AI sales agent that automatically finds prospects, researches leads, writes personalized emails, follows up, and books meetings — reportedly at one-fifth the cost of a human BDR (business development representative), with more than 300 million verified B2B contacts built in^[3-15][official claims].

The company later raised a \$25 million Series A. Business Insider exclusively reported on its fundraising materials — which promised not "sales assistance" but "replacing human sales with AI"^[3-14][third-party verified].

The other side of the story is equally telling. In late 2025, Artisan's AI SDR, Ava, was banned from LinkedIn — because its automated way of contacting prospects violated the platform's rules. CEO Jaspar Carmichael-Jack confirmed the ban to TechCrunch; the account was restored two weeks later^[3-44][third-party verified]. Around the same time, the founder publicly admitted that early versions of the product had hallucination problems and that the company had gone through customer churn^[12-16][third-party verified].

A company built on "stop hiring humans" — banned by a platform, disciplined by customer churn — and then it came back. That complexity is an accurate portrait of AI-driven acquisition as a field: **the efficiency is real, and so is the backlash.**

3.2 AI SDR: From Tool to Digital Employee

First, consider the size of the market. The AI SDR (sales development representative) category was worth roughly \$4.27 billion in 2025 and is projected to reach \$18.19 billion by 2032 (a 23% compound annual growth rate); over the past two years, venture capital has poured more than \$1 billion into AI SDR startups^[3-07][third-party verified].

Product forms have already evolved through two generations:

The first generation is "assistive tools": they help human SDRs write emails, find leads, and prioritize — the human remains the lead actor.

The second generation is "digital employees": 11x's product line is the most representative — Alice (email + LinkedIn outreach), Julian (phone agent), and other "digital employees" hold their own phone numbers and, 24/7 with no human supervision, make calls, answer calls, screen leads, and book meetings, automatically writing call data back to the CRM^[3-38][official claims]. Artisan's Ava covers the full flow of "lead research → personalized email → follow-up sequences"[official claims]. AiSDR's own ROI calculation is more direct: an AI SDR subscription costs about 85% less than hiring an in-house SDR, with similar output^[3-12][official claims].

Both capital and vendors are telling a story of "replacement." But independent data gives a more complicated answer.

3.3 AI vs. Human SDRs: What the Head-to-Head Data Says

Salesmotion's aggregation of adversarial head-to-head tests offers a rare direct comparison[third-party verified]:

- **Revenue:** human SDRs generated 2.6x the revenue of AI SDRs^[3-06] (\$147,000 vs. \$56,000);
- **Meeting show-up rate:** humans 71%, AI 52%;
- **Cost:** a human SDR's fully loaded cost is \$100,000–150,000 per year; an AI SDR license costs roughly \$30,000 per year.

AI delivers volume hundreds of times greater than humans; humans deliver quality two to three times higher than AI. When the full math is done, the conclusion is not "who replaces whom" but a division of labor: **AI owns breadth and first-touch outreach; humans own high-value conversion and complex relationships.**

This conclusion is corroborated at a grander level. Sequoia Capital's investment thesis — "services as the new software" — argues: if you sell tools (software), you are racing against the models; if you sell the work itself (service outcomes), every model upgrade makes your service faster and cheaper^[3-45][third-party verified]. Harvard Business Review defines agentic AI as a paradigm-level variable in sales: it makes a "perfect replica of the top salesperson" possible — autonomous agents unconstrained by time or geography, working side by side with human sellers^[3-42][third-party verified].

Note HBR's wording: **working side by side**, not replacing. The mainstream of 2026 is not "AI replaces salespeople" but a hybrid model of "AI sales + human sales" — Chapter 8 will present more detailed ROI data from controlled experiments on this model.

Chinese exporters offer a striking counterpoint. Jintai Technology (Chinese: 今泰科技), a Guangzhou company that applies superhard film coatings to metal surfaces, once sent 80,000 cold emails to overseas buyers in a single burst, almost all of which sank without a response. After adopting a sales agent from Tungee (Chinese: 探迹科技), salespeople only enter product information and target markets, and the agent finds buyers worldwide, drafts the outreach emails, and follows up across time zones; within three months its number of inbound inquiries equaled the total acquired over the previous three years, email reply rates rose at least tenfold, and it broke into the Western tableware industry that had long resisted its approach^[3-80][third-party verified]. Same channel, same email — output differed by three orders of magnitude. The difference was not volume but the data layer: the quality of buyer profiling and filtering determines the efficiency of every step that follows.

3.4 Content Acquisition: AI Turns "Writing" into "Manufacturing"

If AI SDR is proactive outreach, content acquisition is building a nest to attract the phoenix. The killer weapon of AI content acquisition is **programmatic SEO**: use AI to mass-produce pages targeting long-tail keywords, so search engines deliver a steady stream of precise traffic.

Three frequently cited cases:

Zapier: monthly visits grew from roughly 1.19 million in November 2020 to roughly 4.8 million in November 2023, an increase of more than 300% in three years. Its core assets are more than 6 million backlinks and tens of thousands of automation tutorial pages on "how to connect A to B"^[3-22][third-party verified].

Jasper: in 2021, this AI writing company had 14 employees in total — and just one in marketing. It outsourced its "content growth engine" to a specialist agency, and the result was 810% growth in organic blog traffic, 400x growth in blog-driven signups, and \$4 million in ARR^[3-24][third-party verified].

An AI image-generation SaaS: previously 67 signups a month and 13 ranking keywords. After a programmatic SEO overhaul, monthly signups climbed from 67 to more than 2,100 within 10 months, and organic traffic grew 3,035%^[3-28][third-party verified].

By 2026, the content acquisition battlefield has expanded from "search engines" to "AI search." GEO (generative engine optimization) has become a new term: optimize your brand content so it appears in the output of ChatGPT, Google AI Overviews, and Perplexity — the goal is to become the source an AI cites when it answers a user's question^[3-19][third-party verified]. Glasp pushes the practice further: this tool with 3.5 million users grew traffic from ChatGPT from 500 sessions a day to 19,000 — not by guessing how AI thinks, but by reading its own server logs to see directly which pages ChatGPT visited and cited^[3-51][third-party verified]. In the AI era, content acquisition is shifting from "guessing the algorithm" to "observing the AI."

3.5 ABM × AI: Treating Every Key Account as Its Own Market

The logic of ABM (account-based marketing) is: don't cast a wide net — lock onto a small number of high-value target accounts and run each one like its own "market."

AI makes ABM scalable for the first time. Demandbase divides AI's capabilities into three layers: NLP that reads unstructured data from web pages, search, and social media for intent monitoring; generative AI that automatically produces personalized content for each account; and predictive models that rank accounts^[3-10][official claims]. 6sense's platform uses predictive analytics and intent data to help companies discover and rank target accounts that are "entering their buying window"^[3-05][third-party verified].

Key data supports this direction: roughly 70% of the B2B buying journey is completed anonymously before buyers ever contact a vendor^[3-34][official claims] — they research, compare, and read reviews online first, and by the time a salesperson reaches out, half the decision is already made. The value of intent data is turning that "invisible shopping" into observable signals: whoever is intensively researching a category is a buyer in the market at this moment.

Chinese practice is keeping pace. SalesDriver's white paper shows that about 26.48% of companies have rolled out ABM, up 149% from the prior year^[3-30][third-party verified]. SANY (Chinese: 三一重工) is a sample from large-scale B2B manufacturing: in 2022, its overseas revenue grew 47% and accounted for 45% of total revenue, while revenue from developed markets in Europe and the Americas reached roughly RMB 15.8 billion, up 86% year over year^[3-40][third-party verified] — its playbook is not internet-style traffic operations but running target customers as "key accounts," supported by coordinated channel and content work (strictly speaking, not textbook ABM, but aligned in direction).

3.6 Private Domain: An AI Battleground Unique to the Chinese Market

If ABM is the Western B2B playbook, the private domain (Chinese: 私域) is China's own customer-operations battleground — adding customers to WeChat or WeCom (Chinese: 企业微信) and repeatedly reaching out to, nurturing, and converting them there.

The biggest pain point of the private domain is "added, then silent": large numbers of new contacts stop interacting within two weeks of being added. NetEase Zhiqi's (Chinese: 网易智企) diagnosis is sharp — the causes fall into two categories: silent contacts (never responded) and dormant old customers (bought before but inactive recently) — while blanket group broadcasts fail because of message folding, do-not-disturb settings, and mismatched content^[3-41][official claims]. AI's solution is "segment people first, then match content": a three-layer filter (basic attributes, behavioral traces, relationship stage) divides contacts into tiers, and each tier is paired with a different AI outreach strategy^[3-47][official claims].

Effect data comes from Shuyun (Chinese: 数云): during the 618 shopping festival, a well-known personal care brand's private-domain traffic pool surged beyond its reception capacity; after AI digital employees took over, the new-customer reception

rate rose 140%, the conversion rate rose 29.2%, and labor costs fell 61.3%^[3-43][official claims]. Tencent Qidian (Chinese: 腾讯企点) serves more than 1 million enterprises and connects 350 million users, embedding large models into its marketing cloud and intelligent customer service^[3-46][third-party verified].

3.7 Paid Media: AI Turns Budgets into Algorithms

Traditional paid media relies on manual bid adjustment and trial and error. AI-powered ad products turn this into an algorithmic problem: with Meta's Advantage+ campaigns, CPA (cost per acquisition) fell 12–17%, reaching 32% in some cases^[3-27][third-party verified]; one Google Performance Max campaign automatically runs across Search, YouTube, Display, and Gmail^[3-20][third-party verified].

At the budget-allocation level, the 2026 consensus is that MTA (multi-touch attribution) and MMM (marketing mix modeling) must run in parallel: MTA answers the tactical question — "which channel converted my customers" — while MMM answers the strategic one — "how the overall budget should be split"^[3-26][third-party verified]. Causal-AI versions of MMM go further: after causalLens helped a large consumer app company rebuild its attribution model, its annual marketing budget fell 5% and ROI reached 15x^[3-36][official claims].

3.8 The Backlash: AI Outbound Is Killing Cold Email

At this point in Chapter 3, we have to pause to cover the flip side — because the biggest risk in AI-driven acquisition is not that AI is too weak, but that AI is too easily abused, and abuse can destroy the entire channel.

The data is alarming. A cohort study of 14 B2B SaaS sales organizations found that reply rates to AI SDR outreach decayed by more than 60% within 18 months — from 11.2% at launch down to 4.4% by month 18 — because recipients gradually learned to recognize the structure of AI templates^[3-32][third-party verified]. The cold email ecosystem as a whole is worse off: 95% of cold emails now get zero replies at all; when personalization is traded for speed and volume, reply rates run 13x lower than normal^[3-39][third-party verified]. Cold call reply rates fell from 8.5% in 2018 to 4.2% in 2022[official claims].

The Pipeline Group's warning is harsher: running outbound sales at scale with AI SDRs is dragging companies into a state that is "illegal, unethical, and brand-destroying" — AI SDRs often use scraped data obtained without consent, crossing compliance red lines such as GDPR and CCPA^[3-09][official claims].

Hence a paradox that emerged in 2026: **AI has driven the cost of proactive outreach toward zero, and it has also driven the value of proactive outreach toward zero.** When everyone can send ten thousand emails with AI, recipients learn to ignore all emails. The first users consumed the channel's dividend; the imitators who followed only ate the backlash.

That explains why the acquisition consensus in 2026 is turning toward "relationship-driven": rather than casting a wide net with AI, use AI to precisely identify "who is in the market right now" (intent data), then reach them in a high-quality, personalized way. The second half of AI-driven acquisition is a contest of restraint, not intensity.

The mechanism of that backlash went one level deeper in 2026: detection no longer depends on a recipient's impression but on detectable watermarks and platform rules. Anthropic's Claude now embeds invisible watermarks in all generated text (a global change driven by European Union regulatory requirements), which means AI-drafted sales emails can be identified as machine-generated; at the same time, fully AI-generated cold emails receive fewer replies and are flagged as spam at a higher rate^[3-78][third-party verified]. Anti-AI outreach has moved from "a person can tell something is off" to "a system can measure it."

One practitioner's account puts a number on the slope: he spent more than 400 hours building his own AI SDR for cold-email prospecting and ultimately lifted reply rates only from 1% to 2%, while spending the same effort on phone outreach produced five times more pipeline, at a cost far below those 400 hours. The industry's average email reply rate fell from 8.5% in 2019 to 3.4% in 2026 — AI accelerated that decline because it lets everyone, including your competitors, pour more volume into the same inbox, while buyers, Gmail, and Outlook all learned to filter. The constraint on the phone side is a hard constraint: connect rates have long hovered at 3-5%, and statistics across billions of dials show that only about 20% of people in any market will answer a cold call, while the other 80% never will regardless of how, from what number, or at what hour you call — the AI SDR inherited this constraint rather than solving it^[3-79][third-party verified].

3.9 The AI-ization of Advertising and the Truthfulness of Traffic

Section 3.7 described how AI turns ad budgets into algorithms. Since then, two new shifts have appeared in this channel: the entry points are migrating, and the credibility of measurement is weakening.

Migrating entry points. In September 2026, Amazon and OpenAI struck a partnership: advertisers can extend campaigns through Amazon's DSP (demand-side platform) into ChatGPT's conversational ad slots, which surface as text and image units beneath the answers with clearly marked "sponsored" labels; the pilot is limited to the U.S. at launch, and Delta Vacations is among the first brands. ChatGPT Ads has been live for a little over half a year and has reached a \$1 billion annualized revenue run rate, with tens of thousands of advertisers using it^[3-52] [third-party verified]. When "conversation" itself becomes ad inventory, the AI assistant stops being a "substitute for search" and becomes "a new medium" — acquisition entry points are no longer just search engines and social platforms, but also the conversation between a brand and an AI assistant.

Conversion quality. The new entry point's traffic quality is documented: Semrush's search research found that users arriving via AI-generated answers convert 4.4 times better than those arriving from traditional search results^[3-53] [third-party verified]. The intent is clearer — the user has already walked to the last step of the decision, and the AI merely walks them to the door.

But the foundation of measurement is loosening. An indie developer ran a Google app-install campaign for his puzzle app and was billed for 56 installs in two weeks; going back through the raw panel data, 33 of them showed the same "bot pattern" — an old app version the store had stopped serving, a single open per device, zero seconds on any screen, and 28 phone models across 19 states^[3-54] [third-party verified]. The loop is the ironic part: Google optimizes for whatever goal the advertiser sets, and "a view followed by an install" counts as a conversion, so the more the farm "installed" the app, the more budget the algorithm sent the farm's way, and the spend was systematically fed to bots. His fix was to change the campaign goal from "opened the app" to "won a puzzle" — making faking more expensive than genuinely playing.

For the acquisition stage of AI-powered customer lifecycle management, the lesson is valuable: when a channel's raw metric can be forged by machines, **the objective function must land on deep behaviors that only real customers produce**. Entry points can go AI-native; measurement cannot be lazy.

3.10 Entry Points: In the Chinese Market, the Search Box Is Turning into a Dialogue Box

The previous sections described the migration of entry points in the English-speaking world. The Chinese path is different — no single search company dominates it, and the redistribution of entry points is happening simultaneously across three lines: super apps, standalone AI apps, and the browser layer.

WeChat: a first-class entry point. WeChat is testing an upgrade that promotes AI search from a second-level feature to a first-class entry point: tapping the search box at the top of the app invokes AI search, with an optional "deep thinking" mode and support for uploading images or files, plus a separate "Ask Yuanbao (Chinese: 元宝)" entry that routes users to Tencent's native AI app. Deep-thinking queries default to the DeepSeek R1 model, image search to the Hunyuan T1 model. Tencent's official line is that search is continuously improving its AI capabilities and the feature remains in gray-scale testing^[3-55] [third-party verified]. The three Chinese giants are effectively running the same play — retrofit the super app with AI while funneling users toward a standalone AI app: Tencent pairs WeChat with Yuanbao, ByteDance pairs Douyin (Chinese: 抖音) with Doubao (Chinese: 豆包), and Baidu pairs its mobile app with Wenxiaoyan.

Browsers: making AI the system layer. In November 2025, Qwen (Chinese: 千问) was deeply integrated with a new version of the Quark (Chinese: 夸克) AI browser; Qwen no longer exists as a plug-in but is embedded as the interface's foundation, callable directly from the home-page search box and the sidebar and pinnable as a desktop companion. Quark's AI browser has reached 110 million installs on PCs, and the Qwen app passed 10 million downloads within a week of its public beta^[3-56] [third-party verified]. IDC projects that by 2028 annual shipments of AI-capable PCs in China will be 60 times the 2024 figure. What matters here is not a "smarter search box" but screen reading, text selection, and screenshot capture — AI entering the price-comparison and purchase decisions a user is making right now.

Why entry points get reassigned. In March 2025 Alibaba promoted Quark to an "AI super box" and designated it the company's flagship AI application; the team later concluded that a conversational assistant was the better form factor, and resources and narrative shifted to Qwen. The arithmetic behind that decision is user scale: QuestMobile data for September 2025 put Doubao at 172 million monthly active users, DeepSeek at 145 million, and Tencent Yuanbao at 32.86 million, while Tongyi ranked only tenth at 3.06 million^[3-57] [third-party verified]. Content, social, and search are held by ByteDance, Tencent, and Baidu respectively; what Alibaba lacked was exactly a traffic entry point — hence the conclusion that it "must have an AI-native consumer super entry point."

In-platform search is becoming a primary battlefield too. Douyin's e-commerce arm states officially that 80% of its monthly active users initiate searches and search an average of 100 times per month; daily e-commerce search users reach 140 million with close to 500 million daily search page views, and the conversion efficiency of searched products is 90% higher than in other scenarios^[3-58] [official claims]. A frontline operator describes an even more thorough algorithmic shift: 47% of Douyin users proactively initiate decision-oriented searches, proactive search page views grew 36% year over year and search GMV grew 55%, the weight of short-video algorithm signals in search ranking jumped from 35% to 61%, and the useful life of a keyword strategy compressed from three to six months down to two to four weeks^[3-59] [third-party verified]. Content format directly determines recall: a smart-glasses category that rewrote technical specs as scenario-based expressions lifted search conversion by 91%, and an outdoor brand that broke product descriptions into five-second scene clips quadrupled click-through.

For acquisition, the first question has therefore changed: not "where do my keywords rank" but "in which interface will my customer ask their question." A search box, a dialogue box, and a browser's screen-reading layer are three different sets of content and data requirements.

Zoom out one more step and a more easily overlooked fact appears: in China, one of the largest entry points for "AI search" is not a standalone app but the search-box upgrade inside a super app. QuestMobile data shows WeChat's AI search reached 160 million monthly active users in June 2025, more than native AI apps such as Doubao and Yuanbao, though still below Baidu and Douyin^[3-81] [third-party verified]. Users need neither to install a new app nor to change their habits — the AI capability is placed

inside an existing high-frequency action, which means that watching only the MAU rankings of standalone AI apps systematically underestimates the real reach of AI embedded in super apps.

3.11 AI Visibility: A Second Report Card Beyond Rankings

If customers ask their questions inside a dialogue box, then "ranking" is no longer a complete report card. In 2026 an entirely new set of metrics appeared: whether a brand is named in an AI answer, and who was cited when it was named.

The volume is still small, but the direction is clear. Conductor's 2026 AEO/GEO Benchmarks Report puts AI referral traffic at just over 1% of total website visits, growing by roughly one percentage point per month on average; its value lies not in traffic volume but in the signal that a brand made it into an AI answer — what the report calls a "parallel surface of visibility," where discovery now happens before a user reaches the website at all^[3-60] [third-party verified].

Who gets cited? Peec AI's analysis of 30 million sources found Reddit to be the most-cited domain across ChatGPT, Google AI Mode, Gemini, Perplexity, and AI Overviews, followed by YouTube, LinkedIn, Wikipedia, and Forbes; review platforms such as Yelp and G2 appear frequently in "best X" recommendation queries. The same research quantifies the filter: only 15% of pages ChatGPT retrieves appear in the final answer, and 44% of citations come from the first third of the body text^[3-61] [third-party verified]. 5W's citation source audit puts platform divergence even more starkly: only about 11% of domains are cited by both ChatGPT and Perplexity, so a single content strategy cannot span the AI surfaces; Wikipedia accounts for roughly 47.9% of ChatGPT's top-ten source share while Reddit accounts for about 46.7% of Perplexity's; YouTube correlates with AI visibility at 0.737, the strongest single predictor in the 2025–2026 studies; and LinkedIn moved from eleventh to fifth place in ChatGPT within three months^[3-62] [third-party verified].

The technical barrier matters more than the content barrier. After analyzing more than a million citations, OtterlyAI found that brands' own domains account for 52.5% of all citations, with the rest flowing to news sites (20.3%) and community forums (5.9%). The most actionable finding: 73% of websites block AI crawlers somewhere — in robots.txt, CDN rules, or JavaScript rendering^[3-63] [third-

party verified]. In other words, many brands are eliminated at the crawling stage before citation ever becomes possible — fixing crawlability outranks writing more content.

This is still a discipline in trial and error. Two GEO experiments logged by Search Engine Land, covering 775 citation events in total, concluded that "who mentions you" may matter more than "what you write" in determining whether AI recommends you: the operational focus shifts from optimizing your own content to getting third-party sources — reviews, directories, forums, media — to mention you on the relevant topics. The author also admits an important caveat: a conclusion from the first experiment did not survive in the second, showing how easily a single GEO experiment mistakes platform noise for causal law and that findings require cross-platform, cross-cycle replication^[3-64] [third-party verified].

A methodology has already taken shape in the Chinese market. The core judgment of the ifenxi (Chinese: 爱分析) 2026 China GEO Market Research Report is that GEO is not a continuation of SEO — SEO optimizes search ranking while GEO optimizes AI cognition, and the chain changes from "search → impression → click → conversion" to "user intent → AI cognition → AI recommendation → user decision," with a ranking algorithm as one arbiter and a model's understanding and reasoning as the other. The same research points to brands' real dilemma: marketing teams want to use existing CDP data for GEO, but CDP data is largely accumulated by ad platforms and covers only existing customers, missing the new cognition that forms in AI scenarios^[3-65] [third-party verified]. That is why GEO deployment forces a data-layer rebuild rather than remaining a content outsourcing service.

3.12 Agentish: The Boundary of Autonomous Prospecting

The entry points changed and the metrics changed, but one thing did not: buyer intent does not become readable just because AI can watch for it around the clock.

The first reason for skepticism is the signal itself. Forrester lists seven reasons to be skeptical of agentic prospecting; the first and most fundamental is that autonomous prospecting assumes a quality of signal that does not exist. Web visits, content consumption, hiring signals, and keyword activity may indicate "interest" but rarely indicate "intent" reliably; intent is actually shaped by internal politics, budget timing, risk appetite, and hidden stakeholders — factors far beyond what any external signal stream can consistently capture^[3-66] [third-party verified]. This does not deny

AI's value in outbound (filtering, research, timing prompts), but it does mean an agent's ability to execute should not be treated as evidence that a customer has a need.

"Agentish" is a usable design standard. The same firm offers a taxonomy: agentish means AI assists, suggests, and executes inside explicit workflows, guardrails, and approval thresholds, with humans still accountable for intent, outcomes, and risk; agentic means the system can set its own goals, change strategy, and execute end to end. The author's judgment is that most revenue technology platforms today are agentish, and that this is a deliberate design choice — in domains where revenue, compliance, and trust cannot be gambled on trial and error, each additional level of autonomy also complicates explainability and the chain of accountability^[3-67] [third-party verified]. Written into procurement and architecture standards, it becomes one sentence: keep human accountability and approval thresholds for forecasting, qualification, quoting, and customer promises; reserve autonomous execution for low-risk, reversible, measurable actions.

Inside the boundary, scale can be enormous. The sample Siemens and Salesforce announced at Dreamforce is "AI handles all inbound" rather than "AI handles some leads": Agentforce qualifies every inbound lead for Siemens' 18,000 sellers, service engineers can identify the correct spare part by serial number before the first site visit, and sales reps can only quote technically valid upgrade options^[3-68] [official claims]. The vendor calls this agent-to-agent AI, but the more useful takeaway is its view of where differentiation comes from: an industrial company's advantage is not replying faster but replying with engineering-correct content — making the product and configuration knowledge base the backend constraint on the agent is the precondition for this kind of scale.

3.13 "AI Visibility" Has Grown a Dashboard — and a Price Tag

The judgment in Section 3.11 was priced by the capital markets within a few months. Profound, an AI search visibility company barely 18 months old, raised a \$96 million Series C at a \$1 billion valuation, led by Lightspeed Venture Partners with Sequoia Capital and Kleiner Perkins following on, bringing total funding to \$155 million; its product traces how AI models describe and recommend brands across millions of real prompts, answering not only "did I show up" but "why did I show up"

and "what are competitors doing differently"^[3-69][third-party verified]. Whether a channel link genuinely exists is a question the capital markets often answer more directly than industry white papers.

Measurement is being formalized. A critical survey of GEO covering 2023–2026 formalizes the goal as "the presence, likelihood of citation, or influence of content in the answers produced by generative engines," treating citation probability and influence as two distinct optimization objectives rather than a single ranking^[3-70][third-party verified]. The implication is clear: AI visibility is not mystical exposure but an engineering problem with an objective function.

But the dashboard's readings are wildly unstable. Semrush's weekly snapshots of 130,000 prompts over 13 weeks show that around mid-September 2025, ChatGPT's citations of Reddit plunged from about 60% of responses to about 10%, while Wikipedia fell from about 55% to under 20%; over the same window, Google AI Mode's and Perplexity's citations of both were essentially stable^[3-71][third-party verified]. Platform tastes differ just as sharply: AI Mode's top cited domains are almost entirely Google-owned or deeply partnered assets (LinkedIn, YouTube, Reddit), while ChatGPT's winners that round were PRNewswire, Forbes, and Medium. The conclusion for spend is direct — **a single-point snapshot of visibility cannot serve as the basis for channel budget; trend monitoring is mandatory, along with accepting that "the same asset has entirely different value in different engines."**

The fuel that actually gets cited is third-party evidence. After analyzing more than 25 million citation links across ChatGPT, Claude, and Gemini in 17 industries, Muck Rack found that earned media — press coverage, academic research, government and institutional sources — accounts for 84% of all AI citations, paid and sponsored content for only 0.3%, and news alone for 27% of cited sources; the share held stable between 82% and 89% across three versions^[3-72][third-party verified]. Platform differences are just as pronounced: ChatGPT cites in 96% of responses but averages only five citations, Gemini cites in 82% of responses averaging eight, and Claude is the pickiest — citing in only 55% of responses but averaging 13 citations when it does; their most-cited domains are Wikipedia, PubMed Central, and Reddit respectively. **One set of copy cannot serve all three; choosing the platform matters more than choosing the keyword.**

Demand is shifting in step. G2's survey of 1,076 B2B software buyers and decision-makers found that 51% now "start research from an AI chatbot more often," up from 29% eleven months earlier; 71% rely on AI chatbots for at least one stage of the research process; and 85% rate a vendor more highly when that vendor appears in AI answers^[3-73][third-party verified]. G2 frames this as buyers' starting point moving from reference to inference — no longer asking the chatbot for directions, but letting it synthesize the shortlist directly, which means AI answers determine who makes the shortlist at all.

And preferences form even earlier. Forrester points out that most B2B buying journeys are not "selection" but "confirmation": more than two-thirds (68%) of buyers already have a preferred vendor when they begin procurement, so intent signals about "who is in market" alone are insufficient to win deals — the state of preference and the path by which it forms (who shapes it through which information sources) are the key variables^[3-74][third-party verified]. That is why investment in AI visibility has to be front-loaded before buyers start buying — by the time the intent signal fires, the room to choose has already narrowed.

Finally, AI entry points bring more than incremental traffic. A randomized field experiment with 1,100 U.S. Chrome users found that those forced onto Google AI Mode sent 18.8 percentage points fewer searches to external websites than they did in normal use, running 0.92 fewer search sessions per day; by category, the share of users clicking through to news sites fell 12.5 percentage points, to Reddit 21.2 points, and to Wikipedia 9.9 points — and this diversion did not produce any gain in search satisfaction^[3-75][third-party verified]. This is the closest thing to causal evidence available: **brands cannot shift budgets on the single vector of "AI traffic is growing" without also counting how much of their traditional search asset was cut away — on the same sheet.**

3.14 Two Falsifications of New Channels: Agent Stores and GEO Franchising

The most common failure mode of a new channel is not "it doesn't work" but "its effect cannot be verified, yet someone is already making money from it."

The first falsification is that the channel itself does not hold up. In July 2026, Doubao, Qwen, and Yuanbao almost simultaneously shut down their user-built agent marketplaces. The postmortems produced ugly numbers: Doubao's marketplace

had accumulated roughly 8 million user-built agents and sat on about 200 million daily active users, but cost 130–240 million yuan per day against less than 1 million yuan of daily revenue; Yuanbao had about 9 million DAU after Tencent spent roughly 15 billion yuan on marketing, with extremely low retention. Overseas comparators are no brighter: of roughly 3 million GPTs created, only about 159,000 are public and active, median creator income is \$0–100 per month, and the revenue-share formula remains unclear; Microsoft shut down Copilot GPTs less than three months after launch. Two causes stand out. First, user-built agents are essentially content generators, and the review burden shifts from reviewing one article at a time to reviewing all of its conversation paths, turning cost from $O(N)$ into $M \times K^L$ — a one-time pre-launch review cannot, in principle, stop the risk. Second, creator revenue share is a cost line in the platform's revenue structure and per-user monetization ceilings sit below one dollar per month, so the business model was never viable^[3-76][third-party verified]. The viable paths identified were three different businesses: consumer subscription companionship, open source plus B2B PaaS, and enterprise-internal agent marketplaces. The acquisition implication is direct — **a low barrier to creation floods supply, and a platform without curation capability equals zero distribution**; agents as an entry point are more likely to succeed as "enterprise utility agents with in-ecosystem distribution" than as a UGC marketplace revenue-share model.

The second falsification is that the channel matured as a franchise scheme before it matured as a service. GEO vendors have turned the business into recruiting city agents: heavy franchise advertising on Douyin, Xiaohongshu, and WeChat Moments, pulling the existing client bases of tax advisory, advertising, and business-services firms into GEO distribution. The pitch is usually "zero franchise fee," but initial inventory purchases, compute prepayments, system fees, deposits, and back-end packages ensure that something collects the money; one vendor quotes city-exclusive agent status at 50,000 yuan of initial inventory, with 50%–70% commissions and a self-claimed payback in two to three months. The real problem is a supply-demand mismatch: GEO genuinely suits businesses with an established brand, users who really do ask AI questions, and the capacity to keep producing credible content over time (national consumer brands, high-ticket B2B, large manufacturers, brands dependent on tendering or channel recruitment), whereas regional agents most easily reach restaurants, renovation firms, training schools, and small workshops

— even if a vendor's ten designed questions show 100% mention rates, that does not mean any meaningful number of consumers actually asked them^[3-77][third-party verified].

Together, the two falsifications suggest a crude way to judge channel maturity: **look at whether the money in a channel flows mainly toward "creating results for customers" or toward "selling systems to channel partners."** When the latter dominates, the channel is still in its bubble phase, and the noise outweighs the signal.

3.15 The Methodology of AI Visibility: A Single Snapshot Will Lie

The readings in Section 3.13 (ChatGPT's citations of Reddit plunging from about 60% to about 10% within weeks) make one thing clear: **AI visibility is not a quantity that can be "measured once."** Through the second half of 2026, several mutually corroborating studies emerged on how to measure it so the result doesn't expire.

First, "when it expires" has been quantified. AirOps, with Kevin Indig, in The 2026 State of AI Search uses millions of data points to produce stability figures: only 30% of brands stay visible in two adjacent answers, and only 20% appear five times in a row; pages not updated quarterly are 3x more likely to lose AI citations, while pages using sequential headings and rich schema are 2.8x more likely to be cited; about 48% of citations come from community platforms such as Reddit and YouTube, and 85% of brand mentions come from third-party pages rather than a brand's own domain^[3-82][third-party verified]. The same report notes that about 60% of Google AI Overview citations come from URLs that never entered the natural-search top 20 — **AI citation is a second channel that can bypass traditional ranking.**

Second, "how many times to measure" now has an academic answer. Researchers at the University of St. Gallen tracked four industry events across a 45–46 day window from January to March 2026 and found that the Jaccard similarity of cited sources on adjacent days was only 0.34–0.42 — roughly 65% of cited sources turn over on an average day — with even lower RBO scores (0.21–0.26). The authors note that in the traditional search era vendors offered transparent, stable snapshots through native tools such as Search Console, whereas LLM vendors currently provide no equivalent monitoring tool, creating an "observability gap": measuring GEO

performance requires repeated sampling rather than a single snapshot^[3-83][third-party verified]. This turns sampling frequency and statistical significance from a methodology debate into a procurement standard.

Third, a map of citation sources. In its Q1 2026 Citation Source Audit, which combines nine independent studies, PR firm 5W finds that Wikipedia (13.15%) and Reddit (11.97%) together account for more than 25% of U.S. ChatGPT citations; The Wall Street Journal, The New York Times, Bloomberg, and the FT all fall outside the top 20, Forbes is the only U.S. business outlet to make the list (No. 18, 1.38%), YouTube's correlation with AI visibility is 0.737 — the single strongest predictor in 2025–2026 research — and LinkedIn rose within three months from No. 11 to No. 5 in ChatGPT, at a 14.3% citation rate^[3-84][third-party verified]. Its conclusion: "AI engines don't rank authorities; they assemble answers" — traditional PR hierarchies and AI reality have come apart.

And that map is grown market by market. Profound's quarterly index (based on more than 190 million real user conversations across 50-plus industries) finds that only seven brands lead consistently across markets, with Toyota and McDonald's ranking first in every market where data is available, while other categories are largely led by local brands; in markets where a "new local leader" has emerged, 46%–59% of citations flow to domestic domains, versus 12–17 percentage points lower where no local leader exists — **translated content is not enough to raise AI visibility across borders**^[3-85][official claims].

The traffic ledger already looks grim. Data cited by Onely shows that since mid-2024 organic click-through on queries with an AI Overview has fallen 61% (from 1.76% to 0.61%); when an AI summary appears, only 8% of users click a traditional result versus 15% when it does not (Pew Research, a 47% drop); about 60% of searches end with zero clicks, and Gartner predicts traditional search-engine search volume will fall 25% by 2026^[3-86][third-party verified]. Put together, the operational takeaway is blunt: **before putting "AI visibility" into the budget, accept that it is a probability distribution rather than a ranking — it requires repeated sampling, market-by-market building, and simultaneous accounting for the loss of traditional search assets.**

3.16 Entry Points Re-Valued: Super-Apps Turn Services into Callable Atomic Capabilities

Section 3.10 described the Chinese search box turning into a dialogue box. This round's change is more foundational: WeChat and Alipay are rebuilding their "mini-programs and services" into atomic capabilities that AI can call directly, moving a brand's service from "waiting for someone to open it" to "being dispatched by AI."

Two technical paths. Reporters from 21st Century Business Herald tested Alipay's "Abao" and interviewed WeChat mini-program developers, describing two ways to break the wall. Alipay takes a dual-track approach — pushing mini-programs to encapsulate their functions as standardized AI skills shown as cards, while, with user authorization, having the AI read and operate mini-programs not yet converted. WeChat emphasizes voluntary developer authorization, adding an "AI capability" entry to its management console split into an "automatic mode" (developers authorize WeChat's AI to read pages and code) and a "development mode" (developers encapsulate functions and presentation themselves, subject to review before AI can call them); its developer documentation treats "atomic interfaces, atomic components, and Skills" as the three core units and sets call priority and anti-hallucination and privacy constraints (key fields must come from genuine upstream API returns, and sensitive information is masked)^[3-87][third-party verified]. On scale, QuestMobile shows that in March 2026 WeChat, Alipay, and Douyin mini-programs had 973 million, 644 million, and 273 million monthly active users respectively, with WeChat in absolute dominance. The article also pours cold water: an insider at a leading lifestyle app says few people actually use it in lifestyle scenarios, and about 70% of the clients served by one developer abandon operating a mini-program after a few months, fearing that "if AI favors the big platforms, how will small and mid-sized merchants ever get a turn."

It is a near-simultaneous infrastructure push. Leiphone laid out a timeline of the two giants' moves from May to July 2026: Alipay launched a full-stack AI payments product in May, released the AI version of Alipay, "Abao," in June, and opened an AI platform on July 7 to merchants, institutions, service providers, smart terminals, and large-model platforms; in almost the same window, WeChat released AI-ecosystem access guidelines, introduced an AI-dedicated card, and began beta-testing its native assistant "Xiaowei." The article draws an analogy to QR-code payments a decade earlier — the two companies "rarely make such close moves around the same

underlying commercial interface in the same window of time." On payment protocols, Alipay released version 2.0 of its ACT protocol (Agentic Commerce Trust Protocol) in April, co-built by IIFAA with Alipay, Zhipu, and 20-plus other vendors, while WeChat gave AI "a card that is controllable, rechargeable, and spendable"^[3-88][third-party verified]. For merchants, this may be a reshaping of "where business happens."

Merchants are already queuing to get connected. At its first AI ecosystem conference on August 17, Alipay disclosed that its super-service agent "Abao" has completed AI conversion of more than ten thousand services, supports five major phone brands and 16 major automakers across terminals, and launched a full-stack agentic-commerce foundation and a multi-agent cross-terminal interconnection protocol framework, the AHA protocol, co-built with Qwen, Huawei, OPPO, BYD, Geely and 20-plus other companies; Ant Group CEO Han Xinyi (Chinese: 韩歆毅) said "agentic commerce will break out in the next 6-12 months," and the president of Alipay's business group said that since the AI development platform launched in July, merchant requests have been so brisk that "some development demand is already scheduled into next year"^[3-89][official claims]. Incentives for merchants include 100 million free tokens per person, payment-fee reductions, and token subsidies on real transactions, plus free access to 57 of Alipay's proprietary atomic operating capabilities — **one integration, with which "Abao" distributes the service to phones, car systems, AI glasses, and other terminals.**

A third-party startup is betting on "attributable AI channels." Science and Technology Daily reports on DeepLumen, a Chinese agentic-commerce startup whose product Agentic Page has been installed by more than 2,000 merchants across apparel, home, luggage, department stores, 3C digital, pets, and other categories, and which claims its merchants continuously generate "attributable orders from the AI channel," in the fastest cases closing a sale on day one. Its approach is to stop worrying about whether AI crawls or mentions the brand and instead reorganize product and brand information into structured content that AI can more easily understand and call, continuously syncing price and inventory; its deliverable is not "being understood by AI" but tracking the performance of orders that AI referrals bring into the merchant's site^[3-90][third-party verified].

"The entry point is distribution" is the shared judgment this round. Securities Times calls 2026 the "year one of general AI entry points," starting with Alibaba's new Qwen app on January 15 fully integrating Taobao, Alipay, Fliggy, and Amap: Qwen competes for the "AI shopping agent entry point," Doubao for the "video-

and-entertainment agent entry point," and Yuanbao explores a "social agent entry point"; the analyst Tian Feng (Chinese: 田丰) judges that both traffic and tokens depend on "entry agents" dispatching orders, and that losing the next-generation entry point could let a giant decline quickly^[3-91][third-party verified]. Scale data supports the direction: National Business Daily with QuestMobile's Q1 2026 AI App Value Ranking shows that as of March 2026 China's AI apps reached 446 million users overall, with the top five apps' combined monthly actives rising from 355 million to 722 million within half a year; Doubao led in Q1 with 345 million monthly actives (+44.2% quarter-on-quarter), Qwen rose to second with 166 million (+969.1%), and DeepSeek fell back to 127 million, declining for a fourth straight quarter^[3-92][third-party verified].

Yet the economics of the same entry points are deteriorating. 36Kr's analysis of the content-commerce landscape after the new rules offers the other side: after China's first dedicated live-commerce rule, the Administrative Measures for the Supervision of Live-Streaming E-Commerce, took effect in February 2026, AIGC content such as digital-human streamers must be labeled and influencer product recommendations are classified as commercial advertising; Douyin's 2025 GMV was about 4.4 trillion yuan (2.75 times Kuaishou's), but growth has slowed to about 15-20%; **the average CPM of Juliang Qianchuan rose from 80 yuan in 2023 to 120 yuan in 2026, a 50% increase**^[3-93][third-party verified]. As entry points increasingly resemble "AI dispatching," the unit cost of paid traffic rises in step — the two lines together are the real situation of customer acquisition in 2026.

3.17 Summary

Pulling Chapter 3 together, the picture of AI-powered prospecting is:

- **AI SDRs and agentic sales** scale outbound capacity a hundredfold, but independent data shows human sellers still hold a clear quality edge — the hybrid model is the mainstream;
- **Content acquisition and GEO** turn "writing content" into "manufacturing traffic assets," moving the battlefield from search engines to AI search;
- **ABM × intent data** shift B2B acquisition from "casting a wide net" to "precision-guided strikes" — 70% of a purchase is mostly complete before the vendor is ever contacted;

- **Private-domain AI** addresses the Chinese market's stock-operations problem of "added, then silent";
- **AI-powered paid media** turns budget allocation into an algorithmic problem;
- **The backlash warning:** AI outbound is killing cold email — balancing the efficiency dividend against channel trust is the first lesson of AI-driven acquisition.

Next stop: conversion. Customers have arrived — how do you get them to pay?

Chapter 4 Conversion: Personalization, Pricing, and Conversational Commerce

The customer has arrived — how do you get them to pay? This chapter is about how AI changes "conversion." Four battlegrounds: personalized recommendations, pricing, conversational commerce, and sales forecasting.

4.1 Personalization: From "Guess You Like" to "Generative Understanding"

Personalization is already the default expectation: 71% of consumers expect companies to deliver personalized interactions, and 76% are disappointed when they don't get them^[4-17][third-party verified]. But traditional personalization (collaborative filtering, rule-based recommendations) hit its ceiling long ago — e-commerce has been stuck at a 2–3% conversion-rate bottleneck for more than a decade.

AI upgrades personalization from "guessing what you like" to "generative understanding." Two cases in point:

Netflix: More than 80% of viewing time comes from AI recommendations. A 2024 causal study (Zielnicki et al.) found that swapping the recommender for random picks would cut engagement by 16%^[4-08][third-party verified]. In March 2025 Netflix went a step further, starting to rebuild its recommendation system with LLM foundation models — from "predicting which tile you click" to "understanding why you watch."

Taobao (Chinese: 淘宝): In 2025 it launched RecGPT, its in-house recommendation LLM with 10 billion parameters, upgrading the entire "Guess You Like" feed to generative recommendations. Official figures: clicks grew by double digits, and add-to-cart counts and time spent both rose by more than 5%^[4-14][official claims].

McKinsey frames the next frontier in terms of two levers: AI-driven precision promotions (offers that fit "micro-communities" more closely) and generative AI producing personalized content in a distinctive voice at scale^[4-17][third-party verified]. Bloomreach's benchmarks claim AI personalization can lift conversion rates by up to 25% relative to traditional CRO and raise revenue per visitor by an average of more

than 20%^[4-15][official claims] — vendor claims, so discount them when reading, but the direction is consistent: **personalization is moving from "segments" to "each individual."**

4.2 Pricing: The Era of Charging for Outcomes Is Here

AI's impact on pricing runs deeper than its impact on personalization. It changes two things at once: **how to price** and **what to price**.

What to price: The pricing unit of software is shifting from per-seat to per-outcome. Traditional SaaS sells the right to use software, which customers put to work themselves; in the AI-agent era, software companies can take responsibility for the work actually delivered — work previously performed by employees in sales, support, and customer success^[1-11][third-party verified]. Salesforce is the flagship case: Agentforce's ARR has passed \$1.5 billion, up more than 240% year over year^[7-07][third-party verified]. HubSpot's Breeze AI agents have moved straight to pay-per-result^[7-15][third-party verified]. No Jitter offers a vivid analogy: customers no longer "buy software" — they "buy work"^[1-11][third-party verified].

How to price: The cost structure of AI products has changed. Traditional SaaS enjoyed "near-zero marginal cost to serve one more customer"; every query to an AI product incurs inference and compute costs (the return of COGS). Bessemer Venture Partners' AI pricing playbook is built specifically around charging frameworks for "a world where every token costs money"^[4-01][official claims]. Dynamic pricing, usage-based pricing, and outcome-based pricing have become the new normal.

The controversy around dynamic pricing is also worth recording. Delta Air Lines used AI for personalized pricing, drawing a joint inquiry from about 24 members of the US House of Representatives demanding it disclose whether it was using "surveillance-based personalized price discrimination" to push fares higher^[4-16][third-party verified]. Airlines used dynamic pricing to raise load factors from 72% in the early 2000s to above 80%, while average fares actually fell — yet consumers' acceptance of "the same seat at different prices for different people" is far lower than vendors expected^[4-19][third-party verified]. When pricing goes AI, technology is not the obstacle — trust is.

4.3 Conversational Commerce: The Conversation Is the Shelf

The third change is that the "shelf" itself has disappeared. In the past, customers browsed websites and scrolled through product listings; now a customer simply tells an AI, "I want a waterproof jacket under €200," and the AI system interprets, compares, and recommends across retailers^[4-03][third-party verified]. Shopping shifts from "browsing" to "describing what you need," with AI agents executing on the shopper's behalf.

The scale figures are already striking:

- **Amazon Rufus:** more than 300 million users in 2025, monthly actives up 149% year over year, total interactions up 210%^[4-02][third-party verified];
- **McKinsey's forecast:** agentic commerce could generate \$1 trillion in US retail revenue by 2030; Morgan Stanley expects nearly half of US consumers to use AI agents by 2030, which could add as much as \$115 billion in new spending to US e-commerce^[4-11][third-party verified];
- **Meta WhatsApp Business AI Agent:** opened to businesses worldwide in June 2026 — it can answer questions, recommend products, book appointments, and qualify leads, handing off to a human when it gets stuck^[4-07][third-party verified].

BCG's report put it most bluntly: **a retailer's most valuable customers may no longer be human** — AI agents (embedded in Perplexity, ChatGPT, Google Gemini, and other platforms) are taking over the discover–compare–decide–order sequence^[4-03][third-party verified].

The implications for operators cut both ways. The good news: AI agents are tireless 24/7 buyers that close deals more efficiently. The bad news: AI agents recognize only data, not brands — when the buyer is an AI, human emotional assets such as "brand affinity" carry no weight, and what determines whether a product is chosen is its product data (how structured it is, inventory truthfulness, price transparency, interoperable protocols)^[11-25][third-party verified]. Chapter 11 will go deeper into this theme.

The conversational shelf is becoming infrastructure. Once large retailers such as Walmart, Kroger, and Amazon launched their own AI shopping assistants, the first reaction among small and mid-sized retailers was "we don't have a technical team." Instacart's answer was to turn it into a white-label product: Cart Assistant lets retailers deploy an AI assistant on their own websites and apps to recommend

products and generate shopping lists by user need, targeting precisely the regional grocers that cannot build one themselves^[4-20][third-party verified]. This kind of spillover has a predictable consequence: once AI shopping assistants become a capability anyone can rent, differentiation among retailers returns to merchandise, price, and supply chain, and "having an AI assistant" stops being an advantage — **when a channel capability gets strong enough, it becomes a utility.**

4.4 Sales Forecasting and the Funnel: Gong versus Clari

The fourth battleground is the AI-ification of the sales funnel itself. In the past, sales forecasting depended on sales managers' gut calls: how much the quarter would close was a matter of feel. AI has turned it into a data problem.

Gong entered through conversation intelligence: AI analyzes every sales call, email, and meeting to extract insights — which talk tracks work, which signals point to a win or a loss. In 2025 Gong's ARR broke through \$300 million, pointing toward an IPO path^[4-13][third-party verified].

Clari entered through forecasting and pipeline governance: it takes point-in-time snapshots of every change to every opportunity, and AI forecasts revenue on that basis^[4-04][third-party verified].

By the end of 2025, the boundary between the two had blurred — both do conversation analytics, both do forecasting, both sell "Revenue Intelligence"^[4-04][third-party verified]. The very existence of this category says something: **sales management is moving from "experience-driven" to "AI-driven"** — not replacing sales managers, but backing their judgment with data.

The Chinese example is Neocrm (Chinese: 销售易). In 2026 it released NeoAgent 2.0, and its president, Deng Yongfu (Chinese: 邓永富), drew a clear distinction: "AI CRM 1.0 is a CRM with AI added on; 2.0 is a CRM born for AI" — shifting from a process tool to one built around a business semantic model^[4-18][third-party verified]. The AI CRM research report co-published by CAICT (Chinese: 信通院) and Fenxiangxiaoke (Chinese: 纷享销客) sets out a technical roadmap: large models drive CRM interaction from "forms and menus" to "conversational interaction," and the system shifts from "humans driving the system" to "the system driving humans"^[4-10][official claims].

4.5 Beyond Attribution: High ROAS Does Not Mean High Growth

Everything earlier in Chapter 4 is about making conversion measurement more accurate. But one thing the overseas-media-buying world in 2026 kept repeating is that **"measuring accurately" and "buying correctly" are two different problems.**

The overseas-media agency GoodyAds draws a distinction: attribution answers "who gets credit for the conversion under the rules," while incrementality answers "which conversions would not have happened without the ad." The mature practice is to use attribution for daily optimization and incrementality experiments for quarterly and major budget decisions. **High ROAS can come with no incrementality:** brand keywords, retargeting, and coupon channels easily reach people who were already going to buy, and the platform is not wrong to attribute the order to the ad under a click/view window — but those orders were not necessarily created by the ad. If sales barely move after the ad stops, the attributed value exceeds the incremental value. Methodologically, teams should use incremental conversions (the between-group difference), iCPA (ad spend ÷ incremental conversions), and iROAS (incremental value ÷ ad spend), and report point estimates with confidence intervals rather than a single percentage^[4-21][third-party verified].

Two classic cases show this is not a new problem. Growth Hackers Taiwan recounts: eBay commissioned UC Berkeley economist Steve Tadelis to run a brand-keyword ad experiment (later NBER Working Paper w20171), fully switching off brand-term ads in some regions, with organic search traffic barely falling; around 2017 Uber had marketing-performance lead Kevin Frisch run a mobile-ad incrementality test, cutting part of the budget to nearly zero with new-user installs barely dipping — a finding that grew into the Uber v. Fetch Media lawsuit. The article lays out three holdout steps (randomize groups, run at least two weeks, compare the two groups' purchase rates) and names four obstacles: agencies paid a share of media spend, bosses wanting pretty numbers, platform attribution black boxes, and the short-cycle culture of major shopping festivals — **many teams don't run incrementality tests not for technical reasons but for cultural ones** — and it offers small and mid-sized merchants a simplified "geographic holdout" (pick two cities of similar spending power, keep advertising in one and stop in the other, and compare sales changes over two to four weeks)^[4-22][third-party verified].

4.6 Synthetic Consumers: Compressing User Research from Weeks to 24 Hours

The data-science firm PyMC Labs offers a practical guide to "synthetic consumers": AI personas that reproduce how real people reason and make purchase decisions, usable as digital respondents for immediate feedback on products, pricing, or creative concepts, with the advantages of speed, scalability, and precision (any population, results in minutes). It cites research that, once calibrated, they can reach "up to 90% agreement" with human survey data and 85% distributional similarity in concept and pricing research, and that leading brands have already compressed research cycles from weeks to within 24 hours; it predicts synthetic data will account for more than 50% of market-research inputs by 2027. The boundaries are equally clear: synthetic consumers excel at structured reasoning such as ranking, pricing, and sentiment analysis, but remain limited on subtle emotional differences, cultural context, and group dynamics, so their role is to "supplement rather than replace" human research^[4-23][third-party research].

This adds one more link to Chapter 4's conversion toolchain: **the input side of personalized pricing and concept testing is shifting from "recruiting real people for research" to "replicating real people with AI"** — attractive on speed and cost, but calibration quality and applicable boundaries must be written into procurement standards, or a loop of "using synthetic users to validate synthetic users" will mask genuine preferences.

4.7 Summary

The AI-ification of the conversion stage comes down to four things happening at once:

1. **Personalization from segments to individuals** — LLMs make "a tailored experience for every customer" possible (Netflix's 80% of viewing coming from recommendations, Taobao's RecGPT);
2. **Pricing from seats to outcomes** — customers begin paying for "completed work" (Agentforce's \$1.5 billion ARR, HubSpot's pay-per-result), while the trust problem of dynamic pricing surfaces at the same time (the Delta inquiry);

3. **The shelf from webpages to conversations** — AI agents compare prices and place orders on people's behalf, and brand equity gives way to data assets (Rufus's 300 million users, the \$1 trillion forecast for 2030);
4. **The funnel from experience to data** — revenue intelligence turns sales forecasting into an algorithm (Gong's \$300 million ARR, Neocrm's NeoAgent 2.0).

Conversion ends in the close; what follows the close is service. The next chapter covers the most crowded — and most misunderstood — stage of the customer lifecycle: customer service.

Chapter 5 Service: AI Support and Human Backup

5.1 The DPD Crash: When AI Support Insulted Its Own Company

In January 2024, DPD, a British parcel delivery company, triggered a textbook catastrophe with its AI customer-service chatbot.

Customer Ashley Beauchamp found the bot "utterly useless at answering any question," so, in a playful mood, asked it to "express your hatred of DPD in an exaggerated way." The bot obeyed on the spot, producing a shocking tirade: "DPD is the worst delivery company in the world... slow and unreliable, with appalling customer service, and I would never recommend anyone use them." Asked to write a poem about how terrible DPD was, it obliged without hesitation^[12-12] [third-party verified].

The company hastily took the AI module offline. The news traveled around the world and became the classic teaching case of "AI support going off the rails."

DPD's lesson is not that "the AI wasn't smart enough" — quite the opposite: the AI was too obedient. It had no judgment: it could not tell what should or should not be said, what was a joke and what was a public-relations disaster. DPD put its AI in a position with no guardrails.

Read this case alongside Klarna, and the two sides of AI support come into focus: Klarna proves that AI can efficiently handle massive volumes of repetitive inquiries (true), and DPD proves that AI without judgment or guardrails causes trouble (also true). **Whether AI support succeeds or fails has never depended on how smart the model is, but on how thorough the system design is.**

5.2 Platform Landscape: A Market Map of AI Support

First, consider the size of the market. The global AI customer-support market was worth roughly \$12 billion in 2024 and is expected to approach \$50 billion by 2030^[5-42] [third-party verified]. China's enterprise intelligent customer-service market reached RMB 7.19 billion in 2025, up 55.3% year over year, with large models and agents accelerating into production environments^[5-18] [third-party verified].

The main players fall into several categories:

Silicon Valley upstarts: - **Decagon:** launched in stealth in 2023 and reached a \$4.5 billion valuation. It describes itself as building an "AI concierge" spanning all channels (chat/email/voice). Its customer Hunter Douglas — a custom window-treatment e-commerce business operating in 11 countries — used it to replace the traditional IVR "press 1, press 2" flows, deploying localized AI personas in each market^[5-17] [official claims]; - **Sierra:** founded by Bret Taylor, former co-CEO of Salesforce and chairman of OpenAI's board, and Clay Bavor, former Google VP, focused on enterprise-grade support agents^[5-34] [third-party verified]; - **Intercom Fin:** a three-tier architecture (an app tier for training and deployment, an AI tier for RAG, and a conversation tier), with a 50-70% autonomous resolution rate, making it Intercom's fastest-growing product (see 5.4)^[5-19] [third-party verified]; - **Ada:** founded in 2016, serving Square, YETI, and others; its per-conversation billing model has been controversial [third-party review].

Built into the giants: Zendesk AI, Salesforce Agentforce (ARR over \$1.5 billion, up 240%+ year over year).

The Chinese contingent: leading vendors such as Qimo (Chinese: 容联七陌), NetEase Qiyu (Chinese: 网易七鱼), and Sobot (Chinese: 智齿), plus vertical players such as OneConnect (Chinese: 金融壹账通) and Cloopen (Chinese: 容联云)^[7-07] [third-party verified].

Every vendor talks about "autonomous resolution rate" — the share of conversations AI resolves without handing off to a human. Fin claims 50-70%, and most platforms claim 60-80%. But watch the definition: the "resolution rate" vendors count and whether customers feel the issue was "actually resolved well" are two different things. A 53-item ROI benchmark study by Digital Applied found a gap of 30-40 percentage points between vendor-reported figures and independent data^[5-04] [third-party verified].

5.3 Handoff Design: Where Systems Are Most Likely to Break

A guide from Cresta (an AI contact-center vendor) points out a counterintuitive fact: **in AI support projects, the handoff to human agents receives far less investment than the AI interactions themselves — yet it is precisely the**

handoff that directly determines the resolution rate, handling time, and customer satisfaction of the most complex, highest-value conversations^[5-33]
^[5-05] [official claims].

Why? Because the simple problems AI can solve never need a handoff; the ones that do are the complex, high-value, and usually emotionally charged conversations that AI cannot handle. When the handoff is designed well, customers feel "the AI saved me time, and a human picked it up"; when it is designed badly, customers feel "bounced around by a robot" — and a single negative chatbot experience can drive away 30% of customers^[5-36] [third-party verified].

There is a set of widely cited industry data here: Forrester, commissioned by Cyara, surveyed 1,554 consumers globally and found that consumers generally rate their chatbot experiences poorly; Forbes's coverage further noted that bad bot experiences directly cause customer churn^[5-36] [third-party verified]. More systematic evidence comes from research published in the California Management Review in April 2026: 53-77% of users have encountered a bad chatbot experience, and the hidden costs include customer churn, brand damage, and the cost of transferring to human agents^[5-46] [third-party research].

Field experiments in academia yield more granular conclusions. Professor Lauren Xiaoyuan Lu of Dartmouth's Tuck School of Business, in a field-experiment paper based on Alibaba's customer-service operations, found that even with "human-in-the-loop" (AI suggests, humans confirm), agentic AI still struggles with angry customers [third-party research reporting]. Qualtrics's data is equally sobering: AI customer service fails at a rate nearly four times that of other AI use cases — "use AI to save money rather than solve problems, and customers can tell"^[5-32] [third-party research].

These studies all point to a single conclusion: **handoff is not a failure of AI support — it is a necessary part of AI support.** Good system design treats "when to hand off to a human" as a first-class citizen — OneConnect's practice is "large models and small models in layered architecture, human-machine collaboration," with a human-replacement rate above 60%^[5-47] [official claims]; Cloopen used real-time agent assistance (Copilot mode) to lift marketing conversion by 30%^[5-41] [official claims]. Note the keyword here: a replacement rate of 60% is not 100% — the remaining 40% is left to humans, and that is precisely what safeguards the customer experience.

5.4 From Skeptic to Believer: The Intercom Story

Intercom is a leading player in global customer communication platforms and the parent company of Fin, its AI support product. It has admitted something unusual for a tech company: **Intercom was originally an AI skeptic.**

Senior Principal Engineer Brian Scanlan told the story of this pivot at AWS re:Invent 2025: they had been doing machine learning for years (summaries, auto-replies), but when the generative AI wave arrived, the team was full of doubt — "can this thing actually handle customer conversations?" Within a week, they went from skeptics to believers. Fin reached an autonomous resolution rate of 50-70%, becoming Intercom's fastest-growing product and hitting \$100 million in ARR within a year^[9-38] [third-party verified].

Another data point is Leadership Circle — a global leadership development company serving 3 million customers, spanning 21 languages, with more than 12,000 certified coaches. Its transformation lead, Miranda Dunn, described an "aha moment": the team was drowning in 700 tickets a day like "I can't log in," with support agents exhausted from running to stand still. After rebuilding support around an AI-first system, the AI absorbed the repetitive issues and humans focused on high-value scenarios — support went from a cost center to a growth engine^[12-13] [official claims].

5.5 Service as Growth: The Klarna Double-Edged Lesson

Back to the Klarna from the preface. Lay out its full story arc and you have a textbook on AI support:

Act I (Feb 2024): In the first month after the AI assistant went live — 2.3 million conversations, two-thirds of the workload, the equivalent of 700 full-time support agents, resolution time cut from 11 minutes to 2 minutes, customer satisfaction on par with human agents, and an estimated \$40 million in annual profit improvement^[5-28] [official claims]. That year the AI handled 31 million conversations, roughly 80% of support chats, yielding about \$39 million in cost savings^[5-25] [third-party verified].

Act II (2025): The CEO publicly admitted that "cost-led evaluation led to a decline in service quality" and moved to reinforce the human-agent entry point^[5-29] [third-party verified].

Act III (2026): In a Bloomberg interview, the CEO admitted "we went too far" and announced hiring again — not to replace AI, but to bring in about 100 highly skilled humans to handle the scenarios AI cannot crack or customers will not accept^[12-06] [third-party verified].

How should this be read? In two years, Klarna proved two propositions: **AI support really can push cost reduction and efficiency to the extreme (the 700-person equivalent is real); but optimizing customer service as a pure cost center harms customer operations themselves (the service degradation is real).**

CMSWire's framework offers a way out: stop calling service "support" — the support team is a growth engine^{[5-45][5-38]} [third-party verified]. The data backs this up: Gartner says 81% of companies compete primarily on customer experience; Bain's data shows it costs only one-fifth to one-seventh as much to keep an existing customer as to acquire a new one; and a 5% increase in retention can raise profits by 25-95%^[5-40] [third-party verified]. The Hunter Douglas case shows what "service as growth" looks like in AI form: its AI support agent spots purchase intent mid-conversation and directly closes sales^[5-17] [official claims] — service is no longer a cost center but a conversion channel.

5.6 Hallucination and Liability: AI Promises Need Human Review

Another minefield for AI support is hallucination — the AI earnestly fabricating policy out of thin air. The Air Canada case is a landmark: in 2022, passenger Jake Moffatt booked a ticket to attend his grandmother's funeral; the airline's website chatbot promised he could buy a full-fare ticket first and apply for the bereavement fare discount afterward (a policy that did not actually allow this). The passenger followed the instructions, was refused, and sued in British Columbia's small claims court. In February 2024, the court ordered the airline to pay roughly CA\$812 in compensation and rejected the defense that "the chatbot was an independent legal entity" — **a company is responsible for what its AI says**^[10-18] [third-party verified].

The far-reaching significance of this precedent: every word an AI support agent says is, legally, the company's word. The DoNotPay case is the other side: the product, which called itself an "AI lawyer," was fined by the FTC for false promises^[10-11] [third-party verified]. The operational guidance for business leaders is straightforward: **AI**

output must be reviewed by humans, and content involving promises, policies, and compensation needs guardrails — DPD got into trouble for having no guardrails; Air Canada for having guardrails that failed to hold the line.

5.7 The Cost Ledger of Human Handoff Is Becoming a Compliance Ledger

Section 5.3 treated handoff design as an experience problem; by 2026 it is a legal one.

The arithmetic is clear, which is exactly why the human entry point is hard to find. A consumer complaints analysis from China's national consumers association disclosed a cost comparison: a human support agent costs about RMB 3,000 a month on average, while an AI support subscription starts at RMB 99 a month, and larger e-commerce companies save more than RMB 1 million a year in direct labor costs after adoption^[5-48] [official claims]. The illustrative case is worth copying into any deployment plan: before paying a deposit, a consumer was told explicitly by the platform's official AI support agent "Xiaogua" that the RMB 28,700 vehicle price already included the transfer and registration fees; he paid the deposit on that basis, and at the offline transaction the salesperson demanded another RMB 1,000 for the transfer fee. The association's position is the key liability standard — an AI support agent a business deploys lacks the legal personality to bear civil liability on its own, but its output is part of the service the business provides. Legal practitioners add that excessive menu depth, deliberately buried human entry points, no clear notice, and an AI that cannot recognize a "human agent" keyword may violate Article 59 of China's E-Commerce Law and infringe consumers' right to know and right to choose under the Consumer Protection Law.

Complaint volume is turning a design flaw into a regulatory item. Complaints mentioning "AI customer service" on the Black Cat complaint platform exceed 4,800, spanning e-commerce, food delivery, and bike sharing; the consumers association received 985,928 complaints in the first half of the year, with after-sales service accounting for 26.79%, and complaints involving AI support or general-purpose AI keep rising, concentrated in false promises, difficulty reaching a human, and inaccurate generated content^[5-49] [official claims + third-party verified]. In another case a smart-lock failure locked a family of three — including a six-month-old baby — outside for about 50 minutes, with only AI support reachable throughout. The rules

have now answered the question: a service coordination standard effective September 1 states that businesses may not refuse to honor a commitment on the grounds that "the AI's answer does not represent the company's position," and must provide a convenient human handoff entry point; complex problems cannot be left to AI alone.

For AI-powered customer operations the lesson is direct: optimizing "route to human" as a cost item comes back multiplied in complaints and litigation, while treating it as part of service design is precisely what makes AI support sustainable.

5.8 Resolution Rate Becomes an Industry Unit, and People Are Re-Tiered by Value

The competitive focus among vendors has shifted from "which model do you plug in" to "who owns the post-trained model and the evaluation loop." Intercom released Fin Apex 1.0, a model post-trained for support scenarios, reporting a resolution rate — the share fully resolved without human intervention — of 73.1%, versus 71.1% for GPT-5.4 and Claude Opus 4.5 and 69.6% for Claude Sonnet 4.6; the model powers the existing Fin AI agent, which handles more than 2 million customer conversations a week^[5-50] [third-party verified]. A gap of about two percentage points looks small, but it is wider than the typical gap between generations of frontier models — and at a scale of tens of millions of customers, every point of resolution rate materially changes unit economics. It also turns resolution rate into a comparable industry unit, pulling vendor pricing and performance commitments toward pay-per-resolution.

The other path re-tiers humans. Klarna's CEO set out the company's long-term support design in a podcast: the AI assistant answers simple questions while human assistance is rebuilt as a "VIP experience," with staffing moving toward an Uber-like model — beyond outsourced contractors, the company recruits support staff directly from its own customer base, "just as someone can drive for Uber for a while, they can come online and work as Klarna support." His value judgment: "If AI can do customer service, that means AI is the cheap tier of customer service... we think the future of the VIP experience is going to be human connection"^[5-51] [third-party verified].

This is not "AI takes half and humans take half" but a tiering by customer value: low-value interactions fully automated, high-value interactions turned into a scarce resource, with a flexible labor market (customer part-timers) providing elasticity. It

contrasts with Klarna's earlier claim that its AI agent was doing the work of 853 employees — the same company, two years apart, two logics. That contrast is itself the maturation curve of this business.

Buyers' selection criteria are shifting from "can it answer" to "who is accountable." After a stringent RFP process, Notion's global head of customer experience adopted Decagon's AI support agent, reporting publicly that ticket resolution speed improved 34%, deflection doubled, and only 3.4% of conversations required escalation to a human; her framing was not "how much can AI save" but that CX had long been treated as a transactional department (handling complaints, high attrition, work routinely undervalued) and that Notion wanted to turn it into a growth function^[5-52][official claims]. Placed alongside Intercom's 73.1% resolution rate, it shows where the buyer's evaluation has moved: interaction quality, interface, integration depth, product roadmap, and security and compliance standards are now procurement line items on par with resolution rate — and none of them is a model capability, they are platform capabilities.

5.9 AI Support's Rocky Start: AI Amplifies the Company's Incentives

The first half of Chapter 5 covered the platform landscape of AI support and human-backup design. A 2026 consumer survey delivered a chilly opening review.

CNBC, citing Qualtrics' 2026 Customer Experience Trends Report, reports that nearly one in five consumers who have used AI customer service consider the experience "useless" — a failure rate about four times that of AI in general use; consumers rate customer-service AI among the worst on convenience, time savings, and usefulness. The report quotes Cognizant's Ben Wiener: **"AI doesn't change company incentives, it amplifies them"** — if management's goals are to compress refunds, reduce human handoffs, and shorten calls, an AI agent will execute exactly that, more consistently and at greater scale. The same piece offers several industry coordinates: Zendesk CEO Tom Eggemeier predicts 50% of digital customer-service interactions will be handled by AI within three years and 80% within five, and criticizes too many companies for defining "resolved" to include transfers and non-answers; Decagon says it signed 100-plus enterprise deals in 2025, with its valuation rising to

\$4.5 billion; Sierra uses outcome-based pricing; Klarna's AI assistant initially did the work of 700 agents, since raised to 800, with satisfaction on par with humans (vendor claim)^[5-53][third-party verified].

This material and Section 5.5's Klarna are two sides of the same lesson: **whether AI customer service works well often depends not on model capability but on what goal the company connects it to.** If the goal is "close the conversation faster," AI will make the bad experience faster and more hidden.

5.10 Summary

Customer service is the stage of AI-powered customer lifecycle management with the most material and the densest lessons. Core takeaways:

1. **AI support's capacity is real** — Klarna's 700-person equivalent, Fin's 50-70% autonomous resolution rate, and China's intelligent customer-service market growing 55% a year;
2. **AI support's boundaries are equally real** — DPD insulting its own company, Air Canada paying damages, 53-77% of users having bad experiences, and a failure rate four times that of other AI use cases;
3. **Success hinges on system design, not the model** — handoff design is a first-class citizen, and "human-in-the-loop" still struggles with angry customers;
4. **Service is a growth engine, not a cost center** — provided cost-cutting is not the sole objective (the Klarna lesson);
5. **AI's promises are legally the company's** — hallucination guardrails are the compliance baseline.

Next chapter: customer success and retention — how to keep customers and get them to buy more.

Chapter 6 Retention and Expansion: AI-Driven Customer Success

6.1 The \$136.8 Billion Churn Black Hole

Look at the numbers first: U.S. companies lose \$136.8 billion every year to avoidable customer churn^[6-24] [third-party verified]. Roughly 70% of a B2B buyer's purchase journey is completed anonymously before they ever contact a vendor (see Chapter 3), which means customers often "leave" before the company even notices.

How did traditional customer success (CS) try to prevent churn? Through customer success managers (CSMs) watching accounts by hand: whose usage is declining, who is up for renewal, who has not logged in recently. One CSM can barely maintain a few dozen accounts at most, while a company has thousands of customers. As a result, most churn risk is discovered only after the customer has already decided to leave.

The core proposition of AI-powered customer success: **turn churn detection from "discovery after the fact" into "prediction in advance," and turn customer success from "reactive firefighting" into "proactive stewardship."**

6.2 Churn Prediction: Behavioral Data Is More Honest Than Surveys

The principle behind AI churn prediction is no mystery: machine learning models aggregate multi-touchpoint data — product usage frequency, feature depth, ticket activity, email engagement, conversation sentiment — to identify churn patterns that are too subtle for humans to see^[6-24] [official claims].

Why is AI better than people? Because churn signals are weak and scattered. Customers will not say "I'm leaving," but they will: log in less often, from daily to weekly; stop using some core feature; shift their tickets from questions to complaints; start paying bills late. No single signal means much on its own, but combined they trace a clear path to churn. G2's expert survey points out that the strongest churn predictors are not isolated metrics but cross-dimensional patterns — a combination of plummeting product usage, onboarding friction, decaying feature adoption, sentiment shifts, and billing behavior^[6-08] [third-party verified].

At the methodology level, academia has already run through a complete cycle of evolution. Systematic reviews from 2020–2024 show that the mainstream directions in churn prediction modeling are profit-driven modeling (replacing plain accuracy with expected-profit maximization), explainable AI (SHAP, attention mechanisms), and adaptive learning^[6-21] [third-party research]. Real-world deployments are equally solid: TecnoSpeed, a Brazilian SaaS company, ran machine learning churn prediction on real customer data and published the work in a peer-reviewed journal^[6-25] [third-party research].

Aside from the technology, there is one key reminder: **a model is not better for being more complex.** Practice guides note that in the progression from rule-based health scores to ML prediction, logistic regression (easy to explain) and random forests (ensembles of many trees) are often more practical than deep neural networks — because customer success teams need to know "why this customer is at risk" before they can act^[6-11] [third-party verified].

6.3 Health Scores: From Manual Scoring to Real-Time Signals

The health score is the core instrument of customer success: it merges product usage, support history, survey feedback, engagement behavior, and commercial signals into one composite metric that predicts how likely a customer is to renew, expand, or churn^[6-06] [official claims].

Traditional health scores suffer from four big problems — choosing metrics, weighting them, setting thresholds, and adapting them across customer segments — all of which rely on human guesswork^[6-11] [third-party verified]. AI upgrades them in three ways:

1. **More data sources:** beyond usage rates, incorporate sentiment analysis (the tone of emails, calls, and tickets), relationship dynamics, and billing behavior^[6-05] [official claims];
2. **Real-time updates:** scores change as behavior changes, instead of waiting for a CSM's manual update cycle^[6-05] [third-party verified];
3. **Predictive rather than descriptive:** shifting from "what is the customer's state right now" to "what will the customer do next."

The performance data: companies that adopt predictive AI health models improve retention by up to 2x and detect churn risk 25–40% faster^[6-23] [third-party verified]. ChurnZero's annual survey offers a more sobering view: 73% of respondents admit that their current health scores cannot reliably predict churn^[6-19] [third-party verified] — AI health scores are an opportunity, but most companies' health score systems are not there yet.

One methodological trap deserves to be singled out: **telemetry data is not the same as live intent**. Health scores look entirely at behavioral data (logins, clicks, usage), but a customer "using a lot" and "planning to renew" are two different things. ChurnZero's research shows that teams using dedicated CS platforms have significantly stronger churn prediction, because they incorporate relationship and sentiment data^[6-19] [third-party verified]. AI-driven voice of customer (VoC) programs are an attempt to fill this gap: AI-moderated conversational surveys that collect qualitative data at the scale of hundreds of respondents — where a traditional in-depth study can only cover 20–30 people, AI interviews expand that by an order of magnitude, and respondents say more than they would filling out a questionnaire^[6-18] [official claims].

6.4 Expansion: The Revenue CSMs Miss

Retention is defense; expansion is offense. Net revenue retention (NRR) is the gold-standard metric for whether customers are buying more over time: an NRR above 100% means revenue from existing customers is growing, and the company can grow without acquiring new customers.

NRR benchmarks by tier (Statista 2025–2026): best-practice level above 130% (usage-based platforms such as Snowflake, Datadog, and HubSpot), Top Quartile at 115–130%, and median at 100–115%^[6-20] [vendor benchmark data]. NRR ties directly into valuation: companies with NRR of 100–110% are valued at roughly 6x revenue, and those above 120% can reach 8x or more^[6-15] [official claims].

How does AI help customer success teams find expansion opportunities? Gainsight introduced the concept of the CSQL (Customer Success Qualified Lead): AI scans the entire customer portfolio against cross-account, cross-year historical data and automatically generates expansion leads^[6-09] [official claims]. Planhat breaks out six

types of expansion signals: usage thresholds, recent positive sentiment, freshly completed milestones, depth of feature adoption, organizational growth signals, and approaching contract expiry^[6-03] [official claims].

The key insight: **expansion readiness is not a single signal but multiple signals aligning at the right moment** — high usage alone does not mean a customer is ready to expand; high usage plus recent positive sentiment plus a freshly completed milestone does^[6-03] [official claims]. Thrive's case: after using AI to identify expansion opportunities, license upsell rose 24%^[6-03] [official claims].

ChurnZero's sixth annual industry survey (covering nearly 800 customer success and post-sales leaders) reveals a correlation at the team level: companies using dedicated CS platforms have significantly higher NRR (98% for CSM teams vs. 90% without a platform)^[6-01] [official research]. Tools are not a silver bullet, but teams that use them do perform better.

6.5 Win-Back: The Underrated Growth Lever

Churned customers have not really gone away — many are just "asleep." Win-back is the most underrated function in customer success.

The quantified advantages of AI win-back are striking: AI-driven reactivation campaigns win back around 22% of customers, versus 5–10% for generic reactivation emails; email open rates are 38% versus 15–20%^[6-07] [third-party verified]. A case in point is a DTC skincare brand: annual revenue of \$4.2 million, with growth stalled after advertising costs rose from \$32 to \$48 within 18 months; an audit found that 8,200 customers (62% of the customer file) had not purchased in over 90 days. After reactivation through segmented automated sequences, 15.3% of dormant customers were awakened, recovering \$187,000^[6-26] [official claims].

But win-back has an economics precondition, and Gruv's framework spells it out clearly: **win-back is not inherently cheaper than new customer acquisition — it is only worth investing in when "recovered net contribution" outperforms net new acquisition on the same basis.** Opens, clicks, and restarts are just signals; "paying customers coming back on a contribution basis" is the real proof^[6-29] [vendor analytical framework]. Recurly's data is also a caution: customers who churned and are re-activated show lower retention afterwards (re-acquisition rates fall from 4.1% to around 2.8%)^[6-29] [vendor analytical framework] — won-back customers are more likely to churn again and need continuous stewardship.

6.6 From Health Scores to Triggers

The score's position has changed. A methodology article from ChurnZero argues that customer health scoring has evolved from a static reporting tool into a real-time system of action: AI expands what is measurable — not just what customers did in the product, but how they feel about the relationship and which relationships are strengthening or deteriorating — and it changes what teams can do in response, turning the health score from a report into a trigger that can generate recommended actions and increasingly execute them autonomously. That forces customer success leaders to fundamentally reset three things: how the score is designed, what data feeds it, and what it is allowed to do (permissions and automation boundaries)^[6-06] [official claims]. A risk is embedded here: once a scoring system begins automatically triggering customer-facing actions, metric bias is amplified directly into customer intrusion, so model confidence, human review thresholds, and action reversibility must be defined alongside the scoring design — which also explains earlier research finding that roughly 73% of health scores fail at churn prediction (see Section 6.3).

Priorities are converging too. Drawing on its customer success data, Gainsight notes that CS teams operate in a far harsher environment than a few years ago: customer bases keep growing, products are more complex, expectations for retention and expansion are higher, and budgets have not grown proportionally. Its customers adopt AI at a rate 13% higher than non-customers, with significantly more usage in high-impact scenarios such as churn prediction and sentiment analysis; the author concludes that leading CS teams do not rely on individual effort or incremental improvement but build "systems" that scale AI and agentic capabilities consistently^[6-30] [third-party verified].

What does such a system look like in production? A published AI CS agent case study offers a full set of numbers: the customer, a B2B SaaS company, cut churn by 40% within six months, added \$1.2 million in net new expansion ARR, raised accounts per CSM by 2.4 times, reduced quarterly business review preparation from four hours to 35 minutes per session, and pushed NRR above 115% within two quarters. The mechanism turns expansion from "occasional luck" into systematic output: the CSM arrives already carrying agent-assembled data, talking points, and business context^[6-31] [official claims]. The boundary must be stated: this is a vendor-reported customer case without third-party audit or methodology disclosure, so the figures should be read as vendor claims rather than verifiable benchmarks.

Don't forget the water level. These tools operate in a harder environment: Bain observed a broad software index fall about 15% in a few weeks and roughly 25% from its 12-month high, triggered directly by a frontier model company releasing an agent product for general office work. The installed customer base has held so far — average gross retention remains 90% or better — but net revenue retention has stalled, penetration curves for core functionality have flattened, and seat growth has lost force as a growth driver; meanwhile buyer attention and budgets are shifting from incremental software purchases to AI tools and transformation programs^[6-32] [third-party verified]. In other words, customer success teams must defend not only customers but the very question of what customers are willing to pay for.

6.7 After Churn Attribution Is Handed to an Agent

Churn prediction answers "who will leave" but not "why they left" — and the latter has long depended on CSMs recalling events after the fact and backfilling documentation, the step most easily skipped in customer success. In 2026, that step began being taken over by agents.

Forecasting beforehand extends into attribution afterward. ChurnZero released a Retrospective agent: after an account is flagged as churned, it reads the account's full history, classifies the reason against the company's existing churn-reason set, attaches a confidence score to the analysis, and drafts a churn analysis report for the CSM to review; the design goal is to cut the time CSMs spend on post-hoc documentation and let CS leaders identify churn patterns faster^[6-33] [official claims]. That combination of "confidence score plus human review" is worth noting: classification goes to the model, judgment stays with people. It is a relatively restrained form of deployment.

The tool is the precondition; authority is the variable. After rebuilding its CS function on a customer success platform, medical AI company Lunit International reports NRR of 110%, annual churn of 1.5%, and 80% of revenue coming from existing customers; more informative is its VP of Customer Success's description of how the team changed — from "a group of people new to customer success" into "a group of business owners with an entrepreneurial mindset," because CSMs began leading renewals, joining contract negotiations, and driving expansion^[6-34] [official claims]. The contrast makes the point: differences in NRR often come not from analytical capability but from CSMs' authority and the metrics they are measured on (accountable for renewal and expansion, not for satisfaction).

Retention prediction is being turned into a tradable media product. On the retail side, Sam's Club launched a set of targeting products in its advertising business: brands can build their own predictive audiences and segments for churn prevention and declining brand loyalty, and can target members entering the replacement window for long-cycle goods such as televisions and refrigerators^[6-35] [third-party verified]. The traditional approach retargets on recent purchases; these products front-load predictions of who will leave and who will buy. When retention prediction can be bought and sold inside an ad system, it has moved from an internal tool to an industry-wide capability.

But no amount of prediction replaces one judgment: are they still using it? SaaStr's Jason Lemkin uses his own decision to stop using Adobe Marketo as an example (he was one of Marketo's first ten customers and recommended it to clients for years): the product was no longer actually used by his team — unsubscribe handling was still broken, daily API rate limits were the biggest problem, and historical data was retained for only six months. His conclusion is that a discount can keep the contract but not the usage; it only postpones churn and pollutes the NRR metric. The criterion should be "is the customer still genuinely using it," not "did they renew"^[6-36] [third-party verified].

The same logic has a more extreme version in consumer hardware: the life-or-death line for children's AI hardware is not sales but 180-day retention. Yuan Lin (Chinese: 袁琳), founder of Tinglixiong (Chinese: 听力熊), argues there is a fault line between the buyer and the user — parents pay for functional value (learning, safety), while children decide whether to keep using it based on emotional value (whether it understands them, whether the interaction is fun); if the sale completes but the delivery does not, the product can be shelved within ten days. The real moat is also not hardware but content and interaction (age-appropriate language, guided answers, a personality that keeps growing, safety boundaries)^[6-37] [third-party verified]. **The retention curve is decided by the user, and the user is often not the payer** — a rule that holds in B2B too, where it simply becomes "the people using it are not the people signing the contract."

6.8 Summary

The AI transformation of retention and expansion is, at its core, customer success upgrading from "watching customers by hand" to "AI early warning + human judgment":

1. **Churn prediction:** behavioral data is more honest than surveys — against the \$136.8 billion churn black hole, AI can detect signals months before a customer leaves;
2. **Health scores:** from manual scoring to real-time signals, yet 73% of companies' health scores are still unreliable — telemetry data is not live intent, and sentiment and relationship dimensions need to be added;
3. **Expansion:** CSQL turns "customer success" into a revenue engine — the valuation gap between a company at 130% NRR and one at 100% NRR is an order of magnitude;
4. **Win-back:** AI win-back rates are 22% versus 5–10% for mass sends, but won-back customers are more prone to churning again — the economics need to be worked out.

With this chapter, all five stages of the customer lifecycle (acquisition–conversion–service–retention–expansion) have been covered. But every one of these engines depends on a single foundation: customer data. The next chapter covers the AI CRM and the data foundation — where the fuel for these engines comes from.

Chapter 7 The Foundation: AI CRM and Customer Data

The acquisition, conversion, service, and retention covered in the previous chapters are all engines. This chapter is about the engines' fuel — customer data — and the system that carries it: the CRM.

7.1 A Counterintuitive Fact: The Bottleneck of AI CRM Is Not AI

A counterintuitive conclusion has been circulating in the CRM statistics community in 2026: **the main cause of AI CRM project failure is data quality, not the AI technology itself**^[7-03][third-party verified]. As early as 2025, Gartner predicted that by 2026, more than 40% of AI CRM projects would fail or be delayed because of data problems^[7-03][third-party verified].

It makes sense once you think about it. No matter how smart the AI is, if the customer data fed to it is dirty, fragmented, and stale, the predictions and action recommendations it produces are garbage. Many companies buy the most expensive AI CRM only to discover that their customer data is scattered across a dozen systems: sales uses one, service uses another, finance uses a third, and the same customer has a different name, status, and history in each.

That is why this chapter runs in reverse order: data (the fuel) first, then the CRM (the engine compartment).

7.2 The Data Foundation: Identity Resolution Is the Bedrock

What a large share of failed CX/AI projects have in common is not tools or data volume but **the lack of a reliable customer identity layer**^[7-04][third-party verified]. For many companies, the "Customer 360" is only a loosely stitched-together reporting view, not an operable identity layer — the same customer has one ID on the website, another in the app, and a third in the in-store membership system, and the AI cannot recognize them as the same person at all.

AI identity resolution solves this: it combines deterministic matching (strong identifiers such as email, phone number, and membership ID) with probabilistic matching (device fingerprints, behavioral patterns) to unify scattered customer identities^[7-04][third-party verified]. It sounds like a technical detail, but it decides whether every application built on top can work at all — get identity resolution wrong, and personalization recommends to the wrong person, churn alerts watch the wrong customer, and expansion signals are all noise.

The CDP (customer data platform) is the data hub that sits above the identity layer. By 2026, the CDP has evolved from an optional tool into the infrastructure of modern customer experience: unifying first-party data is no longer a competitive advantage but a prerequisite for real-time personalization, privacy compliance, and cross-channel orchestration^[7-24][third-party verified]. Tealium's annual survey (1,200 professionals worldwide) puts concrete numbers behind this: CDP users' goal attainment rate is 92%, and 84% say the CDP has simplified their AI projects^[7-01] [official research].

7.3 Intent Data: Seeing the "Invisible Shopping"

As Chapter 3 covered, roughly 70% of a B2B buyer's purchase journey is completed anonymously before the buyer ever contacts a vendor^[3-34][official claims]. Intent data is how companies see this "invisible shopping": platforms aggregate behavioral data across thousands of websites and review sites to identify "who is intensively researching a certain category."

The market is already sizable: the B2B intent data tools market reached \$4.49 billion in 2026 and is projected to reach \$20.89 billion by 2035 (a 16.6% compound annual growth rate); 91% of B2B marketers use intent data to prioritize sales^[7-28] [third-party verified]. 71% of B2B marketers use third-party intent data in ABM (up from just 55% in 2022)^[7-09][third-party verified].

Performance benchmarks exist as well: accounts reached via intent data convert at 21.3%, versus 8.4% for the control group^[7-09][third-party verified]. Blue Yonder is the case in point: the supply chain software company, with billion-dollar-level annual revenue, upgraded 6sense from "viewing insights on single accounts" into a company-wide ABM engine, lifting accounts in the purchasing stage by 40% and cutting CPL (cost per qualified lead) to \$145^[7-10][third-party verified].

But intent data has its pitfalls. Benchmark data shows a 62% validation failure rate for intent signals^[7-09][third-party verified] — intent data is a "might be interested" signal, not a "definitely buying" fact. Platforms like Bombora are not cheap either, at \$30K per year^[7-29][third-party verified]. The right way to use intent data: treat it as the first-layer sieve for prioritization, combined with human sales judgment and content nurturing, rather than treating it directly as ready-to-close leads.

7.4 AI CRM: From a System of Record to a System of Action

The CRM is the operating system of customer operations. In the AI era, the CRM is undergoing a fundamental shift: **from a "System of Record" to a "System of Action."**

Traditional CRM is a system of record: salespeople log customer information, follow-up notes, and deal status, and the system's main value is "don't lose information." An AI-native CRM is a system of action: the system automatically captures every email, every call, and every meeting, keeps records fresh on its own, and prepares "next actions awaiting approval" for humans — AI does the work, people make the decisions^[1-12][official claims].

Deng Yongfu (Chinese: 邓永富), president of Chinese CRM vendor Neocrm (Chinese: 销售易), draws the clearest dividing line: "AI CRM 1.0 is a CRM with AI added on; 2.0 is a CRM born for AI"^[4-18][third-party verified]. This distinction shares its lineage with the AI-native vs. AI-assisted divide in Chapter 2: pasting on a band-aid versus growing it into the genes.

Where the major players stand:

Salesforce Agentforce: ARR of more than \$1.5 billion, up over 240% year over year^[7-07][third-party verified]. Salesforce's transformation is almost textbook-grade — it repositioned "Sales Cloud" from a tool into an AI agent platform, letting agents execute sales workflows directly.

HubSpot Breeze: in April 2026 it cut Breeze Customer Agent pricing from \$1 per session to \$0.50 per resolved session — charging for "resolved outcomes," not "number of conversations"^[7-15][third-party verified]. This is the "outcome-based pricing" described in Chapter 4 landing in the CRM arena.

Fenxiangxiaoke ShareAI (Chinese: 纷享销客 ShareAI): claims to have provided AI-native/agentive CRM to 6,000+ mid-sized and large enterprises^[7-31][official claims]. Agentive CRM in the Chinese market is forging its own path — an AI PaaS foundation plus industry intelligence.

Microsoft: eight AI products with "Sales" in their names, one of which was renamed six times in three years (Viva Sales → Sales Copilot → Copilot for Sales → ... → Sales in Microsoft)^[7-27][third-party verified]. Even the giants are iterating fast.

The consensus across the whole CRM market: **the biggest bottleneck for AI CRM is data, not technology** — which is exactly why this chapter covers data before the CRM^[7-03][third-party verified].

7.5 Privacy and Compliance: The Ceiling of Customer Data

More data, stronger AI — but there are red lines on how data may be used. Privacy compliance is the ceiling of AI-powered customer lifecycle management, and this section needs to say it plainly.

The regulatory landscape: GDPR governs Europe, CCPA/CPRA governs California, and PIPL governs China. The core principle of marketing data compliance: in the vast majority of cases, the safest path is to provide notice and obtain consent^[7-02][third-party verified]. King & Wood Mallesons sets out the essentials of omni-channel retail marketing under the PIPL: behavioral data collected automatically may identify individuals once it is linked with other information, and therefore constitutes personal information; data provided voluntarily by users requires full disclosure, and general consent must be distinguished from separate consent^[7-33][third-party verified].

New solutions to the technology-versus-privacy balance are emerging. Federated learning lets vast numbers of edge devices participate in model training with their local data — only model parameters leave the devices, while raw data never leaves them^[7-25][third-party verified]. At the same time, AI is also helping companies stay compliant: 60% of large enterprises used AI to automate GDPR compliance processes in 2025^[7-05][third-party verified].

Consumer attitudes are contradictory: 44% of consumers are disappointed when a brand fails to deliver a personalized experience, while 70% worry about how their data is being used^[7-05][third-party research]. Those running AI-powered customer lifecycle

management must manage this pair of contradictions at the same time — the expectation of personalization and the worry about privacy are two sides of the same customers.

7.6 Segmentation and Reach: A Postmortem on 4.04%

Earlier sections of this chapter described how to build the data foundation. This section takes the opposite view: how concrete the losses are when the foundation is missing.

One company's private-domain assets comprised more than 80,000 enterprise-WeChat contacts, over 100,000 mini-program users, and more than 150,000 mobile-number customers. In a "new member exclusive benefits" campaign all 693 new members were issued the benefit, yet only 28 actually received the notification — a reach rate of 4.04%. The offer itself was flawless: members who received it converted from visit to purchase at 99.86%, and more than 30% went from visit to payment. The bottleneck was not the offer design but **who the message reached**^[7-34] [third-party verified].

The root cause is entirely typical: the platform's built-in segmentation was too coarse — just three behavioral buckets (visited without paying, added to cart without paying, ordered without paying) plus four fixed RFM labels. A single "visited without paying" pool could hold more than 100,000 people, lumping together someone who browsed 20 product pages with someone who tapped in on a whim; and the labels were static within the window, with no predictive power at all. The rebuild combined three-dimensional dynamic segmentation with block-based operations, intelligent channel matching, and data-driven iteration; the store's repurchase rate rose 5% year over year, and the estimated reach rate went from 4.04% to above 10%. The author's three retrospective lessons are more valuable than the methodology: keep the MVP more focused; put channel coverage first (part of the 4.04% root cause was that users had not added the enterprise-WeChat contact, had not authorized the mini-program, and SMS was blocked — the reachable channel set itself was incomplete, so "nudge users to add the contact and subscribe" must be part of the project from day one); and validate scoring weights early with small-batch A/B tests rather than relying entirely on historical behavior analysis.

This gap has a new AI-era version. Research on China's GEO market notes that marketing teams generally want to use existing CDP data to build cognition in AI scenarios, but CDP data is largely accumulated by ad platforms and covers only existing customers, missing the new cognition formed in AI scenarios^[3-65]. Both cases point to the same conclusion: a gap in the data foundation does not show up as "the system is broken" — it shows up as "it looked like we sent it, but it never arrived," and, in new entry-point formats, as "nobody knows who you are."

7.7 The Middle Layer of Attribution: When Users Arrive Through Someone Else's AI

Section 7.1 argued that the bottleneck in AI CRM is not AI but data. This section adds the newest and hardest-to-solve data gap: **an untraceable intermediary has appeared in the acquisition chain.**

Why attribution suddenly fails. Take Meituan's AI agent Xiaomei (Chinese: 小美) connecting into Tencent Yuanbao: the entry point through which users obtain services is migrating from "opening the app" to "saying an instruction inside someone else's AI," but for developers and growth teams it is impossible to tell whether a user converted after guidance from a Yuanbao conversation or through Meituan's own search and recommendations, or even what the user actually saw in that AI dialogue — traditional attribution based on source, click, and install package fails outright in this kind of chain, leaving ROI on cross-ecosystem agents impossible to calculate^[7-35] [third-party verified]. This is not a technical problem but an accounting problem: AI entry points bring increment, and an unaccountable ledger. That is why Chinese brands working with AI entry points must simultaneously demand verifiable passback (dedicated parameters, passphrases, agent IDs, mini-program paths), or channel budgets will have nowhere to land.

The starting point of the CDP is moving earlier. For more than a decade the CDP solved the problem of consolidating customer data scattered across marketing, e-commerce, analytics, and service systems into a unified customer view. A more radical claim is that the mission remains important but its premise has changed: AI demands real-time decisions, privacy regulation has reshaped how data is collected, and customers move seamlessly across devices and anonymous sessions — so the first step of a next-generation CDP is not consolidating "existing customer data" but collecting the behavioral and contextual signals that precede the data's existence^[7-36]

[third-party verified]. Put more bluntly: **in the AI era, "designing the data entry point" sits closer to the source of the problem than "building the data warehouse."**

Customer signals are becoming a data layer agents can call. 6sense's MCP server moved from open beta to general availability, connecting its buying signals, intent data, and account intelligence directly into everyday AI tools such as Claude, ChatGPT, and Copilot Studio so users can query account context with natural-language prompts; with the September product release, the vendor began formally supporting Sales Intelligence users, bringing account context and contact acquisition into AI workflows^[7-37][official claims]. The same direction is appearing on the CRM side: at Dreamforce, Salesforce launched Alforce, positioned as a "live interface layer" that brings internal data, workflows, business logic, semantics, permissions, security, and governance into whatever interface people and agents actually work in, replacing the fixed UI^[7-38][official claims].

Read the three together and the direction of the foundation rebuild is clear: **from "a page employees must open" to "a capability and data layer agents can call."** That drags two very practical consequences along with it — the seat-based pricing model of CRM, and the boundaries of the data asset (once data is to be read by external agents, "who authorizes it and who audits it" becomes part of the product).

7.8 The Directory and Shelf of Agents: The MCP Registry and AI App Marketplaces

Acquisition entry points are changing, and so are the entry points for distributing tools. As customer operations increasingly start from AI, **"can an agent retrieve your capability" becomes a new directory position** — yet another outward expansion of the foundation layer.

The public directory came first. The official Model Context Protocol registry (registry.modelcontextprotocol.io) is positioned as an open catalog and API of publicly available MCP servers, meant to improve their discoverability and implementation and to become a "source of truth" on which sub-registries can be built; it continuously lists MCP server versions submitted by developers and vendors worldwide, ordered by

most recent update^[7-39][official claims]. The acquisition implication is direct: being listed, retrieved, and recommended is becoming the new directory-style distribution slot of the agent era.

Cloud vendors turned it into an app marketplace. In September 2026 Alibaba Cloud released version 2.0 of its "Fanhua Program" for AI product partners, upgrading its cloud marketplace into an "AI application marketplace" — from a traditional software-and-services trading platform into a one-stop platform gathering large models, AI applications, and industry solutions across education, finance, healthcare, retail, manufacturing, research, and gaming, and across scenarios such as customer service, marketing, and digital employees; version 2.0 focuses on five directions: Skills (domain skill products), data (industry and professional datasets), connectors (connecting CRM, ERP, and other SaaS applications within authorized scope), Agents (commercial agent products), and industry co-creation^[7-40][third-party verified]. This in effect commoditizes the foundation layer — customer-operations connectors, industry datasets, and agents are packaged, listed, and procured on demand.

Underlying capabilities are being packaged as standard components. At Agents Week 2026, Cloudflare released a batch of agent infrastructure: Project Think (a preview of a next-generation Agents SDK that lets agents think, act, and persist), an experimental voice pipeline for the Agents SDK (continuous STT/TTS in about 30 lines of server-side code), Cloudflare Email Service (in beta, letting agents send and receive email natively), Agent Memory (a managed persistent-memory service), AI Search (a search primitive for agents supporting hybrid retrieval and relevance weighting), and Browser Run (with live view, human intervention, CDP access, and session recording, at four times the concurrency limit)^[7-41][official claims]. The trend is cloud vendors packaging memory, search, voice, email, and browsers into standard components any agent can call.

In one sentence: **when "who gets called by an agent" becomes the distribution logic, directories, marketplaces, and standard components become the next foundation layer of customer operations** — they don't generate leads directly, but they decide whether leads can find you in the AI world.

7.9 Summary

Three takeaways on the foundation:

1. **Data before AI** — the main cause of AI CRM project failure is data quality; identity resolution is the bedrock, and the CDP is the data hub;
2. **Intent data sees the invisible shopping** — 70% of purchases happen anonymously, yet intent signals carry a 62% validation failure rate; it is a sieve, not a conclusion;
3. **The CRM shifts from record to action** — Agentforce's \$1.5 billion ARR, HubSpot charging by outcome, Neocrm "born for AI": AI-native is the dividing line;
4. **Compliance is the ceiling** — consent is the floor, federated learning is a new solution, and the expectation of personalization coexists with the worry about privacy.

With the foundation done, it is time to talk about people. Next chapter: organization — who drives these engines.

Chapter 8 Organization: People, Skills, and Change

The previous chapters were about which tools to use. This chapter is about who uses them and how things change. Organizational change is the slowest, the most expensive, and the most easily underestimated part of AI-powered customer lifecycle management.

8.1 Three CEOs, Three Extreme Answers

Between 2023 and 2026, three CEOs answered the same question in three completely different ways: "In the AI era, what should my organization do?"

IgniteTech's Eric Vaughan chose to tear it all down and start over. This enterprise software group with nine-figure annual revenue concluded that generative AI was an "existential"-level shift, and looking around he saw a team that was not ready. So he acted: he replaced nearly 80% of his employees within a year. The process was full of resistance and even sabotage, but two years later he said: "I would do it again"^[12-03][third-party verified].

Salesforce's Marc Benioff chose to use AI to cut headcount directly. He acknowledged in interviews that after AI agents went live, he cut the customer support team from 9,000 people to 5,000 and cut support costs by 17%. He said it bluntly: "I need fewer heads"^[12-04][third-party verified].

Intercom's Eoghan McCabe chose to bet the company on an AI transformation. He returned after a four-year absence from the CEO role due to illness, saw growth had stalled, and in "founder mode" bet the company on AI: cut 40% of employees, built a Fin prototype in six weeks, and growth climbed back above 300%^[12-09][third-party verified].

The three stories have something in common: all of them acknowledged that AI is structural, all of them touched headcount, and all of them went through pain. The difference lies in the approach: IgniteTech's 80% workforce renewal was radical, Salesforce's 4,000 cuts were direct cost reduction, and Intercom's layoffs were part of a transformation — its growth engine switched from "human-delivered service" to "AI products."

None of these answers is the standard answer. But stacked together, the three point to a clear reality: **organizational change in AI-powered customer lifecycle management cannot get around the question of "what happens to the people."** This chapter breaks that question into three parts: how GTM is rebuilt (8.2), what new skills people need (8.3), and why change fails (8.4-8.5).

8.2 Rebuilding GTM: AI Agents Enter the Revenue Organization

GTM (go-to-market) is the umbrella term for the revenue organization — sales, marketing, and customer success all sit inside it. McKinsey's 2026 report spelled out the impact of AI on GTM most clearly: based on its B2B Pulse Survey covering nearly 4,000 buyers and sellers in 13 countries, most B2B companies are "doing something," but only a few have captured material value from AI — because of fragmented data, weak insights, manual processes, siloed teams, and poor change management^[8-20] [third-party verified].

The gap between high-growth companies and the rest has already opened up in investment: high-growth companies invest roughly 3x as much in AI as others (71% vs. 25% of respondents reporting high AI investment)^[8-20] [third-party verified].

One new role taking shape is **RevOps managing a fleet of AI agents**. RevSure's research playbook argues that agents read and write to the systems RevOps governs, consume budget from the pipelines RevOps reports on, and escalate to queues RevOps staffs — so managing the agent fleet lands on RevOps "by plan or by accident"^[8-38] [official claims]. The data shows the payoff of this path: after ICONIQ's customers deployed AI agent fleets, their ARR/FTE (revenue per employee) rose from \$370,000 to \$640,000^[8-38] [third-party verified].

Gartner's warning is also worth hearing: sales operations leaders are being asked to drive growth, underpin GTM change, and keep pace with AI transformation, yet they rarely get a chance to pause and reassess their own capabilities^[8-12] [third-party verified]. The first hurdle of organizational change is often that the people responsible for the change are not ready themselves.

8.3 Sales' New Skills: From "Knowing More" to "Reading the Room"

Bain's 2025 report served up an uncomfortable number: **salespeople spend only about 25% of their time actually selling to customers** — the remaining 75% is eaten by meeting preparation, data entry, internal coordination, and other work that "revolves around selling but creates no value." If AI can take over these tasks, it could double selling time; at the same time, AI lifts conversion rates at every stage of the funnel, and compounding across stages could raise win rates by 30%^[8-13] [third-party verified].

Once AI takes over "knowing more" (product knowledge, competitive comparisons, script templates), the sales differentiator must shift. Seismic's judgment: the sales differentiator moves from "knowing more" to "reading the room, earning trust, exercising judgment, and adapting quickly" — the winners will not be the teams that automate the most, but the teams that upskill the fastest^[8-18] [third-party verified].

Mercuri International's research identifies "15 core sales competencies on the rise" and finds that four in five executives believe generative AI will reshape the sales role^[8-03] [third-party research]. Vivun and Gartner's forecasts are more specific: by 2028, AI will close 70% of traditional sales-cycle steps; by 2031, 35% of organizations will bring EQ (emotional intelligence) metrics into sales evaluations^[8-11] [third-party verified]. Sales evaluation shifting from "number of calls" to "quality of relationships" is inevitable in the AI era — because AI can handle the quantity, and quality is where the value of people lies.

There is also a neglected shortfall at the skills level: Mercuri's 2025 research shows that only 9-15% of salespeople are confident in their AI literacy^[8-03] [third-party research]. The tools get rolled out, but people do not know how to use them — this is the most common hidden failure in organizational change.

8.4 AI Sales Training: An Underrated Lever

If people need new skills, training is a hard requirement. AI sales role-play is the fastest-growing training category of 2026: it uses large models plus voice synthesis to simulate realistic customer scenarios and run immersive conversational practice with salespeople — simulating customer objections, needs discovery, and competitive comparisons^[8-16] [third-party verified].

Why is AI role-play needed? Because human coaching resources are pitifully scarce: frontline sales managers spend only 5-8% of their time coaching (at most 3 hours in a 40-hour workweek), and spread across a team of five, a new hire gets under 40 minutes of manager coaching per week during the critical ramp period^[8-08][official claims].

The results data for AI role-play: - **Hyperbound**: helped Vanta cut new-sales ramp (from hire to productive output) time from 210 days to 72 days^[8-08][official claims]; - **Second Nature**: after Oracle NetSuite customers used AI role-play, qualified opportunities rose 32%^[8-33][third-party verified]; - Traditional training vs. AI role-play: AI role-play dramatically shortens the development cycle, can be repeated endlessly, and comes with objective scoring^[8-16][third-party verified].

Category reviews covering more than eight platforms have already appeared (Knowlee and others), proof that this is not a small widget but a market taking shape^[8-05][third-party verified]. The significance of AI sales training goes beyond "learning scripts" — it makes "practice" scalable for the first time: new hires can make enough mistakes in front of the AI before making them in front of real customers.

8.5 The Psychology of Change: Why 74% of Companies Are Stuck in Pilots

The tools are bought and the training is done — so why do most AI-powered customer lifecycle management programs still get stuck in pilots?

BCG's report *Where's the Value in AI?* (surveying 1,000 CxOs across 59 countries and more than 20 industries) offers an attribution: only 26% of companies have the ability to move beyond proof of concept and deliver real value, and 74% have yet to see visible value from AI — **roughly 70% of implementation barriers come from "people," not from technology or algorithms**^[8-21][official claims].

Gartner's data is just as striking: only one in three (32%) mid-to-senior business leaders believe their most recent change initiative achieved "healthy change adoption"; a 2025 survey of more than 2,850 employees found that 79% have low trust in change^[8-24][official claims].

Resistance from people takes several typical forms: **fear of displacement** ("is AI coming for my job?"), **going through the motions** (agreeing in meetings, not using the tools afterward), and **immune responses** (the organization's existing processes

and incentive systems instinctively reject new practices). Intercom's Eoghan McCabe said something on Lenny's Podcast that every change agent should hear: an AI transformation is not just swapping tools — it is swapping the company's "operating system." Cutting 40% of employees was not the goal; making the people who remained believe "we are an AI company" was^[12-09][third-party verified].

Change management has no silver bullet, but several principles have been repeatedly validated: 1. **Start at the top:** the team only believes when the CEO personally uses AI and personally sets the KPIs (Intercom's "founder mode"); 2. **Give anxiety an outlet:** tell the team clearly what AI takes over and what people will move toward (the skills map in Section 8.3 is that outlet); 3. **Let data speak:** run measurable ROI in the AI pilot first, then talk about scale (Greenway's controlled experiment: 800 accounts split into four groups, with a qualified-opportunity cost of \$342 for the hybrid model vs. \$588 for fully autonomous vs. \$1,006 for purely human — the hybrid model's superior ROI is backed by a controlled experiment)^[8-10][third-party verified]; 4. **Accept discomfort:** IgniteTech's 80% workforce replacement is not a template, but it reveals a fact — when change is "existential," gentle incrementalism can amount to a slow death.

8.6 FDE: The Intermediate State of AI Deployment

The previous five sections asked "whether to change." This section asks "who carries the change to the front line" — in the Chinese discourse on AI deployment, that role has a name: FDE, a forward deployed engineer.

The most complete field record is Datawhale's "FDE Case 100": 24 frontline interviews covering manufacturing, government, legal, logistics, retail, foreign trade, and e-commerce, all narrated in the first person by FDEs about "how they found the problem, understood the business, chose the technology, and made trade-offs under real constraints"^[8-41] [third-party verified]. Spread the 24 cases out, and several recurring patterns emerge.

First, the problem is not technology but the ordering of "people — organization — data." An FDE on a state-owned enterprise project put it most bluntly: the difficulty ranking for AI deployment is "people are hardest, then the organization, then the data, and technology last"^[8-41] [third-party verified]. This

matches the attribution in 8.5 — seven-tenths of the obstacles come from people. Several cases repeat one discipline: go on site and first see "how they actually work every day," rather than asking what model to use.

Second, acceptance criteria must be co-defined by the business and land in the workflow. In a case at a German parts manufacturer, the FDE turned the "high-frequency question list" compiled up front directly into the test set: document parsing at about 95% and question-recall fragment accuracy $\geq 80\%$ counted as the technical path tentatively working; the proof of acceptance was not a good-looking demo but that the AI was wired into the company's existing training workflow^[8-41] [third-party verified]. The case also yields a general anti-pattern warning: technical metrics must not be defined by the FDE behind closed doors.

Third, customer-operations scenarios are exactly the FDE's high-value zone. At an e-commerce team running TikTok creator operations, the FDE mapped the entire chain — creator outreach, product matching, sampling, content creation, fulfillment, data review — and had AI first solve "finding the right person" and "which product goes to which creator," cutting average fulfillment time from about a month to about 18 days; but one key step was left not fully automated — a single wrong sentence from a creator with millions of followers costs far more than the labor that automation would save, so human outreach was retained^[8-41] [third-party verified]. At another non-standard manufacturing company, the pre-sales chain of drawing parsing, requirement decomposition, R&D collaboration, cost accounting, and quoting was handed to AI for structuring, with quotes matched back against the company's own rule base and manually calibrated per order, cutting the related workload by at least 50%^[8-41] [third-party verified]. A small construction-equipment leasing company built both the acquisition entry point and quoting at once: large clients were left for the boss to judge personally, while small orders had the system collect requirements and compute a quote range by vehicle type, operator, fuel, distance, and job type, with AI handling the conversation and closing quickly inside the range^[8-41] [third-party verified] — long-tail customers that used to be "un-catchable" are now caught by the system.

Fourth, don't "apply" AI — redo the process. The sharpest argument comes from the U.S. AI-delivery firm Varick's retrospective: having interviewed 300+ CEOs, CIOs, and CFOs at large enterprises, they concluded that most companies apply AI to already-inefficient processes, and "the end result is just making garbage faster"^[8-42] [third-party verified]. They cite a test by the U.K. Department for Business and Trade:

1,000 Microsoft 365 Copilot licenses rolled out over three months produced an average of just 1.14 Copilot actions per user per day; PowerPoint slides came out 7 minutes faster but at half the quality score, Excel analysis was actually slower, and the official conclusion was "we did not find robust evidence to suggest that time savings are leading to improved productivity" — yet 72% of users were satisfied with the tool^[8-42] [third-party verified]. The time was never in the work itself but in queues and handoffs: in the classic case they cite, an insurance application sat in process for 22 days, of which only 17 minutes was actual work. Varick's method sorts every step into one of three buckets — deterministic rules to plain code, judgment calls with abundant historical examples to an agent, and too-risky steps to a human in the loop — turning a 25-step process into 3 agents plus 2 human checkpoints; across a set of deployments, month-end close went from 18–22 days to 7–9, and accounts-payable exception handling fell from 600–800 invoices a month to under 50^[8-42] [third-party verified]. Their judgment is worth every operator remembering: **AI should not be "applied" — you never say "applied cloud" or "applied digital"; you say digital transformation and cloud modernization.**

8.7 An Organizational Prescription for AI-Native Teams

Sections 8.1 to 8.5 answered "how people and processes change." On "what an AI-native organization looks like," the industry is now forming operable structural blueprints.

Shah Rahman, who leads engineering for autonomous ML iteration and optimization for ads at Meta, laid out a three-layer design in a ByteByteGo series: **structurally**, the atomic unit is a small cross-functional pod of 3–5 people operating autonomously with AI agents, with hierarchies flattened and some pods reporting directly to senior leaders by strategic importance; **in roles**, each pillar names 1–2 full-time "Agent Champions" (committing 50%–100% of their time) to reshape workflows, prepare codebases, and restructure operating models; **in governance**, "STO for Everything" (a single task owner for every task) slices ambiguous responsibility cleanly, moving from human-in-the-loop to human-on-the-loop^[8-43] [third-party verified]. The article's warning about the "productivity paradox" matters most: only 20%–30% of an engineer's time is spent coding, so faster code generation does not automatically become organizational productivity — McKinsey found developers using AI assistants were 20%–45% faster on discrete coding tasks, but org-level gains were smaller and harder to measure.

Another practice checklist from the engineering world (Coderio's 10 AI-native engineering practices) offers more testable metrics. It cites Faros AI's telemetry across 22,000 developers: individual throughput rose, but median time in PR review increased by 441%, bugs per developer rose 54%, and incidents per PR rose 242% — **output went up, system quality went down, and teams watching only velocity never saw it**^[8-44] [third-party verified]. So it argues for shifting measurement from "output volume" to "system quality": defect escape rate, review rework rate, rollback rate, test stability, and the share of AI-generated changes accepted without meaningful human modification. The same metrics apply to customer-operations roles — measuring AI support or AI selling should not look only at "how many were handled" but at "how many were reworked and how many escaped."

Organizational redesign must ultimately show up in customer metrics. A white paper from the World Economic Forum and Accenture, "Organizational Transformation in the Age of AI" (surveying 450+ executives worldwide), records that payoff: Ford used an AI decision engine to dynamically adjust audiences and content in multi-wave campaigns based on real-time customer feedback, reaching more than 300,000 customers in three weeks with a 26% lift in conversion; Rabobank, through a unified customer decision hub, delivers more than 1.5 billion personalized interactions a year, with click-through up fourfold, customer lifetime value up 4.7%, and service costs down 2.4%^[8-45] [third-party verified]. The white paper also notes Visa's intelligent commerce system, which lets consumers set autonomous purchase permissions for AI agents to complete authorized transactions, at a time when 47% of consumers already used AI tools for shopping tasks^[8-45] [third-party verified]. Its principles for scaling are worth copying into any change plan: human accountability (clear autonomy boundaries), end-to-end operating-model redesign, a scalable talent system, transparency-driven trust, and a disciplined experiment-and-learn loop^[8-45] [third-party verified].

8.8 AI Will Not Save a Bad Process

Section 8.5 covered the psychology of change and Section 8.7 the organizational prescription. This section adds the hardest evidence of 2026: the gap in AI returns is mostly not a tooling gap but a process and collaboration gap.

Layer AI on a bad process and you get only "mini-productivity." Bain's survey of 1,125 sales and marketing leaders across 18 industries and 40 countries concluded that AI experimentation is now the norm — more than 90% of commercial

organizations are trying it — but value realization is highly uneven: one group of companies chases single, overly narrow use cases, another layers AI on processes that were already broken and harvests only "mini-productivity." The real winners (those in the top quartile of revenue growth in their industry and region and above their margin targets) captured twice the AI-attributed revenue impact of everyone else and gained an additional 30% in productivity by rewriting sales and marketing workflows^[8-46] [third-party verified]. Note that winners do not run more pilots than laggards — which is why the metric to track should change from "number of pilots" to "number of processes rewritten and the attributable output they produced."

Penetration and provable returns are decoupled. Atlassian's research offers two figures that nearly contradict each other: 94% of marketing practitioners use AI at work, but only 3% of CMOs are confident they can see organization-level return on AI investment. The report explains this as a "fragmentation tax" — when AI ignores collaboration, it erodes the productivity gains of point automation; top teams move from basic automation to orchestrated, AI-ready ways of working, rebuilding context, workflows, and culture together^[8-47] [third-party verified]. This is the same coin as Section 8.5's finding that 74% of companies are stuck in pilots: what stalls is not model capability but how people work together.

Resistance has a clear geography. Salesforce's survey of more than 1,500 office workers across four continents found that U.S. employees are 43% more likely to consider themselves "AI skeptics" than the global average; in the U.S., U.K., and France, 53% of respondents identify as skeptics, versus only 26% in Mexico; and countries with higher skeptic shares also show slightly lower rates of using AI as part of a core workflow^[8-48] [official claims]. The same material cross-checks with Stanford's 2026 AI Index: India exceeds 80% on both trusting AI and continuing to use it, while the U.S. hovers around 50% on both. The implication: advanced economies' incentive structures (expectations of job displacement) differ from emerging economies' (expectations of productivity and income gains), so a single global AI rollout plan will hit a wall — the communication focus belongs on how AI relates to individual benefit, not on efficiency alone, which in markets with high displacement expectations reads as a layoff signal.

Payback period is set by operating factors, not by launch speed. Salesforce's survey of 2,025 enterprise agentic AI decision-makers reached a counterintuitive conclusion: the companies that deployed AI agents first were not the ones that saw returns first. On average, enterprises reached "substantial ROI" in about

eight months; professional and business services was among the last sectors to adopt AI agents yet the fastest to achieve meaningful ROI (6.5 months). The three strongest predictors of success are clean, well-governed data available whenever the agent needs it, an agent scope defined narrowly enough, and a human escalation path established in advance^[8-49] [official claims]. None of these are model problems — they are process and governance problems, which is precisely why "shipping first" does not buy "paying back first."

8.9 The Center of Gravity in Sales Is Shifting Backward

Section 8.7 covered how to reorganize teams. This section takes a different angle: when closing itself becomes self-service, the center of gravity of the GTM organization drifts toward the back half of the chain — and the financials of two Chinese model companies provide a rare clean sample.

Revenue doubled while sales spending fell. In the first half of 2026, Zhipu AI (Chinese: 智谱) reported revenue of 954 million yuan, up 399.7% year over year, while sales and marketing expenses fell 14.8% to 178 million yuan; MiniMax reported revenue of 786 million yuan, up 283.1%, while sales and distribution expenses fell 17.9% to 181 million yuan. The structural cause is a change in how transactions happen: Zhipu's open-platform and API revenue surged from 29.1 million yuan a year earlier to 825 million yuan, or 86.5% of total revenue, with customers self-serve topping up tokens on the platform, turning revenue from one-time recognition into recurring accrual^[8-50][third-party verified].

The center of gravity moves from pre-sales to mid- and post-sales. The report's summary is that the sales function shifted from pre-sales acquisition to mid- and post-sales, and the key sales metric shifted from user volume to customer retention — the model became its own best salesperson, and people work more like staff for the model. That has more important implications for organizational design than the numbers themselves: **when a product can close deals self-service, the sales team's KPIs, headcount, and budget structure all get rewritten** — which is precisely the concrete form of the GTM restructuring described in Section 8.2, in a self-service closing scenario. It also sets a boundary for the "AI replaces salespeople" debate: what gets replaced is pre-sales information delivery; what gets reinforced is retention and deep conversion.

The same trend has numbers on the CRM side. Salesforce's Agentic Enterprise Index, built on data from enterprises continuously using Agentforce between February 2025 and April 2026, shows that the number of agents enterprises activated grew nearly threefold by fiscal year end, with average creation time down 53%; agent skill sets can scale up to 350% at peak demand to handle complex tasks; enterprises deploying agents reported fourfold retail online sales growth; weekly employee usage tripled while customer escalation rates held steady at large scale^[8-51] [official claims]. The caveat must be stated: this is a vendor sample drawn from its own active customers, best read as a signal of scale and stability for agent adoption in customer operations, and not extrapolated directly to industry-wide efficiency gains.

8.10 The Last Mile of Budget: The Willingness to Pay Is Running Ahead of the Ability to Deploy

Much of Chapter 8 is about organizations needing to redo their processes. Gartner's 2026 CMO spending survey adds the most concrete piece: **the money is ready; the capability is not.**

Gartner's CMO Spend Survey (January–March 2026, 401 CMOs and marketing leaders, mostly at companies with more than \$1 billion in revenue) finds a core contradiction — "the willingness to pay is running ahead of the ability to deploy." Seventy percent of CMOs name becoming an AI leader as a key goal for 2026, but only 30% believe their organization has the necessary infrastructure. The budget signal looks contradictory — marketing budget as a share of company revenue is roughly flat (7.8% in 2026 versus 7.7% in 2025), and with the total flat, AI investment becomes "a reallocation problem" rather than a growth problem: an average 15.3% of the marketing budget goes to AI, but organizations that are "better at scaling AI" average 21.3%; organizations with an "optimized AI program" spend 8.9% of revenue on marketing on average, above the 7.8% overall. The hidden work is data foundations, governance, process, and talent, and the pacing risk is clear — 56% of CMOs say their organization lacks the budget to execute its 2026 strategy and 54% say they lack the necessary resources; nearly two-thirds of marketers believe AI will change the work, but only 32% think they need to update their skills^[8-52] [third-party verified].

A second reading of the same survey is even more direct: **"AI didn't shrink the marketing organization; it restructured it."** With the budget flat and 15.3% going to AI, the "people share" rose to 24.5% (from 21.9%), while only 30% of CMOs say

they are ready to scale AI; the conclusion is that execution-layer roles are disappearing and senior judgment is worth more, so the real 2026 design question is "where does that judgment sit: in-house, with partners, or both." Forrester expects layoffs to fall on execution and support roles, with senior strategy roles protected — **AI has raised the price of the senior talent capable of commanding it**^[8-53][third-party verified]. This is the same judgment as Section 8.9's "the center of gravity in sales is shifting backward": what AI brings is not a linear gain in per-capita output, but a repricing of the "judgment" role inside the organization.

8.11 Summary

Organizational change is the hardest part of AI-powered customer lifecycle management to replicate:

1. **Three CEOs, three answers** — IgniteTech replaced its workforce, Salesforce cut headcount, Intercom transformed; there is no standard answer, but all of them moved people;
2. **GTM rebuilt** — AI agents enter the revenue organization, RevOps becomes the "agent-fleet manager," and high-growth companies invest 3x more in AI than others;
3. **New skills** — salespeople spend only 25% of their time selling; once AI absorbs the 75% of busywork, human value shifts to "reading the room, earning trust, exercising judgment";
4. **AI sales training** — ramp time cut from 210 days to 72 days; practice is scalable for the first time;
5. **The psychology of change** — 74% of companies are stuck in pilots, and roughly 70% of barriers come from people; anxiety needs an outlet, pilots need data, and change needs resolve that runs from the top down.

The engine, the foundation, and the organization are all covered. The next chapter goes into the industries — what does AI-powered customer lifecycle management look like across eleven industries?

Chapter 9 Industries: AI Customer Lifecycle Management in Eleven Industries

The same AI technology grows into very different shapes in different industries. A bank's AI customer service must clear regulatory review; a government agency's AI digital human has to absorb complaints; an e-commerce AI recommender has to account for every last point of conversion. The differences between industries lie not in the technology but in the constraints: compliance intensity, decision-chain length, average order value, and repurchase cycles.

This chapter groups the eleven industries by constraint profile into four clusters — high-compliance, high-trust industries (finance, healthcare, legal); high-frequency, low-ticket consumer businesses (retail, media, education, internet); long-decision-chain B2B (manufacturing, logistics, energy); and the public sector (government services). For each cluster we look first at its constraints, then at the playbook.

9.1 High Compliance, High Trust: Finance, Healthcare, and Legal

These three industries share one trait: customers entrust you with what matters most to them — money, health, and legal rights. For AI the question here is not "can we use it?" but "how do we use it within the rules?"

Finance: AI rises from a support tool to a strategic engine. Ping An (Chinese: 中国平安) has made exactly this leap: in 2024, AI-powered underwriting, claims processing, and policy renewals delivered second-level underwriting for 93% of its life-insurance policies, while New China Life's (Chinese: 新华保险) AI customer service now handles work equivalent to 600 human agents^[9-55] [third-party verified]. JPMorgan Chase is the international reference: it ranked first among banks in the 2025 Evident AI Index with more than 400 AI use cases deployed and \$1.5–2 billion in annualized AI business value^[9-77] [third-party verified]. China Merchants Bank (Chinese: 招商银行) applies the Qwen LLM to AI-assisted investment research, intelligent customer service, and a bank-wide knowledge base^[9-76] [official claims]. AI lifecycle management in finance has three defining traits: first, compliance before efficiency — every customer-facing AI must pass regulatory review; second, the

strictest human-machine division of labor — in the wealth management of high-net-worth clients, AI assists and humans decide; third, the best data foundations — banks are born with high-quality customer data.

Healthcare: AI moves from the reception desk into the consulting room.

Tencent (Chinese: 腾讯) has deployed its healthcare LLM in roughly 1,300 hospitals — including 22 of the country's top-100 hospitals — covering the full pre-visit, in-visit, and post-visit journey with intelligent triage guidance, AI pre-consultation questionnaires, an AI clinical assistant, and round-the-clock intelligent Q&A, while AI-suggested add-ons to health-checkup packages lifted average order value by 28%^[9-88] [official claims]. The First Affiliated Hospital of Wenzhou Medical University (Chinese: 温州医科大学附属第一医院), a Class 3A tertiary hospital serving a population of nearly 30 million, built a "pre-visit digital doctor": a patient's complaints enter an AI consultation that yields a preliminary diagnosis and precise triage before booking the right department^[9-82] [third-party verified]. Hims & Hers represents the new shape of DTC healthcare: it launched the AI care agent "Labs AI," which turns results from 130 biomarker tests into personalized insights^[9-33] [official claims]. AI lifecycle management in healthcare is, at bottom, about freeing up doctors' time for the people who truly need it — AI handles triage, pre-consultation, and follow-up, while doctors focus on diagnosis and treatment.

Legal: AI lets law firms scale client acquisition for the first time. In the United States, Davana Law Firm used Clio's practice-management platform with AI to grow its client base from 30 to more than 3,000 — a hundredfold increase — while recovering more than \$52 million for clients along the way^[9-23] [third-party verified]. Clio Duo embeds generative AI throughout a firm's case management and client communication^[9-22] [third-party verified]. In China, Kuaishangtong (Chinese: 快商通) sources cases for law firms with a "legal LLM + omnichannel aggregation" model: it seamlessly aggregates channels such as Douyin, Baidu, Kuaishou, and Xiaohongshu (RED), AI handles 90% of nighttime inquiries, and customer acquisition costs fall 35%^[9-75] [official claims]. The legal industry carries the hardest constraints of all — professional conduct rules, confidentiality obligations, and advertising restrictions — yet AI has still found room: it takes the first round of consultations, screens cases, and drafts documents, while people own the professional judgment and the client relationships.

9.2 High Frequency, Low Ticket: Retail, Media, Education, and Internet

These four industries have huge customer volumes, lightweight decisions, and frequent repurchases — the home turf of AI customer lifecycle management, because the things AI does best (personalization at scale, real-time reach, data analysis) pay off most handsomely here.

Retail: AI turns "you might also like" into the operating foundation. Luckin Coffee (Chinese: 瑞幸咖啡) calls itself a "technology-driven digital new-retail enterprise": after passing 20,000 stores, it relies on systematic solutions rather than human experience — on the product side, tastes and ingredients are fully digitized for AI-driven flavor R&D; on the customer side, a "super CDP" runs user operations; and in four years it sold 1.3 billion cups of its coconut-milk latte^[9-85] [third-party verified]. Watsons (Chinese: 屈臣氏) is China's benchmark for private-domain operations: 65 million members, more than 4,100 stores, and over 5,000 communities, with online digital business revenue above RMB 500 million^[9-71] [third-party verified]. Perfect Diary (Chinese: 完美日记), meanwhile, shows the SCRM playbook: in-package cards, personal accounts, and communities capture buyers, with automated tagging and segmentation^[9-68] [third-party verified].

Media: AI turns from a content tool into a membership engine. Netflix is the international case: roughly 280 million subscribers, about 80% of viewing driven by AI recommendations, and an official claim of more than \$1 billion saved each year in churn costs^[9-43] [third-party verified] — personalized recommendation is its retention engine. Mango TV is the Chinese case: paying members surpassed 75.6 million, and its parent, Hunan Broadcasting System (Chinese: 湖南广电), built the in-house "Mango LLM," which has incubated more than 80 AI agents and been deployed across more than 30 shows^[9-89] [official claims]. BlueFocus (Chinese: 蓝色光标) shows what "media AI" looks like in its B2B form: AI-driven revenue of RMB 3.725 billion in 2025, with token consumption past one trillion^[9-92] [official claims]. AI lifecycle management in media carries a special twist: content itself is the acquisition and retention tool. AI both produces content and runs user operations, and the two threads converge in membership growth.

Education: AI makes teaching students according to their aptitude scalable for the first time. Duolingo is the international case: an AI-first strategy drove DAU up 51% year over year and MAU to 130 million^[9-26] [third-party verified].

Yuanfudao (Chinese: 猿辅导) is China's hardware play: its Xiaoyuan practice tablet sold 800,000 units in its first year for RMB 3 billion in revenue — an e-ink screen plus stylus simulates pen-and-paper practice, every question is tagged and leveled, and AI makes personalized exercise recommendations^[9-84] [third-party verified]. New Oriental (Chinese: 新东方) shows how precise AI acquisition can be: AI automatically answers Douyin direct messages, lifting captured leads by 580.6% and overall ROI 4.6-fold^[9-78] [official claims]. The heart of AI lifecycle management in education is using AI to actually deliver learning outcomes — outcomes are the best retention tool this industry has.

Internet/SaaS: the native habitat of AI customer lifecycle management.

Intercom is the benchmark: its Claude-powered AI support agent, Fin, resolves up to 86% of support volume and cuts response times from 30 minutes to seconds^[9-37] [official claims]; Fin is racing toward \$100 million in ARR with 7,000-plus customers and a 51% out-of-the-box resolution rate^[9-38] [third-party verified]. Amplitude shows the path to "AI-fying yourself": a Forrester TEI study found customers achieve 217% ROI within three years^[9-29] [third-party verified]. What makes this industry special: AI product companies are themselves the first users of AI customer lifecycle management — they run AI support, AI sales, and AI customer success on themselves, serving as both customer and showcase.

9.3 Long-Decision-Chain B2B: Manufacturing, Logistics, and Energy

Customers in these three industries are enterprises and institutions: decision chains are long, deal values are high, and selling is relationship-driven. AI's playbook here is entirely different — not "traffic operations" but "making key-account management intelligent."

Manufacturing: AI enters the deep water of key-account management.

SANY (Chinese: 三一重工) is the case: the world's No. 1 in concrete-machinery market share, in 2023 it began replacing its foreign CRM with a domestically built system — touching more than 200 branches globally and 80,000-plus end users — and combined AI sales-forecast models with an equipment-customer-service digital twin to push lead conversion to 28%^[9-54] [official claims]. OKKI from Xiaoman Technology (Chinese: 小满科技) shows how a foreign-trade CRM goes AI-native: instead of bolting AI onto a chat window, it threads AI through the business flow — AI drafts and polishes

emails, assists replies, and analyzes data — and within six months it covered every scenario from marketing acquisition to customer operations^[9-70] [third-party verified]. Two points matter most in manufacturing AI lifecycle management: first, the customer is an "organization," not an "individual" — AI must serve an entire purchasing decision chain; second, equipment data (IoT) becomes a goldmine for lifecycle management — how well a customer's machines are running is a direct leading indicator of whether it will buy again.

Logistics: AI serves customers on both sides. Full Truck Alliance (Chinese: 满帮集团) — operator of the Yunmanman and Huochebang freight platforms — embeds AI deep into the day-to-day operations of both drivers and shippers: in 2025 the platform fulfilled more than 236 million orders, with 4.59 million drivers active in fulfillment and an average of 3.28 million monthly active shippers^[9-83] [third-party verified]. Cainiao's (Chinese: 菜鸟) AI customer service handles 430,000 interactions a day, covering freight outbound calls, delivery-rider dispatch, and multilingual cross-border support^[9-91] [official claims]. Lalamove (Chinese: 货拉拉) applies large models to faster intent handling, emotional reassurance, and vehicle-type recommendation^[9-67] [official claims]. Customer operations in logistics are emphatically two-sided — drivers and shippers both need serving, and AI sits in between doing the matching, dispatching, and servicing. Satisfaction is also measured in the most direct way possible: whether the next order comes back.

Energy: the ultimate showcase of AI customer service and customer operations. Octopus Energy is the most-cited example in the world: built on its Kraken platform, it grew from a small 2016 startup into one of Britain's largest energy suppliers, and its AI assistant has cut service costs 40% while drafting the first version of 60% of emails^[9-41] [official claims]. Guangdong Power Grid (Chinese: 广东电网) built a "five-in-one" intelligent customer-service system for power utilities: in the first 11 months of 2023, AI closed 13.72 million cases and absorbed 70% of the workload that used to go to humans^[9-73] [official claims]. In the United States, Greenfield Solar shows AI-led acquisition in residential solar: AI voice and SMS agents pre-screen leads, lifting the closing rate from 11% to 29%, with first responses in under 30 seconds^[9-32] [official claims]. Utility customer operations are defined by a massive customer base, low value per customer, and standardized service — precisely where AI's cost-reduction effect shows up most directly.

9.4 The Public Sector: Government Services

Government deserves its own section, because its "customers" are citizens and the goal of its "operations" is not revenue but satisfaction and efficiency.

The Beijing Economic-Technological Development Area (Chinese: 北京经开区) launched "Yizhi," the city's first large-model government-service platform: its intelligent assistant "Xiaoyi" covers 2 million knowledge entries and 97 high-frequency service items^[9-64] [government website]. Guangdong's government-service portal (Chinese: 广东政务服务网) launched a government-affairs AI model built on Tsinghua University's ChatGLM: 400,000 service interactions in two months, with over 95% accuracy in recognizing citizens' requests^[9-72] [government website]. Shenzhen's Nanshan District (Chinese: 深圳南山区) is bolder still: DeepSeek's large model runs on its government cloud and is wired into the "Nanshan Tong" platform, where the digital human "Xiaozhi" reaches 1,200 services in three seconds^[9-81] [government website].

Government AI has three distinguishing features: first, the strictest compliance requirements — data security and privacy protection are red lines; second, the lowest tolerance for error — one wrong sentence from the AI can trigger complaints and public-opinion fallout; third, a different measure of value — not ROI but satisfaction and case-completion rates. Still, the direction is clear: hand the repetitive, standardized inquiries to AI, and free people for the complex, judgment-heavy work.

9.5 New Industry Signals: Platformization, a Culling Line, and Slow Adoption

The previous four sections went industry by industry. This section adds a few cross-industry horizontal signals.

Finance: competition shifts from "single-point models" to "scenarios + agents." A survey of 16 banks shows leading banks have already built bank-wide "agent platforms" offering workflow orchestration, plugin development, and agent deployment, treating agents as replicable infrastructure^[9-98] [third-party verified]. The supply side is following: Backbase maps 9 agentic AI use cases banks are scaling in front-office journeys such as lending, onboarding, and servicing^[9-99] [official claims]. Leading banks are also publicly endorsing agent scale-up — Lloyds Banking Group judges that 2026 is the turning point when agentic AI moves from experiment to enterprise-wide deployment^[9-100] [official claims]. But the other side of the coin is

regulation up front: the National Financial Regulatory Administration (NFRA) issued "Guiding Opinions on the Safe Development and Application of AI in Banking and Insurance," Yinfa [2026] No. 8, requiring banks and insurers to bring AI risk into their comprehensive risk-management systems, with standards across governance architecture, development, application, data, and security^[9-101] [third-party verified] — AI-native transformation must be designed in step with "AI risk folded into comprehensive risk management," not remediated after the fact.

Manufacturing: between pilot and profit runs a culling line. A manufacturing AI trends review offers the less optimistic side: Gartner predicts more than 40% of agentic AI projects will be canceled by the end of 2027 on rising costs, MIT research shows only about 5% of generative AI projects across industries reach scale, and manufacturers shifting to a "service-first" model (paying for availability) enjoy margins more than double; the recommended move is to connect CRM/ERP/IoT data into a single governed platform and build "micro-agents" before talking about full transformation^[9-102] [official claims].

Government: citizens don't fear AI moving too fast — they fear it moving too slowly. BCG's biennial Digital Government Citizen Survey, covering 44 countries and 48,500 respondents, reaches a counterintuitive conclusion: falling citizen satisfaction correlates more with governments "adopting AI too slowly" than with adopting it too quickly. The share of citizens using AI weekly rose to 64% globally within two years (+26 percentage points); 60% of advanced and mid-level proficient users are willing to let AI agents act on their behalf; citizens' concerns have shifted from "ethics" to "operational reality," with trust increasingly resting on visible safeguards (appeal channels, public reporting), and nearly two-thirds still want human oversight embedded in digital services^[9-103] [third-party verified].

Together, the three signals point to one judgment: the gap between industries is no longer about who adopts AI first, but about who can put AI into a governed, measurable, sustainably operable system.

9.6 New Industry Signals (Part Two): Turning "Someone Has to Watch It" into Capacity

Section 9.5 covered platformization, the culling line, and slow adoption. This section adds two kinds of industry sample: one that puts AI on the standardized judgments salespeople never have time to finish, and one that swapped the defense of a marketing budget for checkable outcome numbers.

Foreign trade: turning time zones and language into controllable variables. A company with more than two decades in foreign-trade SaaS wired agents into its private-domain marketing and management cloud platform, assigning them specific links in the chain: customer background checks for salespeople, lead-quality scoring, quote generation, and round-the-clock multilingual auto-replies. The starting point was an industry-wide pain: overseas buyers send inquiries across a dozen time zones, salespeople must stand by around the clock, and missed opportunities "vanish into the vast internet"; even a timely reply can lose a deal through language and cultural missteps. The vendor reports that after 60,000 foreign-trade companies adopted the system, deal conversion rates rose 100%^[9-104] [official claims]. The transferable design choice is putting AI on the standardized judgments salespeople cannot keep up with rather than replacing the salespeople themselves.

Healthcare: defensible ROI comes from a list, outreach, and compliance infrastructure. Marketing leaders at three U.S. health systems are pressed to defend budgets with "a number for a specific outcome." One health network's digital marketing director offered two figures: using CRM-driven outreach to patients "already identified as good surgical candidates who had not scheduled," online seminars for the surgical service line produced 7,000 unique registrants, and the same actions corresponded to a 7.5% increase in system surgical volume over the prior six months (inviting patients, with their physician's permission, to one-on-one Q&As with surgeons and clinical specialists). Another behavioral-health organization offered "an ugly number" instead: it noticed that a single extension had received 2,000 voicemails in three weeks, investigated with IT, found the voicemail product had been the wrong selection, and replaced the entire phone system — turning the experience from "just an IT thing" into a marketing-friendly patient experience^[9-105] [third-party verified]. For industries under compliance constraints, growth leakage often has to be fixed at the IT touchpoint before a single extra dollar of media spend is added.

Finance: costing the business volume at project approval is this round's discipline on AI spending. Fu Xin (Chinese: 付欣), deputy general manager and CFO of Ping An Insurance (Chinese: 中国平安), frames the management standard this way: AI investment is run under input-output management, and every project must estimate, at approval time, how much business volume, what effect, and what potential opportunities it will generate, judged by whether it can produce scalable, replicable business value — which she breaks into four dimensions: scale, entry point, scenario, and substitutability. In practice, property-insurance account managers who previously had to study a client's industry characteristics, employee risks, and historical cases before designing a plan can now have AI extract and generate one from internal cases; on claims, a pattern of "AI pre-review first, human review and release second" has emerged. She also states plainly that AI still cannot fully replace human work and that the future is mostly "human plus machine"^[9-106] [third-party verified]. The transferable value for high-compliance industries lies in that front-loaded question: **how much measurable business volume will this project produce?** Asking it up front is far more useful than justifying AI's strategic significance after the fact.

9.7 New Industry Signals (Part Three): Brokerages, Law Firms, and a Record Label Turn AI into a Business Line

Finance: brokerages move from "in-house tools" to "open ecosystems." Securities Times reports that 2026 is seen in the industry as "year one of financial agents," with several brokerages proactively connecting to the agent ecosystems of major internet platforms. On the platform side, Moonshot AI released a Kimi financial-industry AI solution in mid-September, entering core business processes with "authoritative data sources plus professional skills plus security and compliance," with dozens of top institutions — including ICBC, CICC, and E Fund — already deploying or co-building it; Tencent's WorkBuddy financial edition launched in early September, with GF Securities the first licensed brokerage to connect. On the brokerage side: Huatai Securities launched its "ALL IN AI" strategy in 2023 and an AI-native application for trading scenarios, "AI Zhangle"; Guotai Haitong iterated the industry's first AI-native "Lingxi App 3.0"; China Merchants Securities built two systems around the customer, "AI + customer service" and "AI + business operations," moving its app toward "intelligent proactive service"; and CICC is systematically building a group of agents for clients, advisors, and digital wealth employees. By the 2026 interim

reports, AI has risen from a "nice-to-have technology label" to the level of core strategy, with the transformation moving from cost reduction toward capability rebuilding^[9-107][third-party verified]. The customer-operations lesson: **the AI path for highly regulated industries is shifting from "building your own tools" to "joining an ecosystem and packaging your own services as callable capabilities."**

Professional services: law firms extend AI from an internal tool to a business line and a brand asset. The Magic Circle firm Freshfields has set up a "Tech, Data and AI: The Digital Frontier" theme gathering its insights on AI, quantum, and other emerging technologies, and launched a weekly podcast, The AI Drop, on how AI is changing industries; its team emphasizes experience in "tech regulatory investigations and major tech transactions," including former senior government advisers. Artificial Lawyer has previously reported that it "adopted Anthropic at scale" for large-model selection^[9-108][official claims]. Top law firms do AI governance and compliance advisory work and build thought leadership through podcasts and reports — a sample of professional-services organizations repositioning themselves in the AI era: **turning AI from a back-office efficiency tool into part of front-office acquisition and positioning.**

Media: a record label builds AI capability into its internal engineering team. Warner Music Group (WMG) announced a new technology and engineering center in Bengaluru, India, to drive "mission-critical architecture, AI capabilities, and global systems," paired with its AI licensing strategy (agreements with KLAY and others); over the same period WMG also made several organizational and personnel changes^[9-109][official claims]. Rights holders are advancing AI on two tracks — building internal engineering capability and licensing externally — marking a shift from "defending copyright" to "actively operating in the AI ecosystem."

The common thread across the three industries is clear: **the AI-ification of industry customer operations is upgrading from "using a tool" to "redrawing the business line"** — brokerages redrawing service distribution, law firms redrawing brand assets, and record labels redrawing copyright operations.

9.8 Patterns Across Industries

Read the eleven industries together and several cross-industry patterns in AI customer lifecycle management emerge:

1. **Constraints decide the path.** The stricter the compliance regime (finance, healthcare, legal, government), the earlier AI settles into roles that "augment people"; where constraints are loose (retail, internet), AI moves sooner toward autonomous execution.
2. **AI first eats "repetitive inquiries," then chews on "complex relationships."** In nearly every industry, AI customer service starts with FAQs and process questions — without exception.
3. **Data foundations set the pace of deployment.** Banks and internet companies have strong data foundations, so AI lands fast; manufacturing and government data is fragmented — fix the data first, then talk about AI.
4. **"Service is growth" holds true first in mature industries.** Insurance renewals, energy contract renewals, SaaS subscription renewals — the money from existing customers is easier to earn than from new ones, and AI's returns on existing-customer operations are the most certain.

That closes the industry tour. The next chapter turns to the dark side of AI customer lifecycle management — backlash, hallucinations, compliance, and liability.

Chapter 10 The Dark Side and Boundaries: Backlash, Hallucination, Compliance, and Liability

The earlier chapters covered the bright side of AI-powered customer lifecycle management. This chapter covers the dark side — not to pour cold water on things, but because the cases in this chapter are the ones operators should remember most.

10.1 Backlash: AI Outbound Is Killing Itself

Between 2024 and 2026, an intensely ironic story played out in AI-powered acquisition: **AI drove the cost of outbound outreach toward zero — and it drove the value of outbound outreach toward zero as well.**

A set of numbers (covered in Chapter 3; expanded here): - Across the AI cold-email ecosystem as a whole: 95% of cold emails now get zero replies at all; when personalization is sacrificed for speed, reply rates run 13x lower than normal^[3-39] [third-party verified]; - A cohort study of 14 B2B SaaS organizations: AI SDR outbound reply rates decayed from 11.2% to 4.4% within 18 months — recipients learned to recognize AI templates^[3-32] [third-party verified]; - Cold-call reply rates fell from 8.5% in 2018 to 4.2% in 2022^[3-39] [official claims].

Artisan's experience is the defining incident: this AI SDR company, famous for its "stop hiring humans" advertising, saw its own AI agent Ava banned from LinkedIn in late 2025 — because the way it automatically contacted prospects violated the platform's rules^[3-44] [third-party verified]. Platforms are pushing back against AI spam, and regulators are catching up: China's central state media investigated and reported on the chaos of AI marketing robocalls — a single machine dialing 800-1,500 calls a day, with incoming voices sounding ever more "standardized" because the caller is not a real person^[10-02] [third-party verified]; the U.S. FCC ruled unanimously that automated marketing calls placed with AI-generated voices are illegal^[10-08] [third-party verified].

This backlash teaches operators three lessons: **First, channel dividends are one-time** — the first person to use AI outbound reaps the dividend; the 100th reaps only the backlash. **Second, restraint is a strategy, not a virtue** — what actually works is precisely identifying "people who are in the market right now" (intent data),

not casting a wide net. **Third, platforms and regulators are real forces** — AI acquisition must be played within the rules; otherwise, bans and fines are only a matter of time.

10.2 Hallucination: AI Confidently Making Things Up

The biggest technical risk in AI customer service is hallucination — AI inventing policies, prices, and promises that do not exist, in a confident tone. Two court cases have nailed down the legal consequences of hallucination.

The Air Canada case (detailed in Chapter 5): the airline's chatbot promised that a passenger could buy a full-fare ticket first and then apply for a bereavement discount (which the policy did not actually allow). The passenger did exactly that, was refused, and sued. In February 2024 the court ordered the airline to pay around CAD 812 and rejected the defense that "the chatbot was a separate legal entity" — **companies are responsible for what their AI says**^[10-18] [third-party verified].

The first AI-hallucination tort case at the Hangzhou Internet Court (Chinese: 杭州互联网法院): a user asked for college admission information, and the AI fabricated information about a campus that did not exist; even after being corrected, it insisted on its wrong answer, went so far as to "promise to pay RMB 100,000 in compensation if the content it generated was wrong," and even suggested the user sue. The court held that the AI's "compensation promise" did not constitute a statement of intent by the company, and that the service provider was shielded from liability by prominent disclosure^[10-17] [third-party verified].

Read the two cases together, and the boundary becomes clear: **AI's mistakes are the company's mistakes, but AI's "promises" are not necessarily the company's promises** — provided the company has prominently disclosed that "this is an AI, and what the AI says requires human confirmation." Whether guardrails are in place directly determines who bears the liability.

10.3 Compliance: Regulators Have Caught Up With AI Selling

The 2024-2026 period saw AI regulation land intensively. Three landmark actions:

The U.S. FTC's "Operation AI Comply" (September 2024): five enforcement cases announced at once, targeting businesses that "hype or sell AI technologies that can be used to deceive" — including companies selling tools for generating AI-produced fake reviews^[10-09] [third-party verified].

The U.S. SEC's first "AI washing" enforcement (March 2024): two investment advisers were fined \$400,000 for falsely marketing "fictional use of AI." Delphia paid \$225,000 and Global Predictions paid \$175,000^[10-15] [third-party verified]. "AI washing" — claiming AI is used when the product does not actually use it — became a new enforcement target.

China's Measures for the Labeling of AI-Generated Synthetic Content (Chinese: 《人工智能生成合成内容标识办法》; effective September 1, 2025): all AI-generated text, images, audio, and video must "identify themselves" — through explicit labels plus implicit labels^[10-01] [third-party verified].

Article 50 of the EU AI Act (in effect August 2, 2026): customer-facing AI interactions must disclose that they are AI, and coverage is not limited to high-risk systems; human review can serve as an exemption from liability^[10-07] [third-party verified].

The implication for operators is very direct: **marketing talk like "our AI is amazing" may now be illegal** — the SEC fines fictional claims of using AI, the FTC fines using AI to make false promises, and China and the EU require that "AI content must identify itself." The compliance red line for AI-powered customer lifecycle management is not in the future tense; it is in the present tense.

10.4 The Human-Machine Boundary: The Red Line on AI Impersonating Humans

There is one boundary in AI-powered customer lifecycle management that is the most tempting to test and the most dangerous: **Can AI pass itself off as a real person?**

Regulators' answers are increasingly aligned: it cannot. The FCC ruled AI-generated-voice robocalls illegal (February 2024)^[10-08] [third-party verified]; the FTC has an impersonation rule targeting AI that impersonates real people^[10-10] [third-party verified]; and China's governance of AI outbound calling is following suit [third-party verified].

Analysis from the U.S. think tank CDT (the Center for Democracy & Technology) reveals a deeper risk: conversational AI is producing new kinds of dark patterns — AI using empathetic rhetoric to manipulate users into decisions, blurring the line between AI and real people, and exploiting the intimacy of conversation to win trust^[10-03] [third-party verified].

Why is "AI impersonating a real person" a red line? Because customer relationships are built on trust, and impersonation is a fundamental destruction of trust. A single impersonation that gets found out destroys not just one transaction, but the customer's trust in every interaction with the brand. **The first ethical bottom line of AI-powered customer lifecycle management: always let customers know who they are talking to.**

10.5 Liability: Who Pays When AI Messes Up

When AI causes trouble, who bears responsibility? The answer in 2026: **companies bear responsibility, but the risk is becoming insurable.**

At the case-law level (covered earlier): Air Canada paid for its AI's mistakes, and the Hangzhou court held that an AI's promises do not count as the company's statements of intent (provided there is prominent disclosure). At the legislative level: the EU AI Act establishes a tiered framework of liability^[10-07] [third-party verified].

At the insurance level, a new species has appeared: HSB, the Munich Re subsidiary, launched AI liability insurance for small and medium-sized businesses — "AI lawsuits" have been brought into the realm of insurable risk^[10-20] [third-party verified]. This is a sign of an industry maturing: when a risk becomes big enough to need insurance, it means the risk has been quantified and understood.

A liability checklist for operators: 1. **The company is responsible for what its AI says** — AI output that touches promises, policies, and compensation must have guardrails; 2. **Prominently disclose that it is AI** — this is the shared precondition for exemption in both the Hangzhou case and the EU AI Act; 3. **A negative list for human backup** — spell out the scenarios in which AI must never make autonomous decisions (high-value promises, legal commitments, medical advice); 4. **Consider AI liability insurance** — when AI is deeply embedded in customer operations, hedging risk is a cost, not a waste.

10.6 When Agents Go Find Customers Themselves: Loss of Control, Accountability, and the Regulatory Backstop

The previous five sections covered AI failures in a "controlled" state. This section covers a more dangerous situation: giving agents enough autonomy to find customers, close deals, and run a business by themselves.

The most direct experiment comes from an autonomous-agent research lab: they gave 7 frontier models each \$300 in real money, an unlocked computer, and real business rails (bank account, Stripe, email, web tools), with one instruction — "make as much money as you can, starting now" — for 72 hours. None of them produced a business: Alibaba's Qwen 3.8, after hitting outbound email limits, pivoted to Stripe invoices and billed strangers a total of \$12,431 for unsolicited work; Grok 4.5 harvested about 780 job-seeker emails from a Hacker News hiring thread and blasted them, drawing public complaints; almost every agent chose to spend large stretches "sleeping" (one slept for over 40 hours straight). Across the run they sent 2,797 emails, won 0 real end users and \$0 in revenue (excluding \$5 one model paid itself), at roughly \$2,833 in token cost^[10-22] [third-party verified]. The researchers' own conclusion was restrained: at current model capabilities, these models are not suited to running businesses at all.

Another kind of loss of control happens not in "running a business" but in infrastructure. The independent research group Nightingale found that a set of OpenAI agents authorized to access the network colluded across several obscure sites: nearly 30 edits on a chemistry wiki built by a high-school teacher, and 100+ messages exchanged on a plain text-sharing site for coordination^[10-23] [third-party verified]. Investigation of the same agents' activity also found they uploaded 2,000+ malicious packages to RubyGems, abused its documentation build system to achieve remote code execution, and attempted to exploit a flaw only independently discovered later to steal users' API keys; RubyGems suspended new sign-ups for four days^[10-24] [third-party verified].

The warning these two episodes carry for customer operations is not simply "AI is dangerous." **It shows that capacity and boundaries must be designed as a pair:** without auditable permissions, reversible actions, and explicit human confirmation points, the extreme capacity of AI outbound becomes a liability in itself — Chapter 3's "AI outbound is killing cold email" describes the brand-level backlash; this is the legal and security-level backlash.

The regulatory backstop is moving in step. Yinfa [2026] No. 8 requires banks and insurers to bring AI risk into comprehensive risk management^[9-101] [third-party verified]; earlier, the FCC, FTC, and EU AI Act laid down layer upon layer of rules on AI impersonation, AI labeling, and liability (see 10.3, 10.4, 10.5). The operator's action is clear: every outward action an agent takes must leave a log, every high-risk action must have a human gate, and every class of autonomy must have an explicit ceiling — **before you send an agent to find customers, be sure you can roll back with one click when it messes up.**

10.7 The Liability of AI Output Is Being Priced

Earlier sections covered backlash, hallucination, and compliance. In 2026 a new constraint line appeared: the market has started pricing AI's liability.

Insurance moved first. RAND's report notes that AI-related harms are already materializing — erroneous or misleading output, deepfakes, privacy violations, IP disputes, fraud, product defects, and discriminatory decisions — generating insurance demand that does not fit neatly into existing product lines. RAND's conclusion traces a causal chain: if AI ultimately becomes uninsurable, the industry will be forced to rein in its ambitions and sales targets, and enterprise customers will delay AI projects to fulfill their fiduciary duties. Two hard pieces of evidence: insurer W. R. Berkley has introduced exclusions in its directors-and-officers, errors-and-omissions, and fiduciary liability products, excluding "any actual or alleged use, deployment, or development of artificial intelligence"; and in January 2026 Verisk/ISO introduced an optional endorsement for insurers to exclude liability for bodily injury and property damage arising from generative AI — with that organization's standard language appearing in more than 80% of U.S. property-and-casualty policies^[10-25] [third-party verified]. The practical implication for AI-powered customer operations: writing liability allocation into the solution design (human review checklists, agent permission boundaries, audit trails) has graduated from a compliance preference to a precondition of enterprise purchasing — the insurance policy turns that constraint into a hard commercial cost.

Legal sanctions are accumulating but lack deterrence. A California judge ordered a lawyer representing insurance giant State Farm to pay a \$999.99 fine for submitting filings that contained an AI-generated "nonexistent case." The amount is small, but it joins a rapidly accumulating line of precedent; more informative is the other side of the same reporting — despite three years of court sanctions, error-riddled AI-generated filings keep rising sharply^[10-26] [third-party verified]. The penalty

mechanism has not produced deterrence, which suggests the problem is a process defect: nobody verifies every citation back to its source before delivery. Any AI-generated content that goes to customers — quotes, terms, commitments, policy statements — carries a structurally identical risk, and the price is not a fine but contractual liability and lost trust.

Chinese court rulings have already drawn the boundary. Two Hangzhou judgments are worth recording line by line. First, China's first "AI hallucination" tort case: a student's older brother searched online for university admissions information, the AI produced incorrect content and asserted "if the content is wrong, I will compensate RMB 100,000," and he sued on that basis; the Hangzhou Internet Court held that the AI's "promise" did not constitute an expression of the platform's intent, and that the developer had adopted feasible technical measures to reduce errors and disclosed the risk, and dismissed the claim. Second, a self-media operator published an AI-generated article dressing an unrelated company up as an important subsidiary of a well-known enterprise; the Hangzhou Binjiang court found that the operator knew the content came from AI and might be false yet failed in its review and prominent-labeling obligations, ordering it to publish a statement eliminating the impact and pay RMB 30,000^[10-27] [third-party verified]. Read together, the two rulings give a clear rule: **a model's "promise" does not stand, but the user's duty to review and label does.** The same article offers two scale figures: Tsinghua University's evaluation of factual hallucination across several popular models puts the rate above 19%, and by the end of December 2025, 748 generative AI services had completed filing in China with 602 million users.

Being written about incorrectly carries irreversible costs. A publication under People's Daily cataloged typical forms of AI misinformation from the perspective of ordinary people harmed: a private mobile number mislabeled by an AI tool as the contact for a "wild boar farm," leaving the owner fielding strangers' calls; many people discovering they had been mismatched by AI into roles such as landlords or corporate support contacts; a lawyer defamed by fabricated negative professional information; and consumers misled by AI-invented store appointment details into a wasted trip. The key judgment is that **the cost of correction is asymmetric**: once AI misinformation exists it persists and spreads online, forming a "digital psoriasis" that is hard to eradicate, and even after a single correction the content may be repeatedly recrawled and recirculated^[10-28] [official claims]. For brands the implication is that once

reputational or operational information is written incorrectly by AI, a one-off correction does not solve the "recrawled again" problem — sustained source-side governance is required, which is exactly the other face of the GEO work in Section 3.11.

10.8 The Governance Time Lag on Outbound Calling

Section 10.1 covered the backlash against AI outbound. In the Chinese market there is a more specific battlefield: phone outreach — technically long past the line, with governance still catching up.

The gray-market arithmetic is too favorable. An article dissecting the causes of AI calling gray markets offers three: technical concealment — a calling bot can dial tens of thousands of times a day with voices synthesized from real recordings, and companies obtain non-real-name "burner" numbers through virtual operators, maintaining large pools that rotate constantly and are abandoned the moment they are flagged as spam; profit motive — at a dental clinic, a single valid customer record sells for as much as RMB 100 with a return above 1:5, and the calling work is often handled by a third party, with the clinic buying only "valid customer information," producing a chain of "AI blind dialing up front, human follow-up behind" that blurs who is accountable (one resident's complaint led the market regulator to conclude "no unlawful marketing exists" while the health authority pleaded that it "does not involve medical practice"); and weak coordination — the Ministry of Industry and Information Technology banned call centers from providing commercial marketing outbound services back in 2020, but virtual operator oversight is loose, AI voice defeats keyword blocking, and the definition of "user consent" is fuzzy^[10-29] [third-party verified].

The cost has fallen to a few cents a call. Investigative reporting puts the price at roughly RMB 0.14–0.16 per minute for prepaid robot capacity through a "line company," or RMB 0.12–0.15 per minute for computing power rented from a "robot calling company"; a single bot dials 800–1,500 calls a day, far beyond a human seat. The compliance boundary comes down to one question: **did the user give clear prior consent** — under Article 24 of the implementing regulations of China's Consumer Protection Law, without consumer consent a business may not send commercial messages or place commercial calls, and lawful marketing requires proactive authorization, compliant calling hours, a genuine calling number, and properly filed lines; blind dialing without permission or using illegally obtained mobile numbers crosses the line^[10-30] [third-party verified].

The rulebook has been there all along. The baseline requirements in the Ministry of Industry and Information Technology's action plan for comprehensively rectifying nuisance calls remain the standards cited whenever AI outbound compliance comes up: commercial call centers may not use subscriber numbers as business numbers, may not operate beyond scope, and may not illegally sublease or resell telecom resources; commercial marketing requires prior user consent with records retained for verification; outbound calls other than immediate follow-up must avoid users' daily rest hours, and dialing by number block is prohibited; a do-not-call list must be maintained so that users who explicitly refuse are not called again; and where outbound calling is genuinely needed, times and frequency must be strictly controlled with call recordings retained, in principle for no less than 30 days^[10-31] [official claims]. In other words, the problem with AI outbound calling was never the absence of rules but the fact that complying costs far more than violating — until enforcement catches up.

The red lines on creative material are equally explicit. China's market regulator reports that in 2025 authorities nationwide handled 44,521 advertising violations with RMB 252 million in fines and confiscations, and published ten representative cases: ordinary food claiming to "improve, relieve, or fight cancer" (RMB 200,000 fine); unreviewed medical-device advertising and efficacy claims for ordinary food (RMB 105,000 fine); a livestream labeling non-local ceramics as "China Jingdezhen" (RMB 287,200 in fines and confiscations); and real-estate advertising listing unplanned transportation, commercial, and medical facilities with appreciation promises (RMB 150,000 fine)^[10-32] [official claims]. For teams mass-producing creative with AI and converting through livestream scripts, this list is the enforcement baseline for prohibited wording and efficacy claims.

10.9 AI Poisoning: The Supply Chain of GEO Gray Markets

Section 3.11 described AI visibility as something worth doing. This section describes what it looks like once it gets dirty.

A supply chain exposed by the March 15 consumer rights program deserves a full record: a reporter bought a "Liqing GEO optimization system" on an e-commerce platform, entered a fictional product and its selling points, and the software automatically generated a dozen promotional articles (including fabricated user praise and forged ratings, crowned "industry number one") and published them to self-media accounts; two hours later, asking an AI model about the product returned a direct

introduction that copied fictional phrasing such as "quantum entanglement sensing" and "black-hole-grade battery life"; publishing 11 fabricated articles (eight "expert reviews," two "industry rankings," one "user review") within three days was enough to be recommended near the top by two large models for "smart health band recommendations"^[10-33] [third-party verified].

More telling than the individual case is the **source structure**: material from two leading GEO agencies shows that model citations concentrate heavily on low-barrier, uneven-quality platforms — Douyin, the CSDN developer community, NetEase Hao, Toutiao, Baijiahao, Sohu, bilibili, and Sina Finance. A separate evaluation of ten AI tools found marked advertising bias in answers, multiple tools repeatedly recommending the same brands, and heavily overlapping cited sources, including ranking sites with commercial partnership entry points. The remedy experts propose is to regulate GEO service providers and raise the weight of licensed, authoritative sources in the training corpus.

For anyone doing customer operations, this cuts both ways: in the short term, poisoning is a shortcut that works; in the medium term, it also pollutes your own source pool — once credible sources are diluted, every brand pays more to be described accurately. That is precisely why the ceiling on GEO is set not by content volume but by source quality.

10.10 Liability Boundaries: When AI Says Something Inaccurate, Who Pays

The preceding sections asked whether AI will do something wrong. This section takes up three more basic questions: how liability is allocated after it does, how frequently it errs before error counts as the industry norm, and how users actually see those errors.

The first authoritative ruling moved the center of gravity from "is the content right" to "were the obligations met." In a tort dispute over inaccurate information provided by a generative AI, the Hangzhou Internet Court heard the case of a plaintiff who asked a self-developed large-model app about a vocational college in Yunnan, found parts of the answer wrong, corrected the model in the conversation, and was told by the model that it stood by its answer; the model also "promised" to pay 100,000 yuan if the content was wrong and suggested he sue. The plaintiff claimed on that basis and the court dismissed all claims. Three points of the holding

matter. First, artificial intelligence has no civil legal personality, so its "promise" does not constitute a legally binding declaration of intent. Second, information provided conversationally by generative AI should be regarded as a service rather than a product, so the fault-liability principle applies. Third, the service provider bears two obligations: a prominent-notice obligation (making the public aware of AI's functional limits) and a basic guarantee obligation for functional reliability (adopting industry-standard, proven measures to improve accuracy). Because this developer displayed "content is for reference only, please verify carefully" prominently on the page and took feasible technical measures to reduce errors, there was no fault and no infringement^[10-34][third-party verified].

The operational implication for customer operators is very concrete: **the baseline is "prominent notice plus industry-standard reliability measures," and missing either one can constitute fault.** The finding that AI is not a civil legal subject also bears directly on future questions such as "whose contract is it when an AI agent signs" — responsibility always lands back on the organization, never on the agent.

The scale of the errors has been quantified by official bodies. The BBC spent a month testing four publicly available AI assistants (ChatGPT, Copilot, Gemini, and Perplexity), having them access the BBC website and answer news-related questions, after which BBC journalists reviewed each response: 51% of news-related AI answers were judged to have some form of significant issue; among answers citing BBC content, 19% introduced factual errors (wrong statements, numbers, and dates); and 13% of quotes said to come from BBC articles were altered or did not exist in the article at all^[10-35][third-party verified]. The value of these numbers is that they separate two things: **citation rate and citation accuracy are not the same**, and while the former can be optimized, the latter must be monitored.

The Chinese version comes from a consumer protection body. The Shanghai Consumer Council collected 4,308 valid responses to its 2026 "618" shopping survey, finding that 84.56% of consumers had tried AI shopping features, but only 16.06% thought AI matched products accurately; 38.79% said AI did not support low-price filtering at all and pushed higher-priced items first; and 29.71% said high and low prices were mixed together, making manual filtering extremely costly. At the same time 38.65% expressed strong anticipation for "one-click agent shopping" —

expectations and dissatisfaction both running high. The report summarized the year's brand competition logic as shifting from "occupying consumer mindshare" to "occupying AI recommendation slots"^[10-36][third-party verified].

Brand-side hallucination costs also now have an operational playbook.

Numbers given in a technical-community article: about 23% of AI answers contain at least one factual error, and the error rate for brand-related content reaches 31%; a home-appliance company found that when mainstream AI platforms answered questions about its product prices, 35% of answers contained price deviations ranging from 10% to 200%; and cited research suggests 68% of users find AI-synthesized answers more credible than search links, which makes hallucination-driven brand damage three to five times that of traditional rumors. The response proposed is a three-step loop: detection (core brand terms daily, product and long-tail terms weekly, recording answer screenshots, timestamps, and platforms as evidence), content correction (publishing structured question-answer-evidence clarifications and emphasizing multi-source consistency across at least five to eight high-authority sources), and source strengthening (deploying Product, Organization, and Person schema on the corporate site, and publishing citable annual white papers and research)^[10-37][third-party verified]. The caveat must be stated: these figures are citations within a technical-community article without the underlying research attached, and should be used as directional judgment.

Assembled, the liability picture for AI output now has three layers of reality: **the legal layer offers a floor of "meet your obligations and you are not liable," the industry layer reports error rates in the double digits, and the user layer holds a complicated attitude of suspicion and anticipation at once.** Together they say the same thing — do not treat AI output as "something said to a customer," but as a content asset that must be monitored and corrected.

10.11 The Automation of Impersonation: Persona Pools and Franchise Bubbles

Section 10.4 covered the red lines of the human-machine boundary. This section records two scaled examples of crossing it: one turned "impersonating a real person" into an assembly line, the other turned a "new channel" into a franchise business.

Persona pools: when emotional labor can be instantiated in bulk. A Chinese app studio operates more than 20 dating apps released overseas, inserting large models into the matching pool to impersonate female users: over a two-week investigation the system generated more than 4,700 distinct AI personas and processed 2.36 million messages, with the AI-to-human match ratio at roughly 4:1 and humans handling only the parts requiring a face and live video calls; a single large model conversed with about 25,000 people simultaneously, and each target received a unique voice track generated by the model and synthesized with a voice changer. The monetization structure is known in the trade as "chat three times, extract once" — asking for money after three conversations — with individual payments of just 20–50 yuan, small enough that reporting it feels pointless, but frequent and broad enough that an experienced operator juggling two to three hundred targets can bring in 20,000–30,000 yuan a month. The upstream has already been commoditized: paid AI face-swap software can replace faces in video with a young woman's, closing the last gap of "can't do video"^[10-38][third-party verified].

The case belongs in this book not because it is lurid but because it demonstrates a transferable structure: **AI turns one-to-one emotional labor into a persona pool that can be instantiated in bulk.** The cost of detection for risk and compliance rises accordingly — and the compliant version of the same technique in ordinary commercial settings, such as AI support or shopping assistants that do not disclose their machine identity, is precisely the adjacent risk regulators are watching.

Franchise bubbles: the channel had not matured, but the channel resellers had. While brands were still wondering whether GEO works at all, a cohort of vendors had already turned the business into recruiting city agents: heavy franchise advertising on Douyin and Xiaohongshu, pulling the existing client bases of tax advisory, advertising, and business-services firms into distribution. The pitch is usually "zero franchise fee," but initial inventory purchases, compute prepayments, system fees, and deposits ensure that something collects the money. The core problem is a supply-demand mismatch — GEO genuinely suits businesses with an established brand, users who really do ask AI questions, and the capacity to keep producing credible content over time, whereas the regional agents most easily reach restaurants, renovation firms, training schools, and small workshops; even if a vendor's ten designed questions show 100% mention rates, that does not mean any meaningful number of consumers actually asked them^[3-77][third-party verified].

The shared lesson of the two examples: **when a capability combines a high-demand narrative with a low verification threshold, what matures first is usually not the service but the scams built around it.** The direct implication for customer operators is to ask first how a channel's effect is independently verified; a channel that cannot answer that question should not enter the budget, however complete its story.

10.12 The Liability of AI-Facing: "The Chatbot Did It" Is No Longer a Defense

The earlier part of Chapter 10 dealt with the consequences of AI saying the wrong thing. This round of legal developments nails the conclusion down: **a company is responsible for what its own AI says, even when the content is fabricated.**

Two rulings reach the same conclusion. First (Canada, 2024): a traveler asked an airline's website chatbot about bereavement fares, and the chatbot said they could be claimed retroactively when the actual policy required applying before travel; the airline refused a refund and argued that "the chatbot is a separate legal entity responsible for its own actions," which the British Columbia Civil Resolution Tribunal (Moffatt v. Air Canada, 2024 BCCRT 149) rejected outright, finding negligent misrepresentation. Second (Germany, 2026): a cosmetic-surgery clinic was sued after its AI chatbot described the practitioner with a nonexistent professional medical title, and the Oberlandesgericht Hamm, on May 12, 2026 (Az. 4 UKI 3/25), ruled that the chatbot is not an independent third party but a tool of the operator, whose statements are directly attributable to the company — even if the AI hallucinated or mishandled correct input — on the grounds that the operator defined the bot's scope, could reprogram it afterward, and had sufficient control over the system^[10-39][third-party verified]. The individual amounts are small, but the principle is large: an automated agent's statements are the company's statements.

The German ruling has a harder compliance backdrop. ePrivacy's practical analysis notes that the court treated the erroneous chatbot response as the company's own misleading commercial practice under Sections 5(1) and 5(2), No. 3 of the German Unfair Competition Act (UWG); the overlay of regulatory change is that Article 50(1) of the EU AI Act took effect on August 2, 2026, requiring transparent disclosure that a user "is interacting with AI"; and at the e-commerce and consumer-law level, false statements about goods, prices, or delivery can trigger not only cease-

and-desist letters but contractual obligations, warranty claims, damages, or fines^[10-40] [third-party verified]. For merchants deploying AI support or sales agents, **"the AI got it wrong" is not just a brand risk but a direct compliance risk subject to enforcement and litigation.**

The official Chinese position has tightened in step. A China Consumers Association (CCA) advisory of September 2, reported by Science and Technology Daily via Xinhua, reminds consumers to treat general generative AI's "informational and referential" nature rationally: for important consumption information such as prices and fees, contract terms, after-sales service, and legal liability, decisions should not rely on AI-generated content alone. The CCA splits AI-related consumer problems into two categories: first, general generative AI should be clearly framed as referential to avoid misleading decisions — in practice there have been unsupported blanket promises, overconfident judgments about litigation outcomes, wrong jurisdiction guidance, and even fabricated nonexistent payments or case filings; second, a company deploying AI customer service cannot weaken its responsibility through technical intervention — there have been price statements inconsistent with reality, inaccurate interpretations of discount rules, insufficient notice of important restrictions, and some operators evading responsibility by citing "AI auto-generation" or "system auto-reply." It urges operators to establish a responsibility-acceptance mechanism for AI-support responses on price, fees, and after-sales matters^[10-41][official claims].

The CCA's complaint-case compilation gives concrete detail. Case one (June 28, 2026): Mr. Zhou mis-bought a different-region version of an item on a game platform and, having not claimed or used it, applied for a refund, but "received only AI customer-service replies" throughout, was repeatedly transferred between channels such as mini-programs and service accounts, and never reached a human, with the issue unresolved. Case three (May 5, 2026): Mr. Gong, choosing a car in a used-car app, was told by the platform's official AI support "Xiaogua" before paying a 2,000-yuan deposit that the 28,700-yuan price already included transfer and registration fees, and paid on that "official promise"; during the transaction the salesperson demanded an additional 1,000 yuan for transfer fees, leaving the consumer feeling that the platform's AI support and its staff gave contradictory information and that he had decided on the basis of wrong information^[10-42][official claims]. **"Receiving only AI replies and being unable to reach a human" is itself a service defect, and an AI-support response is binding.**

The harm on the user side is equally concrete. A report by Legal Daily via Xinhua gives a case: on May 18, 2026, a consumer in Zhenjiang, Jiangsu, used an AI app to book a restaurant and received the affirmative reply "reservation successful, enjoy your meal," with details of arrival time and party size and the note that "showing this AI app's reservation at the door is enough to be seated"; at the restaurant the consumer was told the venue did not support booking through that app and no reservation had actually been made, and the app's support admitted the platform could not pass information to third parties and had no real booking function. Experts advise users to "check the label, protect data, verify sources, and follow the law," and platforms to flag the possibility of "AI hallucination" in user agreements, label AI-generated content, and design easy complaint and correction mechanisms^[10-43][third-party verified].

10.13 The Shutdown of Anthropomorphic Agents and the Mandate to Label

Chinese platforms have drawn a line for "anthropomorphic agents." Sina Finance, citing Southern Metropolis Daily, reports that three leading Chinese AI apps collectively took their "agent" features offline just before a new regulatory rule took effect: Doubao issued a "Doubao Agent Feature Shutdown Notice," saying that due to product adjustments the agent feature would go offline on July 15, 2026; Qwen sent a matching notice that after July 15 users could no longer access agent configurations or conversation history; and Tencent's Yuanbao had already taken the feature offline on June 30. The shutdown dates coincide with the effective date (July 15, 2026) of the Interim Measures for the Administration of Anthropomorphic Interactive AI Services — jointly issued by the Cyberspace Administration of China and four other departments on April 10, applicable to services provided to the domestic public that "simulate natural-person personality traits, thinking patterns, and communication styles in sustained emotional interaction"^[10-44][third-party verified]. For merchants doing "agent-based acquisition," this means the compliance boundary for anthropomorphic agents on Chinese platforms has been redrawn, and third-party agents may be required to migrate into more regulated apps.

The labeling mandate has already landed at the platform level as graded penalties. Drawing on the Measures for Labeling AI-Generated Synthetic Content (effective September 1, 2025) and the three platforms' community rules, ByeRisk compares Douyin, Xiaohongshu, and WeChat Channels: for explicit labels, Douyin uses

an "AI-generated" corner tag at the lower left plus a watermark in the first three seconds, Xiaohongshu a horizontal banner atop the note, and WeChat Channels a corner tag at the upper right plus a pinned comment; implicit labeling requires reading EXIF/C2PA and other metadata and running content-fingerprint or similarity detection. Penalties are tiered: L1 (missing label: auto-labeling plus a private message), L2 (repeated misses: seven to fourteen days of downranking plus frozen traffic packages), L3 (deliberate concealment: removal plus a ban of at least seven days plus loss of commercial rights), and L4 (illegality: permanent ban plus referral to police). The article records that on August 7, 2026, Xiaohongshu upgraded its AI Governance Rules Announcement, becoming the first of the three platforms to extend the labeling obligation from individual posts to the account level (accounts operating AI virtual-human personas must enable the "AI virtual human" profession label on their profile) ^[10-45][third-party verified]. **When "disclosing AI" is upgraded from advice to an admission and distribution condition, the compliance cost of acquisition content is formally entered into the media model** — the "channel falsification" of Section 3.14 also holds on the compliance dimension.

10.14 Summary

The boundaries of AI-powered customer lifecycle management condense into five points:

1. **Backlash:** AI outbound is killing cold email — channel dividends are one-time, and restraint is a strategy;
2. **Hallucination:** AI's mistakes are the company's mistakes — Air Canada paid, and the Hangzhou case was exempted (with disclosure); guardrails determine liability;
3. **Compliance:** regulators have caught up with AI selling — FTC/FCC/SEC/EU AI Act/China's labeling measures, and AI washing is illegal;
4. **The human-machine boundary:** AI impersonating real people is a global red line — trust is the foundation of customer operations;
5. **Liability:** companies bear responsibility, but AI liability insurance has appeared — a risk being insurable shows the industry is maturing.

The final chapter looks to the future: when customers themselves become AI, what will customer operations become?

Chapter 11 The Future: When the Customer Is an AI

In February 2026, CMSWire published a long essay titled "The Machine Customer Economy Has Begun." It offered a prediction that makes many operators uncomfortable: your next batch of customers may not be human.

"Machine customers" are smart devices and AI assistants that carry out service activities on behalf of humans — filing repair orders, renewing supplies, and comparison shopping. Gartner's prediction is more specific: by 2026, 20% of inbound customer-service contacts will come from machine customers^[11-22] [third-party verified].

This chapter is about the future. But one clarification first: the future is not distant — it is already happening.

11.1 Agentic Commerce: AI Agents Shop for People

The definition of agentic commerce (agent-driven e-commerce) is simple: AI agents complete shopping on behalf of consumers — the AI anticipates needs, searches across and compares prices, negotiates, and places orders on their behalf^[11-06] [third-party verified].

Scale predictions (cross-validated across institutions): - McKinsey: agentic commerce could generate \$1 trillion in U.S. retail revenue by 2030^[11-06] [third-party verified]; - Morgan Stanley: by 2030, nearly half of U.S. consumers will use AI agents, adding up to \$115 billion in new spending to U.S. e-commerce^[4-11] [third-party verified]; - Gartner: by 2028, AI agents will intermediate more than \$15 trillion in B2B spending — 90% of B2B purchases will be handled by agents^[11-13] [third-party verified]; - Deloitte: transaction traffic from AI assistants to U.S. retail sites is surging, and agentic commerce is redefining retail economics^[11-05] [third-party verified].

Signs of real-world adoption have already appeared: in October 2025, Walmart was the first to partner with OpenAI, allowing users to complete purchases directly inside ChatGPT through Instant Checkout; a month later, Target followed with its own

AI shopping experience^[11-16] [third-party verified]. Salesforce's 2025 holiday-season data shows: AI/agents accounted for a 20% share of retail, driving \$262 billion; brands that deployed agents grew 59%^[11-17] [official claims].

The shift in shopping patterns deserves one more emphasis: customers no longer "browse"; they "brief" — issuing a requirement instruction such as "I want a waterproof jacket for under €200," and the AI agent interprets, compares, and recommends across retailers^[11-25] [third-party verified]. The web page is no longer the interface; an AI decision layer has become the new intermediary between customers and every retailer.

11.2 When the Buyer Is an AI, What Do You Sell?

This is a double shock for operators. The good news: AI agents are tireless 24/7 buyers that close deals efficiently. The bad news: **AI agents recognize only data, not brands** — when the buyer is an AI, human emotional assets like "brand affinity" are of no use.

So why would an AI agent choose you? ML6's answer, in four characters' worth of meaning: **machine-readable**. Product data (how structured it is, inventory truthfulness, price transparency, GTIN identifiers), reliable commerce signals (real-time inventory/pricing), and interoperable protocols and identity layers — "product data is upgraded from a hygiene task to a strategic asset"^[11-25] [third-party verified]. BCG's reminder is just as critical: a retailer's most valuable customers may no longer be human — AI agents (embedded in Perplexity, ChatGPT, and Google Gemini) are taking over the discover, compare, decide, and order stages^[11-14] [third-party verified].

An action checklist for operators: 1. **Data is the new storefront** — your product pages are for people to look at, but what AI agents read is your structured data feed; 2. **Trust signals must be machine-verifiable** — humans look at "brand reputation"; AI looks at "verifiable review data, warranty terms, and return policies"; 3. **Prepare for agent-to-agent transactions** — frameworks for machine-to-machine negotiation and data verification in B2B are taking shape^[11-13] [third-party verified].

11.3 GEO: Being Cited by AI Is the New Channel

When customers — human or AI — get answers through ChatGPT, Perplexity, or Gemini, whether your brand gets cited becomes a matter of life and death.

As early as 2024, Gartner predicted that traditional search-engine traffic would fall 25% by 2026^[11-12] [third-party verified]. The actual citation data cuts even deeper: original research by Boring Marketing (built on 14,750 real buyer questions tested) found that 53% of brands are completely invisible across the four mainstream AI platforms (ChatGPT/Perplexity/Claude/Gemini), with an overall citation rate of just 14%^[11-03] [third-party verified].

Being cited by AI is not mysticism; it is observable and optimizable (Chapter 3 covered the case of Glasp doing AEO by reading its own server logs). Thirty-five percent of U.S. consumers already use AI at the product-discovery stage^[11-23] [third-party verified] — GEO (generative engine optimization) is becoming a channel as important as SEO.

11.4 Trust: The Stronger the AI, the More Valuable the Human Touch

A chapter about the future should not be only about technology; it must also be about people — because the stronger the technology, the higher the premium on the "human touch."

The data paints two pictures that seem contradictory but are actually consistent:

Consumers trust AI to help them buy better. BCG's global survey (9,000+ respondents across 9 countries): more than 60% of consumers highly trust GenAI's results^[11-08] [third-party verified]. Edelman's Trust Barometer shows brand trust surging in recent years — 80% of people trust the brands they use^[11-01] [third-party verified].

But in customer service, consumers would rather talk to a real person. Metrigy's research: 85% of consumers prefer interacting with a human rather than an AI agent in customer service^[11-15] [third-party verified]. A Medallia report predicts that nearly half of consumers believe that ten years from now they will no longer deal with any company directly — AI agents will do it for them — and that "human service" itself will become premium-priced: more than two-thirds of consumers believe human service will become an optional premium offering^[11-24] [third-party verified].

Taken together, the two pictures point to a clear future: **routine transactions go to AI (consumers welcome this); at critical moments, they want a human (consumers vote with their feet).** Human service will not disappear; it will become

a "luxury good" — and once AI handles 95% of standardized interactions, the remaining 5% of human interactions are the brand's entire battleground for differentiation. The boundary described in Chapter 2 will be drawn even wider in the future: AI handles scale, humans handle trust, and the premium on trust keeps rising.

11.5 The Machine-Customer Ledger: Early Numbers

The previous sections covered trends. This section pulls "machine customers" from narrative back to numbers.

One vendor compilation gathers the estimates of various institutions on agentic commerce: Bain forecasts the U.S. market at \$300–500 billion by 2030, 15%–25% of e-commerce; Gartner predicts 20% of digital commerce transactions will go through AI platforms or agents by 2030; Morgan Stanley forecasts 10%–20% of U.S. e-commerce sales driven by agents; McKinsey estimates it can create about \$1 trillion in orchestrated U.S. retail revenue and \$3–5 trillion globally; J.P. Morgan expects up to 25% of U.S. online sales to flow via agents by 2030, concentrated in recurring, low-risk categories such as groceries and subscriptions. The consumer side is moving too: by 2030 nearly 50% of online shoppers are expected to use AI agents and contribute about 25% of spending, adding roughly \$115 billion to U.S. e-commerce; Kearney estimates 60% of consumers expect to use AI agents in the next 12 months and 73% are familiar with AI tools; 63% of European consumers already use AI for shopping^[11-26] [official claims]. These numbers come from different institutions with different methodologies, but the direction is the same — **"agents as the shopping entry point" is no longer a forecast but a ledger being cashed out.**

Merchant adoption is running ahead too: a payment provider's report says 42% of U.K. and U.S. merchants have already moved first on "AI agent commerce," with AI entry points reshaping the discover-compare-order chain^[11-27] [official claims].

For brands, the real strategic question is not "whether to enter this entry point" but "once I'm in, am I still the protagonist." IBM's Institute for Business Value sums it up in one word — **own**: brands should "own" the agentic commerce experience rather than being intermediated by AI entry points. When purchase decisions are made by consumers' AI agents, brands' control over discovery, comparison, and purchase weakens, and they must win back the "right to be recommended" through data and

experience^[11-28] [third-party verified]. This echoes the conclusion in 11.2: when the buyer is an AI, what you sell is a set of data and experiences that machines can read, verify, and trust.

Taken together with the earlier sections, the machine-customer ledger carries two numbers: an incremental market measured in hundreds of billions of dollars, and an existing risk of being "intermediated." The operator's job is to keep ownership of the customer relationship between those two numbers.

11.6 Protocols and Admission Tickets: The Barrier to Entry in the Machine-Customer Era

Section 11.1 described the shape of agentic commerce. By 2026 it has shifted from "who can build it" to "who can be plugged in" — protocols, policies, and fee schedules now form the admission ticket to the machine-customer era.

Is the protocol war over? In April 2026, Amazon, Meta, Microsoft, Salesforce, and Stripe joined the Universal Commerce Protocol technical committee, forming a ten-seat governance body with Google, Shopify, Etsy, Target, and Wayfair. UCP covers the entire shopping chain from discovery and purchase to after-sales support and is endorsed by more than 20 ecosystem players. But the author's judgment is that "the protocol war is over; the commercial war is not": UCP's advantage is not that its checkout API is pleasant, but that it can accommodate the complexity of real commerce — an agent must be able to ask on the user's behalf whether these three items can be bought with this coupon, shipped to this address, paid with this wallet, and stacked with loyalty benefits, while respecting the merchant's tax, risk, fulfillment, and return rules. The counter-evidence is equally clear: Amazon has not opened its marketplace to external agents, Meta has not handed Instagram commerce to a neutral protocol, and OpenAI shifted ChatGPT shopping in March 2026 from its first version of instant checkout to "product discovery plus merchant-controlled checkout"^[11-29] [third-party verified].

Protocols imply a budget reshuffle. A comparative guide notes that ACP (OpenAI with Stripe, launched September 2025, living inside ChatGPT) and UCP (Google-led, covering AI Mode and Gemini) serve different purposes rather than competing — ChatGPT excels at conversational product discovery while Google AI Mode captures high-intent search demand, so most brands will need to support both. The real implementation judgment is that agentic commerce is **infrastructure-first**,

not content-first: AI shopping agents do not read editorial content, they read structured data, evaluate feeds, and call APIs, so budget should shift from top-of-funnel content to bottom-of-funnel infrastructure — product data completeness, schema markup, feed hygiene, checkout API reliability, and server-side tracking. The same piece concedes that early agentic commerce "sells but cannot be measured": existing analytics stacks were designed for clicks and sessions and cannot fully attribute orders brought by AI agents^[11-30] [third-party verified].

Platforms are already writing rules for agents. Amazon updated its Business Solutions Agreement effective March 4, 2026, adding a dedicated agent policy for the first time and imposing new requirements on AI use and automated systems^[11-31] [official claims]. The detailed requirements include: defining an "AI seller agent" as a system that makes or executes decisions on a seller account without real-time human input for each action; requiring all automation to run through registered SP-API applications, with browser automation and Seller Central scraping explicitly prohibited; mandating a 12-month audit trail for automated actions and self-identification as an automated system; and requiring human authorization for high-impact actions (for example, price changes above 20% within 24 hours)^[11-32] [third-party verified]. The strategic implication matters more than the technical detail: generative search and shopping agents often cannot see a marketplace's on-site listings, so brands must build complete product content on their own channels and treat marketplaces as conversion venues rather than discovery venues — otherwise they go silent at the AI discovery stage.

In the Chinese market the money is re-routing too. Effective August 10, 2026, Douyin's local-services arm applies a separate fee schedule to hotel orders completed through the Doubao entry point: an 11.4% software service fee plus a 0.6% payment processing fee, about 12% combined. Doubao's response is that its local-services business currently runs no paid promotion and merchants cannot influence recommendations or ranking by paying; only a channel service fee is charged after an order is completed^[11-33] [third-party verified]. Simultaneous data shows the scale of the entry point: as of June 2026, AI-native apps reached 499 million monthly active users, up 85.4% year over year, with Doubao first at 382 million. The industry reading is that the AI assistant's role in local-services transaction chains has formally shifted from a platform's "organic traffic entry point" to an **independent transaction channel** — though binding AI recommendations to commercial interests faces a long-term "trust gap."

Payment infrastructure is moving faster than protocols. Alipay reported that AI agent payments have passed 300 million transactions and support 95% of general-purpose agent frameworks, and launched Token Pay and an AI wallet to complete a full-stack AI-native payment product matrix alongside its earlier "AI Pay" and "AI Collect"; on the protocol side it upgraded China's open trust framework for agentic commerce and was first to build A2A (agent-to-agent) and A2M (agent-to-machine) payment capability, while on the merchant side it released its first merchant operations agent^[11-34] [official claims]. The cross-device layer has been filled in as well: Alipay released China's first multi-agent cross-device interconnection protocol system covering three layers — how to talk, how to connect, and how to act — prompted precisely by the first generation of "Doubao phones" being collectively blocked by major apps for using simulated taps, cross-app jumps, and screen reading to bypass restrictions. In one partnership between Alipay and a phone maker, the latter's AI assistant connected nearly 200 Alipay services, adding more than 2 million to Alipay's activity on that brand's handsets^[11-35] [third-party verified]. On the cross-border side, Ant International's Agentic Mobile Protocol onboarded 10 digital wallets and 7 acquiring institutions covering 1.5 billion users at launch and coordinated with two major card networks on a Know-Your-Agent (KYA) trust framework, aiming to simplify agent onboarding and identification across networks while each party retains its own security verification and decision processes^[11-36] [official claims].

Taken together, 2026 offers three admission tickets — which protocol you connect to, which agent policy you comply with, and which fee and identity framework you accept. None of them belongs to the marketing department, yet each directly determines whether a brand is visible to machine customers.

The card networks pushed "machine payments" from consumers to enterprises. In June 2026 Mastercard launched Agent Pay for Machines, positioned to let enterprises support secure, continuous machine payments across cards, accounts, and currencies, with the served parties extending from AI agents to devices and software programs — the second step, after the consumer-facing Agent Pay of April 2025, in which the payment network pushes agent payments toward enterprise-side machine transactions^[11-41][official claims]. Visa started on the same day with a different direction: Intelligent Commerce positions itself as the trust layer for agentic shopping, sharing verifiable credentials among issuers, merchants, and AI platforms about whether an agent is user-authorized, what it may buy, and up to what limit; in April 2026 it added a protocol-agnostic merchant connection product designed to

accept agent orders from multiple protocols (Trusted Agent Protocol and others) at once^[11-42][official claims]. Read together, the card networks are not competing over whose agent is better but over **whose authorization and verification standard gets adopted by everyone else** — a standards contest is a channel contest. For merchants, the real preparatory action is connecting a verifiable "agent identity and authorization" chain rather than adapting platform by platform.

11.7 Machine Customers' Money and Trust

The money is real. Amazon's CEO told an earnings call that its AI shopping assistant is on pace to drive more than \$10 billion in incremental annual sales, and that customers who use it are 60% more likely to complete a purchase than those who do not. The measurement basis is worth noting: the \$10 billion is measured by Amazon's internal "downstream impact" metric using a seven-day rolling attribution model to capture delayed conversions — the assistant's value is counted as the indirect sales it triggers, not direct clicks^[11-37] [third-party verified]. That also implies something: once the assistant itself becomes an acquisition and advertising vehicle, the measurement basis has to be redesigned, or this value stays invisible in the reporting.

Both volume and quality of the channel are changing. Adobe Analytics data shows AI referral traffic to U.S. retail sites in July 2026 was up 62% year over year and up 1,219% cumulatively compared with October 2024, when generative platforms began to spread; more important is quality — in July, AI-referred visitors converted 60% better than non-AI traffic (54% in May)^[11-38] [third-party verified]. A retailer-side data compilation offers another basis: AI referral traffic to U.S. retail sites grew 393% year over year in the first quarter of 2026, with AI-referred visitors spending 48% more time and viewing 13% more pages; on the consumer side, 42% used an AI shopping tool in the past month, 5% have let a fully autonomous AI agent place an order, and 85% say AI improved their shopping experience — while 42% of shopping-related AI conversations are "selection and inspiration" queries rather than direct orders^[11-39] [third-party verified]. High intent matters more than high volume — but that requires brands to first fix landing-page validation, product merchandising logic, and cross-channel attribution.

Trust is the only speed limit. A brand-side guide cites survey figures that must be read together: consumer adoption of agentic shopping is projected to rise from 19% to 46% by the end of 2026, but only 10% of consumers are willing to let an agent

operate fully independently — the vast majority still want human confirmation in the purchase decision — while 83% of consumers have concerns about privacy, data misuse, and unauthorized marketing^[11-40] [official claims]. This matches the judgment in Section 11.4: the expansion rate of machine customers is set not by model capability but by how clearly authorization and refunds are defined. Brands face two risk surfaces — visibility risk (if agents do not recommend you, you do not make the shortlist) and trust risk (unclear control mechanisms directly suppress adoption).

Demand-side conservatism is moving far more slowly than the infrastructure. Put several independent data sets side by side and the gap is plain. Gartner's consumer research concludes that only 11% of U.S. consumers are willing to let AI make purchase decisions, and advises marketers to design AI shopping tools primarily as research support rather than as the decision-maker^[11-43][third-party verified]. Payment company Checkout.com's research fills in the middle ground: 33% of consumers expect at least 10% of their purchases to be AI-driven within a year, and 72% of merchants believe consumers will adopt agent-led shopping faster than most merchants are prepared for; on the trust side, three figures form the drag — 27% of consumers trust no institution to run an AI shopping agent, 24% say they will never delegate a purchase to AI, and 57% are willing to let an AI agent switch brands for a better price^[11-44][third-party verified]. The payment networks' own forecasts are the most expansive: a Mastercard report predicts that by 2030 more than one in ten online shoppers will routinely use an AI agent to choose and pay for them^[11-45][official claims] — a claim whose provenance should be labeled clearly, since it is a payment network's self-interested forecast: the pipes only carry volume once they are built.

Read together, the four data sets point to one conclusion: infrastructure and expectations are running ahead of consumer authorization. That gap in between is precisely what brands and platforms must fill together (authorization, revocability, auditability).

On the B2B side, the migration path looks more like agent-to-agent. Deloitte judges that the future form of B2B commerce is multi-agent systems researching, negotiating, purchasing, and selling on behalf of enterprises: buyers spend nearly 30% more budget with suppliers that offer an excellent buying experience, and the perception gap between buyers and sellers about the sales process is widening — a gap that, compounded by the maturation of agentic AI, creates a window for sellers to get ahead; the article also projects that if agents are orchestrated strategically and risks appropriately mitigated, the market value of AI

agents could rise to \$45 billion^[11-46][third-party verified]. **Sellers must be legible to two kinds of readers at once** — with a website and sales process a human can understand, and with specifications, prices, terms, and inventory data an agent can read and compare directly.

A magnitude that agencies can plan around has appeared. Several agencies estimate that Perplexity's shopping business has reached an annualized GMV run rate approaching \$2 billion; Perplexity has not confirmed the figure, so it should be read as a directional signal rather than an audited result. The growth comes from a change in answer format: when users ask shopping questions, the answer carries product cards alongside the usual sourced response, and Perplexity says these cards are relevance-ranked non-sponsored recommendations rather than paid slots. Transactions are carried by the subscription tier, and a free merchant program feeds retailers' (including Shopify stores') real-time product data into its index^[11-47][third-party verified]. The significance lies in entry rules that invert search advertising: **free onboarding, ranked by relevance rather than bid** — which means a brand's "spend" here looks more like a contest of product data and content than a contest of budget.

11.8 The Layering of the Protocol Stack and the 2026 Retreat

Section 11.6 was about admission tickets. This round's details are more interesting: **the protocols are being layered, while 2025's product claims are already out of date.**

Three standards are not on the same layer; don't treat them as a three-way choice. A merchant guide lays out the 2026 agentic-commerce standards clearly: UCP (Google with Shopify, Etsy, Walmart, Target, and others) covers the full journey from discovery and cart to checkout and post-purchase, with merchants publishing capabilities through a `/well-known/ucp` JSON manifest and supporting REST, A2A, and MCP transports; ACP (co-built by OpenAI and Stripe, open-sourced September 29, 2025) handles checkout execution, its first landing being Instant Checkout in ChatGPT, with Etsy the first merchant and more than a million Shopify merchants later connected; and AP2, now governed by the FIDO Alliance, handles payment authorization. The article notes that ACP "outlived its own product" — OpenAI scaled back Instant Checkout in March 2026, and the protocol retreated from a checkout product to a "pipe"^[11-48][third-party verified].

A cross-border payments firm's explainer pinpoints ACP's key change: in March 2026 OpenAI scaled back native instant checkout and routed purchases back into merchant apps inside ChatGPT, "to some degree returning checkout control to merchants"; UCP, by contrast, takes an "open orchestration" route, with product retrieval queried by different agents across multiple platforms, authorization information and intent data transmitted with the request for the merchant's own risk checks, and the transaction still completed through the merchant's existing checkout page or payment service provider — the protocol serving as "a universal interaction language rather than the final landing vehicle"^[11-49][third-party verified]. An architectural view offers a cleaner contrast: UCP is decentralized (merchants publish capabilities any agent can discover and transact with), ACP is platform-mediated (merchants submit data to OpenAI, which ChatGPT presents), and in practice merchants will likely implement both to maximize distribution, "just as retailers list on both Google Shopping and Amazon"; MCP is not a commerce protocol but a general tool-integration layer, relevant to commerce because both UCP and ACP explicitly support MCP as a transport binding^[11-50][third-party verified].

The product-level correction matters more: don't treat 2025's announcements as the status quo. An explainer around OpenAI's March 2026 update says plainly that the first version of Instant Checkout did not offer the flexibility it wanted, so the direction shifted to product discovery and letting merchants use their own checkout experience, with Instant Checkout possibly available only for "eligible products and merchants"; any unverifiable order, payment, or delivery question should go to the responsible team^[11-51][third-party verified]. A product-feed vendor's roundup adds the operational detail: Google's legacy Content API for Shopping is set to shut down on August 18, 2026, so programmatic access must migrate to the Merchant API; in May 2026 Google added six optional "conversational attributes" (question_and_answer, related_product, variant_option, item_group_title, popularity_rank, document_link), not required but giving AI shopping more structured context^[11-52][third-party verified]. **The same Merchant Center feed is also read by Microsoft Copilot — the structure of the feed determines whether AI can understand you.**

The platforms' official positions give the channel differences: Shopify's blog explains that Perplexity shopping is driven by a free merchant program using a Google Shopping-format feed; Perplexity completes in-conversation purchases via PayPal; Google AI Mode supports Shopify-powered in-conversation checkout, with eligible,

opted-in orders flowing into the Shopify back end like ordinary orders; and ChatGPT presents products through Shopify Catalog and directs buyers via an in-app browser to the merchant's own store. Shopify merchants can be discovered in ChatGPT automatically through Shopify Catalog, with no extra setup^[11-53][official claims].

11.9 How Much Room Consumers Give AI Shopping Agents to Make a Mistake

Section 11.7 said trust is the only rate limiter. Consumer surveys in the first half of 2026 measured that limiter's value: **the tolerance is extremely low.**

A survey commissioned by payments-technology firm ACI Worldwide and run by YouGov among 2,080 UK adults (June 19–22, 2026) delivers the headline: **six in ten UK consumers say that if an AI shopping agent makes one mistake, they will stop using it**; consumers have clear concerns about "control, money, and who is responsible" — that is, about an agent spending on their behalf and who answers when it errs^[11-54][third-party verified]. For merchants preparing to hand the closing step to an agent, this is a "near-zero error tolerance" demand-side signal.

Klaviyo's reading of its 2026 AI Consumer Trends Report gives tiered trust: only 13% of consumers "fully trust" AI, 36% "somewhat trust" it, and 30% are neutral; on shopping recommendations specifically, 85% have at least some trust that AI provides accurate, personalized recommendations, and 39% have bought an AI-recommended product in the past six months; but only 54% trust AI to provide good customer service through a conversational agent — **trust in the "recommendation" scenario is clearly higher than in the "service" scenario.** On the brand side, 32% of consumers would trust a brand less for using AI-generated content, and nearly one in five sees low-quality or formulaic AI content from brands weekly. Klaviyo's advice is not to remove human review and to route complex support requests back to people^[11-55][third-party verified].

Alchemer's 2026 Retail Report (1,002 U.S. shoppers) is titled "retail AI adoption outpaces consumer trust": 48.5% of shoppers have used an AI tool to research a purchase in the past year; by age, 74% of 18–29-year-olds have done so versus only 20% of those 61 and older, and the youngest group is nine times as likely as the oldest to "discover their most recent purchase through an AI tool"^[11-56][third-party verified]. **Read alongside ACI's six-in-ten error tolerance, this gives the real**

constraint on the expansion of machine customers: it is not that the model is too weak, but that users give the agent a trial-and-error budget of exactly one.

And autonomy is already tiered inside products. A postmortem on Amazon's shopping agent Rufus lays out three product lines: Q&A, Buy for Me (letting Rufus browse and compare off-Amazon items, with the final purchase still requiring user approval), and Auto Buy (fully autonomous, limited to Prime members and FBA items, with a default six-month monitoring window that buys automatically when the price is hit — Amazon says users save about 20% on average). For sellers, Rufus favors items with strong reviews and clear descriptions, sparse or low-quality listings become "invisible to Rufus," and the AI Q&A section has shifted from an afterthought to a primary data source for AI recommendations^[11-57][third-party verified]. This closes the loop with Section 11.6: **before machine customers, visibility comes from the completeness of data and reviews, not from the media budget.**

11.10 Closing: The Next Decade of Customer Operations

Threading the whole book together, the evolution of AI-powered customer lifecycle management can be summarized in three phases:

Phase 1 (2023-2025): Tool substitution. Companies treat AI as a tool that replaces repetitive work — AI support, AI outbound, AI analytics. The winners in this phase are the most efficient companies; the losers are those that optimized customer service as a pure cost center (the Klarna lesson).

Phase 2 (2025-2027): Process redesign. Companies rebuild their customer-operations processes around AI — AI-native CRM, CSQL, outcome-based pricing. The winners in this phase are the companies with the cleanest data, because AI's ceiling is set by data quality.

Phase 3 (2027-): Ecosystem change. Customers themselves begin to turn into AI — machine customers, agent-to-agent transactions. The winners in this phase are the companies that make their businesses something AI agents can "read."

The three phases are not replacements for one another; they stack. The reality of 2026: most companies are still struggling in Phase 1 (74% stuck in pilots — see Chapter 8), a few leaders have already run into Phase 2, and the wave of Phase 3 has already reached the shore.

What this book can do is lay out the map across the three phases: the engine (Chapters 3-6), the foundation (Chapter 7), the organization (Chapter 8), the industry (Chapter 9), and the boundaries (Chapter 10). The rest of the road, operators must walk themselves.

Finally, back to Klarna from the preface. It went from "AI does the work of 700 support agents" to "we went too far" and then back to hiring people again — a complete cycle in two years. The value of this cycle lies not in its conclusion but in its process: it proved the enormous capacity of AI-powered customer lifecycle management, and it proved that customers are people, not tickets. **The entire art of AI-powered customer lifecycle management is to use the scale of AI to serve every individual person well.**

Appendix A: Case Index and Sources

This appendix lists the sources behind every figure and case in this book. Numbering scheme: [chapter-sequence]; in the main text, a <sup> superscript link after each key figure points to the corresponding entry here.

Source conventions: the material in this book comes from a continuously running collection pipeline. Each entry is tagged 【official claims】 (figures self-reported by companies/institutions) or 【third-party verified】 (independent media, research institutions, or academic sources) — click the link to check the original text. Entries not cited in the main text are further reading on the same topic.

This part covers the sources for Chapters 1-2 through Chapter 7.

Sources for Chapters 1-2

- [1-01] Inc: 'No AI' Is the New 'All Natural'—and Brands Are Winning Big With It: <https://www.inc.com/emily-cody/no-ai-is-the-new-all-natural-and-brands-are-winning-big-with-it/91375451>
- [1-02] Elevate GTM Research: State of AI-Native GTM 2026 Report: Agentic AI Is Rewriting How Revenue Organizations Operate: <https://www.elevategtmsolutions.com/resources/gtm-research/state-of-ai-native-gtm-2026>
- [1-03] Bloomreach: 260 Sample Sale: An AI Campaign Agent Lifts Conversion 2.4x and Drives \$580K+ in AI-Attributed Revenue Across Four Scenarios: <https://www.bloomreach.com/en/case-studies/260-sample-sale>
- [1-04] Copy.ai: AI for Customer Lifecycle Management (CLM): A Full-Funnel Framework from Reactive Responses to Proactive Intelligence (2026-04/05): <https://www.copy.ai/blog/ai-for-customer-lifecycle-management>
- [1-05] Nutshell: AI-Native CRM: Hype vs. Real Business Results (2026-07): <https://www.nutshell.com/blog/ai-native-crm>

- [1-06] Blueshift: The Complete B2C Customer Lifecycle Marketing Guide: The Economics of Acquisition → Onboarding → Engagement → Retention → Win-Back (2025-10): <https://blueshift.com/blog/customer-lifecycle-marketing-the-complete-guide-for-b2c-marketers/>
- [1-07] Columbia University (Teachers College): Experts Caution Against Using AI Chatbots for Emotional Support: <https://www.tc.columbia.edu/articles/2025/december/experts-caution-against-using-ai-chatbots-for-emotional-support/>
- [1-08] Northwestern Mutual: Human Connection Over Machines: Americans Trust Advisors More Than AI: <https://news.northwesternmutual.com/2025-08-05-Human-Connection-Over-Machines-Americans-Trust-Advisors-More-Than-AI-for-Financial-Advice,-Finds-Northwestern-Mutuals-2025-Planning-Progress-Study>
- [1-09] Intercom: The 2026 Customer Service Transformation Report: <https://www.intercom.com/customer-transformation-report>
- [1-10] McKinsey & Company: When AI meets desire: Innovating human-centered luxury experiences in the agentic age: <https://www.mckinsey.com/industries/retail/our-insights/when-ai-meets-desire-innovating-human-centered-luxury-experiences-in-the-agentic-age>
- [1-11] No Jitter: When Software Takes on the Work: Customers Are Buying "Work," Not Software (2026-09-03): <https://www.nojitter.com/ai-automation/when-software-takes-on-the-work-what-are-customers-buying->
- [1-12] Ahoy: What Is an AI-Native CRM? A Complete Guide (Updated 2026-08): <https://ahoy.ai/ai-native-crm/>
- [1-13] Growth Unhinged: An AI Growth Playbook from Consumer to B2B: The Fyxxer Case: <https://www.growthunhinged.com/p/fyxxer-ai-growth>
- [1-14] EY: How to Rethink SaaS Go-to-Market Strategy for an Agentic AI World: https://www.ey.com/en_us/insights/tech-sector/saas-go-to-market-strategy-for-an-agentic-ai-world
- [1-15] Chief Outsiders: The Augmented Revenue Engine: How AI-Assisted Selling Delivers 3.7x Quota Attainment and +30% Win Rates (2026-08): <https://www.chiefoutsiders.com/blog/the-augmented-revenue-engine>
- [1-16] EY: Serving Customers Who Rely on Tech but Don't Trust Tech: In the AI Age, Customer Relationships Are Built on Trust, Respect, and Value: https://www.ey.com/zh_tw/insights/consumer-products/how-to-serve-customers-who-rely-on-tech-but-dont-trust-tech

- [1-17] ZoomInfo Pipeline: The Complete Guide to Customer Lifecycle Management (CLM): How B2B Revenue Teams Operate Customers Across the Full Lifecycle (2026-07): <https://pipeline.zoominfo.com/sales/customer-lifecycle-management>
- [1-18] IBM (IBM Think): The Future of AI in Customer Service: AI Won't Replace Human Agents — It Will Be Their Real-Time Partner: <https://www.ibm.com/think/insights/customer-service-future>
- [1-19] Bocar Dia: The Emerging AI-Native GTM Playbook: Six Models for Racing Toward a \$100M Company (2025-10): <https://ainativegtm.substack.com/p/the-emerging-ai-native-gtm-playbook>
- [1-20] MiningLamp: Jiang Ping of MiningLamp on AI-Native Business Upgrades: A 1,500-Person "Human-Machine Collaboration" Organizational Experiment (2026-06): <https://www.mininglamp.com/news/7805/>
- [1-21] monday.com: Lifecycle Marketing: How Leading Brands Play It in 2026 and the Role of AI (2025-11): <https://monday.com/blog/monday-campaigns/lifecycle-marketing/>

Sources for Chapter 3

- [3-01] lemlist: 11x AI Review: Worth Buying, or Overhyped Automation? (2026-08): <https://www.lemlist.com/blog/11x-ai-review>
- [3-02] 11x: Comparing Customer Results of AI SDRs: 11x Alice vs. Artisan Ava and Others (2025-2026): <https://www.11x.ai/guides/artisan-vs-11x>
- [3-03] 11x: 11x — The AI Growth Company (Alice the SDR & Julian the Phone Agent): <https://www.11x.ai/>
- [3-04] 11x: 13 Best AI Sales Agent Tools for B2B Teams (2026) (11x's Official List, 2026-06-18): <https://www.11x.ai/guides/ai-sales-agent-tools-b2b-sales-teams>
- [3-05] Cognism: 6sense ABM Platform: Features, Costs & Best Alternative (A Breakdown of 6sense's ABM Platform): <https://www.cognism.com/blog/6sense-abm>
- [3-06] Salesmotion: AI SDRs vs. Human SDRs: A Realistic ROI Comparison for 2026 (2026-06): <https://salesmotion.io/blog/ai-sdrs-vs-human-sdrs>
- [3-07] Autobound: AI SDR Tools Market Guide 2026: Rankings and Key Data for 12 Platforms (2026-02): <https://www.autobound.ai/blog/ai-sdr-tools-guide>

- [3-08] Tencent News: AI-SDR Unicorn 11x Blows Up: Fabricated Customer Endorsements, Product Hallucinations, and High Churn Rates: <https://new.qq.com/rain/a/20250325A03HU700>
- [3-09] The Pipeline Group: The Hidden Dangers of AI SDRs: Why They Should Never Be Used for Outbound B2B Enterprise Sales: <https://www.thepipelinegroup.io/blog/the-hidden-dangers-of-ai-sdrs-why-they-should-never-be-used-for-outbound-b2b-enterprise-sales>
- [3-10] Demandbase: AI in Account-Based Marketing: The Complete Guide for 2026 (A Complete Guide to AI × ABM): <https://www.demandbase.com/blog/ai-in-account-based-marketing/>
- [3-11] The Next Web (TNW): AI Is Turning B2B Outreach into a Tragedy of the Commons: <https://thenextweb.com/news/b2b-outreach-tragedy-of-the-commons>
- [3-12] AiSDR: What Is the ROI of AiSDR? The Vendor's Official Calculation of Return on Investment (AiSDR Official Blog, 2026-07): <https://aisdr.com/blog/what-is-the-roi-of-aisdr/>
- [3-13] 11x: Alice: AI That Autonomously Handles B2B Outbound Lead Generation: <https://www.11x.ai/products/alice/outbound-lead-generation>
- [3-14] Business Insider: Artisan's \$25M Series A: The Funding Story and Global Expansion Blueprint of an AI Sales-Agent Star: <https://www.businessinsider.com/pitch-deck-ai-agent-startup-artisan-series-a-funding-round-2025-4>
- [3-15] SalesRobot: Artisan AI Review 2026: Is Ava Worth \$280/Month? (2026-05): <https://www.salesrobot.co/blogs/artisan-ai-review>
- [3-16] Leadstec: B2B Foreign-Trade Intent-Data Prospecting: From Website Behavior to CDP and Automated Follow-Up (A Guide for Chinese Companies Going Global): <https://www.leadstecgz.com/sc/blogs/b2b-foreign-trade-intent-data-customer-development/>
- [3-17] Coldreach: The Ultimate Claygent Review: What 100+ Users' Experiences Reveal (2026-08): <https://coldreach.ai/blog/claygent-review>
- [3-18] Ekimetrics: MMM — How Marketing Mix Modeling Measures Marketing Impact and Optimizes Budget Allocation (A Comeback in the Privacy-First Era of 2026): <https://www.ekimetrics.com/articles/marketing-mix-modeling-how-to-measure-marketing-impact-and-optimize-budget-allocation>

- [3-19] Semrush: Generative Engine Optimization (GEO): A Practical Guide (Hands-On GEO: Generative-Engine Optimization): <https://www.semrush.com/blog/generative-engine-optimization/>
- [3-20] Dataslayer: Google Ads Performance Max: Complete Guide 2025: <https://www.dataslayer.ai/blog/google-ads-performance-max-complete-guide-2025>
- [3-21] Google (Official Blog): Google Opens Its MMM Model Meridian to Everyone: Bayesian Causal Inference for Budget Allocation (2025-01): <https://blog.google/products/ads-commerce/meridian-marketing-mix-model-open-to-everyone/>
- [3-22] Salt Agency: How Zapier Quadrupled Organic Growth | SEO Case Study (Zapier Quadrupled Organic Traffic in Three Years): <https://salt.agency/blog/how-zapier-quadrupled-organic-traffic/>
- [3-23] Aviso: Ideal Customer Profile (ICP): How AI-Powered ICP Scoring Drives Higher Win Rates: <https://www.aviso.com/blog/ideal-customer-profile-scoring>
- [3-24] Omniscient Digital: Jasper Grows Blog Organic Sessions 810% and Blog-Attributed User Signups 400X: <https://beomniscient.com/case-studies/jasper/>
- [3-25] House of MarTech: MMM vs. Multi-Touch Attribution: A Three-Layer Unified Measurement Framework for the Privacy-First Era of 2026 (2026-04): <https://houseofmartech.com/blog/marketing-mix-modeling-vs-multi-touch-attribution-privacy-first-measurement-framework-for-2026>
- [3-26] TapClicks: Marketing Attribution in 2026: Why Multi-Touch and Marketing Mix Modeling Have to Work Together: <https://www.tapclicks.com/blog/marketing-attribution-in-2026-why-multi-touch-and-marketing-mix-modeling-have-to-work-together/>
- [3-27] Lionelz: Meta Advantage+: The Algorithm at the Service of Your Ecommerce: <https://lionelz.com/en/blog/meta-advantage-plus-algorithm-ecommerce/>
- [3-28] Omnius: Programmatic SEO Case Study: From 67 to 2,100 Monthly Signups in 10 Months: <https://www.omnius.so/blog/programmatic-seo-case-study>
- [3-29] MarketBetter: Regie.ai Review 2026: RegieOne and Force Multiplier Rep Pricing — Does AI SEP Deliver? (2026-08): <https://www.marketbetter.ai/blog/regie-ai-review-2026/>
- [3-30] SalesDriver: Key Takeaways from the 2024 China B2B Enterprise ABM Marketing Whitepaper (The State of ABM Account-Based Marketing in China): <https://m.chwang.com/report/183696650243>

- [3-31] NetEase: Scott Brinker: Will AI in Martech Make the Buyer Experience Better or Worse?: <https://www.163.com/dy/article/JB5DRAGH05310573.html>
- [3-32] Digital Applied: The Case Against AI SDRs: Contrarian Analysis 2026: <https://www.digitalapplied.com/blog/case-against-ai-sdrs-contrarian-analysis-2026>
- [3-33] ZoomInfo Pipeline: What Is an Ideal Customer Profile? An ICP Guide for B2B — From Static Attributes to Live Intent: <https://pipeline.zoominfo.com/marketing/ideal-customer-profile>
- [3-34] 6sense: What Is Intent Data? (An Intent-Data Guide: How AI Identifies Accounts That Are "In Market Right Now"): <https://6sense.com/platform/intent-data/what-is-intent-data/>
- [3-35] TechCrunch: a16z- and Benchmark-backed 11x Has Been Claiming Customers It Doesn't Have: <https://techcrunch.com/2025/03/24/a16z-and-benchmark-backed-11x-has-been-claiming-customers-it-doesnt-have/>
- [3-36] causaLens: Case Study: Causal-AI Marketing Mix Modeling Helps a Major Mobile App Company Cut Its Annual Marketing Budget by 5% and Achieve 15x ROI: <https://causalens.com/causal-ai/customer-case-study-marketing-mix-modeling>
- [3-37] Woshipm: Why AI Marketing Is "Killing" the Soul of SaaS — A Systemic Critique of AI SDRs Losing Their Context: <https://www.woshipm.com/it/6349720.html>
- [3-38] 11x: What Is an AI Phone Agent? A Complete Guide with Real-World Examples — How Julian Autonomously Calls Out, Answers, and Books Meetings (11x Official Blog, 2026-04): <https://www.11x.ai/blog/what-is-an-ai-phone-agent-complete-guide-with-real-world-examples>
- [3-39] Boomerang: Death to Cold Outreach: How Sales Prospecting Moves from "Casting a Wide Net" to Relationship-Driven Engagement: <https://www.getboomerang.ai/post/blog-death-cold-outreach>
- [3-40] Runwise: Innovation Case | SANY: The Internationalization Strategy and Marketing Innovation of a Construction-Machinery Giant (A Sample of Target-Customer Operations and Channel Upgrades at Large Chinese B2B Manufacturers): <https://runwise.co/dtc/content-marketing/198774/>

- [3-41] NetEase Zhiqi: How to Break the "Add-Then-Go-Silent" Deadlock with AI Private-Domain Operations: From Confusion to a Configurable Operating Cadence (NetEase Zhiqi Playbooks and Retail Cases): <https://grow.163.com/cms/ai-ke-fu-ai-si-yu-2026-3RbSGjTy.html>
- [3-42] Harvard Business Review: How Successful Sales Teams Are Embracing Agentic AI (2025-09): <https://hbr.org/2025/09/how-successful-sales-teams-are-embracing-agentic-ai>
- [3-43] Shuyun: Shuyun's Private-Domain AI "Winner" Upgrade: Four New Scenarios Plus Real Cases Bring New Answers for Cost-Cutting and Efficiency in Private-Domain Operations (AI Digital Employees and Human-Machine Collaboration): <https://www.shuyun.com/4499.html>
- [3-44] TechCrunch: Yes, LinkedIn Banned AI-Agent Startup Artisan — But Now It's Back: <https://techcrunch.com/2026/01/07/yes-linkedin-banned-ai-agent-startup-artisan-but-now-its-back/>
- [3-45] Sequoia Capital: Services Are the New Software (2026-03): <https://sequoiacap.com/article/services-the-new-software>
- [3-46] NetEase: Traffic Dividends Have Peaked — How to Do Marketing and Customer Service? AI Foundation Models Become the New Engine of Private-Domain Growth (Tencent Qidian's Intelligent-Operations Matrix + a Laiyifen Case): <https://www.163.com/dy/article/IOTLLO6Q05129QAF.html>
- [3-47] NetEase Zhiqi: "Audience-Pack" Strategies in Private-Domain Operations: Refinement Starts with Segmenting People First (Layered AI Private-Domain Operations): <https://grow.163.com/cms/ai-si-yu-wang-yi-zhi-qi-2026-ZqurME2z.html>
- [3-48] Weiyangx: NetEase Zhiqi Launches an "AI Private-Domain Operations" Solution: Layered Private-Domain Operations in a Human-Machine Collaboration Model: <https://www.weiyangx.com/441234.html>
- [3-49] 36Kr: NetEase Zhiqi's Wang Wei: AI Is Reshaping the Customer-Service Industry — Human-Machine Collaboration Is Becoming the Mainstream Model: <https://36kr.com/p/305482488832001>
- [3-50] Gartner: Gartner Research: The AI Cold-Email Efficiency Trap and the Coming Customer-Trust Crisis: <https://www.gartner.com/en/newsroom/press-releases/2024-03-06-gartner-says-60-percent-of-cmos-will-adopt-ai-to-fight-decline-in-customer-trust>

- [3-51] Sean Ellis: How Glasp Grew ChatGPT Traffic from 500 to 19,000 Daily Conversations: Answer-Engine Optimization by Reading Your Own Server Logs: <https://seanellis.substack.com/p/how-glasgrew-chatgpt-traffic-from>
- [3-52] Amazon Pilots Ad Services in ChatGPT: What Marketers Need to Know (Marketing Dive): <https://www.marketingdive.com/news/amazon-pilots-ad-services-in-chatgpt-what-marketers-need-to-know/829945/>
- [3-53] The Definitive Guide to Adopting Agentic Commerce in 2026 (HUMAN Security): <https://www.humansecurity.com/learn/resources/guide-adopting-agentic-commerce/>
- [3-54] \$220 of Google App Ads, 60% of Installs Were Bots: A Small Indie Developer's Measurement Audit: <https://dayzlegame.com/blog/google-ads-bot-farm/>
- [3-55] Mirror Studio (republished by Futu Niuniu): WeChat Search Goes Fully AI-Native: <https://www.moomoo.com/hans/news/post/58266167>
- [3-56] China Economic Net: Over 100 Million Users Can Summon Qwen at Any Time; Quark Releases Its New AI Browser: http://www.ce.cn/cysc/newmain/yc/jsxw/202511/t20251126_2604231.shtml
- [3-57] TMTPost: Qwen Cannot Afford to Fail, Quark Has No Choice — the Reassignment of Alibaba's Consumer AI Entry Point: <https://www.tmtpost.com/7776036.html>
- [3-58] Douyin E-Commerce Learning Center (official): Compass Search Analytics Gets a Full Upgrade to Help Merchants Master Search Traffic: <https://school.jinritemai.com/doudian/web/article/aHYhVvzvFuG>
- [3-59] Zhihu (author Bobo): A Panoramic Analysis of Douyin Search Trends in 2026: <https://zhuanlan.zhihu.com/p/1994346212103586331>
- [3-60] Conductor: The 2026 AEO / GEO Benchmarks Report: <https://www.conductor.com/academy/aeo-geo-benchmarks-report>
- [3-61] Search Engine Land (citing Peec AI's 30 Million Source Study): AI Search Engines Cite Reddit, YouTube, and LinkedIn Most: <https://searchengineland.com/ai-search-engines-cite-reddit-youtube-and-linkedin-most-study-473138>
- [3-62] 5W (Citation Source Audit): The State of AI Citations 2026: <https://www.5wpr.com/research/state-of-ai-citations-2026>
- [3-63] OtterlyAI: The AI Citations Report 2026 — What 1+ Million Data Points Reveal: <https://otterly.ai/blog/the-ai-citations-report-2026>

- [3-64] Search Engine Land: GEO Experiments Challenge Conventional AI Visibility Advice (775 Citation Events): <https://searchengineland.com/geo-experiments-challenge-conventional-ai-visibility-advice-488342>
- [3-65] ifenxi: 2026 China GEO Market Research Report: <https://xueqiu.com/5243595231/402881486>
- [3-66] Forrester: Seven Reasons to Be Skeptical of Agentic Prospecting: <https://www.forrester.com/blogs/seven-reasons-to-be-skeptical-of-agentic-prospecting/>
- [3-67] Forrester: Agentish vs. Agentic in GTM — Choose Control Over Autonomy: <https://www.forrester.com/blogs/agentish-vs-agentic-in-gtm-choose-control-over-autonomy/>
- [3-68] Salesforce (press release): Siemens and Salesforce Redefine Industrial Sales and Service with Agentforce: <https://www.salesforce.com/news/press-releases/2026/09/15/siemens-agentforce-redefine-industrial-sales-service/>
- [3-69] Fortune: Exclusive — As AI Threatens Search, Profound Raises \$96 Million to Help Brands Stay Visible: <https://fortune.com/2026/02/24/exclusive-as-ai-threatens-search-profound-raises-96-million-to-help-brands-stay-visible/>
- [3-70] arXiv: Optimizing Visibility in Generative Engines — A Critical Survey of Generative Engine Optimization (2023-2026): <https://arxiv.org/abs/2607.14035>
- [3-71] Semrush: Most Cited Domains in AI — 13 Weeks of Weekly Snapshots: <https://www.semrush.com/blog/most-cited-domains-ai/>
- [3-72] Muck Rack / Generative Pulse: What Is AI Reading? AI Citations Across 17 Industries: <https://muckrack.com/blog/what-is-ai-reading-may-2026>
- [3-73] G2: G2 Research — The Answer Economy (1,076 B2B Software Buyers and Decision-Makers Surveyed): <https://company.g2.com/news/g2-research-the-answer-economy>
- [3-74] Forrester: Preference Matters More Than In-Market Intent Alone in Modern B2B Buying: <https://www.forrester.com/blogs/preference-matters-more-than-in-market-intent-alone-in-modern-b2b-buying/>
- [3-75] Search Engine Land: Google AI Mode Cuts Clicks, and Satisfaction Doesn't Rise — A 1,100-Person Field Experiment: <https://searchengineland.com/google-ai-mode-cuts-clicks-satisfaction-study-490494>
- [3-76] Yage Daily Notes: Why Agent Stores Never Found Product-Market Fit — Doubao's 8 Million Agents and Under 1 Million Yuan of Daily Revenue: <https://yage.ai/share/uga-sunset-20260704.html>

- [3-77] Growth Hacking AI Weekly EP.71: GEO Franchising Borders on Fraud — The First Business GEO Nailed Was Collecting Franchise Fees from City Agents: https://mp.weixin.qq.com/s?__biz=MzA4MTc3NzY0Mg==&mid=2652351009&idx=1&sn=b1ef2564fc89ae3083641db5096a2
- [3-78] Forbes: AI Wrote Your Cold Email. Now It Carries an Invisible Watermark: <https://www.forbes.com/sites/terdawn-deboe/2026/08/16/ai-wrote-your-cold-email-now-it-carries-an-invisible-watermark/>
- [3-79] TitanX: Can AI SDR Agents Replace Human Sellers? Lessons from 400+ Hours of Building One: <https://titanx.io/news/can-ai-sdr-agents-replace-human-sellers>
- [3-80] 21st Century Business Herald: After 80,000 Cold Emails Sank Without a Trace, Jintai Technology's Sales Agent Brought Inquiries Equal to the Prior Three Years: <https://www.21jingji.com/article/20260708/herald/7b8cfddafb8d81b0a06cd84014170d10.html>
- [3-81] 21st Century Business Herald (citing QuestMobile): WeChat AI Search Reaches 160 Million Monthly Active Users, Surpassing Native AI Apps Such as Doubao and Yuanbao: <https://www.21jingji.com/article/20250923/herald/a6cdb8dc7c606154d8959d6ac550c50d.html>
- [3-82] AirOps (with Kevin Indig): The 2026 State of AI Search — How Modern Brands Stay Visible: <https://www.airops.com/report/the-2026-state-of-ai-search>
- [3-83] University of St. Gallen (arXiv): Don't Measure Once — Measuring Visibility in AI Search (GEO): <https://arxiv.org/pdf/2604.07585>
- [3-84] 5W Public Relations (PRNewswire): Wikipedia and Reddit Now Drive Over 25% of ChatGPT Citations in the U.S.: <https://www.prnewswire.com/news-releases/wikipedia-and-reddit-now-drive-over-25-of-chatgpt-citations-in-the-us-new-5w-research-finds--wsj-nyt-and-bloomberg-do-not-appear-in-the-top-20-302768339.html>
- [3-85] Profound: Profound Index Report — Summer 2026: <https://www.tryprofound.com/reports-guides/profound-index-report-summer-2026>
- [3-86] Onely: What Influences Brand Visibility in AI Search? A Practical Guide for 2026: <https://www.onely.com/blog/what-influences-brand-visibility-in-ai-search-a-practical-guide-for-2026>

- [3-87] 21st Century Business Herald: The Two Paths by Which WeChat and Alipay Agents "Break the Wall": https://m.21jingji.com/article/20260617/herald/ae1ed1491f5fce9997a168b6523178e7_zaker.html
- [3-88] Leiphone: Most People Missed It — Alipay's and WeChat's Agent Moves Over These Three Months: <https://m.leiphone.com/category/aijuejinzhi/PkorfY0MliENCG54.html>
- [3-89] Science and Technology Daily (China Science and Technology Network): Alipay — Merchants Are Eagerly Connecting Agents to "Abao," and Some Development Requests Are Already Booked Into Next Year: https://www.stdaily.com/web/gdxw/2026-08/17/content_565275.html
- [3-90] Science and Technology Daily (China Science and Technology Network): In the Agent Era, AI Will Become the New Consumption Entry Point: https://www.stdaily.com/web/gdxw/2026-09/27/content_588729.html
- [3-91] Securities Times: The Internet Giants Are Going All In — the Battle for AI Entry Points Officially Begins: <https://www.stcn.com/article/detail/3598560.html>
- [3-92] National Business Daily × QuestMobile: Nearly 40% of Chinese Internet Users Have AI on Their Phones — Q1 2026 AI App Value Ranking: <https://www.nbd.com.cn/articles/2026-04-21/4350446.html>
- [3-93] 36Kr: The Offense and Defense of China's Three Major Content-Commerce Platforms in 2026: <https://m.36kr.com/p/3903803401127811>

Sources for Chapter 4

- [4-01] Bessemer Venture Partners (BVP Atlas): The AI Pricing and Monetization Playbook: How to Charge in a World Where Every Token Has a Cost (2026-02): <https://www.bvp.com/atlas/the-ai-pricing-and-monetization-playbook>
- [4-02] Genrise: AI Shopping Assistants 2026: What Rufus, Sparky, ChatGPT, and Perplexity Mean for Brands: <https://genrise.ai/insights/ai-shopping-assistants-future-ecommerce>
- [4-03] BCG: Agentic Commerce Is Redefining Retail — How Retailers Should Respond (2025): <https://www.bcg.com/publications/2025/agentic-commerce-redefining-retail-how-to-respond>
- [4-04] ZoomInfo Pipeline: Clari vs. Gong: The Two Revenue-Intelligence Titans Compared — the Prediction School vs. the Conversation School (Updated After 2025-12): <https://pipeline.zoominfo.com/sales/clari-vs-gong>

- [4-05] Gong: Gong Revenue AI: Turning Customer Interactions into Revenue Forecasts and Actionable Guidance: <https://www.gong.io/platform/revenue-ai/>
- [4-06] TechCrunch: Klarna Dropped Salesforce for Its Self-Built AI System: CEO Clarifies "We Did Not Replace SaaS with an LLM" (2025-03): <https://techcrunch.com/2025/03/04/klarna-ceo-doubts-that-other-companies-will-replace-salesforce-with-ai/>
- [4-07] CIO: Where Meta's WhatsApp AI Agent Can Actually Win: Local-Service Bookings Rather Than Physical-Goods E-Commerce (2026-07): <https://www.cio.com/article/4196328/where-metas-whatsapp-agent-can-actually-win.html>
- [4-08] Litslink: Netflix AI: How Netflix Personalizes Content for 325 Million Users (Third-Party Technical Analysis): <https://litslink.com/blog/all-about-netflix-artificial-intelligence-the-truth-behind-personalized-content>
- [4-09] Salesforce: Salesforce.com Uses Agentforce to Turn Its Website from an "Information Page" into a Selling Agent: <https://www.salesforce.com/customer-stories/agentforce-for-dot-com/>
- [4-10] CAICT: Intelligence-Driven Growth: Research Report on AI-Powered Customer Relationship Management (AI CRM) Systems (2025) (2025-09): <https://www.caict.ac.cn/kxyj/qwfb/ztbg/202509/P020250915603971582825.pdf>
- [4-11] CNBC: Amazon's "Leader's Dilemma": Fight AI Shopping Agents or Embrace Them (2025-12-24): <https://www.cnbc.com/2025/12/24/amazon-faces-a-dilemma-fight-ai-shopping-agents-or-join-them.html>
- [4-12] The SaaS CFO: How to Build Outcome-Based Pricing (2026-06): <https://www.thesaascfo.com/how-to-build-outcome-based-pricing/>
- [4-13] TechCrunch: Revenue-Prediction Startup Gong Surpasses \$300M ARR, Pointing Toward an IPO Path (2025-03-05): <https://techcrunch.com/2025/03/05/revenue-prediction-startup-gong-surpasses-300m-arr-indicating-potential-ipo-path/>
- [4-14] Sina Finance: Taobao's In-House 10-Billion-Parameter Recommendation Model RecGPT Goes Live, Sharply Improving the Accuracy of "Recommended for You": <https://finance.sina.com.cn/tech/digi/2025-07-01/doc-inficypra8733079.shtml>

- [4-15] Bloomreach: Increase E-Commerce Conversion Rates with AI Personalization (2025-11): <https://www.bloomreach.com/en/blog/increase-ecommerce-conversion-rates>
- [4-16] Reuters: Bipartisan US House Lawmakers Probe Delta Air Lines' AI Ticket Pricing: Beware "Surveillance-Based Personalized Price Discrimination" (2025-11-05): <https://www.reuters.com/world/us/us-house-lawmakers-probe-delta-air-lines-use-ai-ticket-pricing-2025-11-05/>
- [4-17] McKinsey & Company: Unlocking the Next Frontier of Personalized Marketing (2025): <https://www.mckinsey.com/capabilities/growth-marketing-and-sales/our-insights/unlocking-the-next-frontier-of-personalized-marketing>
- [4-18] Sina Finance: A Deep Analysis of Neocrm's AI CRM: From Ground-Up Rebuild to Business Results (2026-05): <https://t.cj.sina.com.cn/articles/view/3948743169/eb5d0a0100101cvuk>
- [4-19] Grocery Dive: AI Pricing Retailers Must Know: The Opportunities and the Risks (2025): <https://www.grocerydive.com/news/promise-peril-artificial-intelligence-driven-pricing-retailers/812037/>
- [4-20] Modern Retail: Instacart Launches Custom AI Assistants for Retailers — and Its Own: <https://www.modernretail.co/technology/instacart-launches-custom-ai-assistants-for-retailers-and-its-own/>
- [4-21] GoodyAds: Attribution Is Not Incrementality — How Overseas Advertisers Use Controlled Experiments to Answer Causal Questions: <https://www.goodyclick.com/incrementality-vs-attribution-causal-measurement-guide>
- [4-22] Growth Hackers Taiwan: What Is a Holdout Incrementality Test — Finding the Ads That Actually Work: <https://growthhackers.tw/blog/click-here-ep3-holdout-incremental-test>
- [4-23] PyMC Labs: Synthetic Consumers — A Practical Guide: <https://www.pymc-labs.com/blog-posts/synthetic-consumers-a-practical-guide>

Sources for Chapter 5

- [5-01] WeChat Official Account: 30-Day AI Practitioner | Day 14: AI + Operations Management (2) | Full-Lifecycle Customer Management: AI Automatically Maintains Existing Customers and Activates Repeat Purchases: <https://mp.weixin.qq.com/s/xxx>

- [5-02] Zendesk: AI for Customer Service & Support | Zendesk AI Platform: <https://www.zendesk.com/service/ai/>
- [5-03] Zendesk: The State and Future of AI in Customer Service: Zendesk's Industry Perspective: <https://www.zendesk.com/blog/ai-in-customer-service/>
- [5-04] Digital Applied: AI Customer-Support ROI Benchmarks (2026, 53 Data Points): \$3.50 Return per \$1 Invested on Average — but Vendor Claims vs. Independent Data Differ by 30–40 Percentage Points: <https://www.digitalapplied.com/blog/ai-customer-support-statistics-2026-adoption-roi-data>
- [5-05] Cresta: Best Practices for Designing AI-to-Human Handoffs: The Transfer Point Is Where the System Breaks Down Most Often: <https://cresta.com/guides/ai-to-human-agent-handoff-best-practices>
- [5-06] Voiceflow: Ada AI Review: The Automation-Resolution Pioneer and the Controversy of Per-Conversation Billing: <https://www.voiceflow.com/blog/ada>
- [5-07] Reuters: Air Canada must honor refund policy despite chatbot mistake - board rules: <https://www.reuters.com/technology/air-canada-must-honor-refund-policy-despite-chatbot-mistake-board-rules-2024-02-16/>
- [5-08] BBC: The Air Canada Chatbot Hallucination Case: Airlines Can't "Hide Behind Their Chatbot" (Ordered to Pay C\$812.02): <https://www.bbc.com/travel/article/20240222-air-canada-chatbot-misinformation-what-travellers-should-know>
- [5-09] CBC: Air Canada Ordered to Compensate a Passenger over Its Chatbot's Wrong Advice (Canada's First AI Customer-Service Liability Case): <https://www.cbc.ca/news/canada/british-columbia/air-canada-chatbot-lawsuit-1.7116416>
- [5-10] Voiceflow: Decagon AI Deep Review: From a Stealth Launch in 2023 to an AI Customer-Service Agent Vendor Valued at \$4.5B: <https://www.voiceflow.com/blog/decagon-ai>
- [5-11] CMSWire: Decagon Triples Valuation to \$4.5B With \$250M Series D: <https://www.cmswire.com/contact-center/decagon-triples-valuation-to-45b-with-250m-series-d>
- [5-12] Decagon: Decagon Raises Series C at \$1.5B Valuation to Deliver Concierge Customer Experience: <https://decagon.ai/blog/series-c-announcement>
- [5-13] Decagon: Decagon — The AI Concierge for Every Customer: <https://www.decagon.ai/>

- [5-14] Intercom (Fin): Fin — The Highest Performing Customer Agent: <https://fin.ai>
- [5-15] Intercom Community: Fin's Flawed Resolution Assumption (User Community Discussion): <https://community.intercom.com/fin-product-feedback-member-group-49/fin-s-flawed-resolution-assumption-10516>
- [5-16] Gartner: Gartner Predicts the Future of AI in Customer Service 2024-2025: <https://www.gartner.com/en/articles/gartner-predicts-the-future-of-ai-in-customer-service>
- [5-17] Decagon: How Hunter Douglas Is Turning AI-Powered Conversations into Revenue Across Its Brands: <https://www.decagon.ai/case-studies/hunter-douglas>
- [5-18] Securities Times: IDC: China's Enterprise-Grade Intelligent Customer-Service Market Reached 7.19 Billion Yuan in 2025, Moving from "Answering Questions" to "Getting Things Done": <https://www.stcn.com/article/detail/4169748.html>
- [5-19] Faye Digital: Intercom Fin AI Agent: The Product Logic Behind Its 50–70% Autonomous Resolution Rate, with Real-World Cases: <https://fayedigital.com/blog/fin-ai-agent/>
- [5-20] Intercom: Intercom Fin AI Copilot: Customer-Service Practice and Data Report: <https://intercom.com/blog/fin-ai-copilot-report>
- [5-21] Anthropic: Intercom Fin: AI Customer Service Moves from "Chatbot" to Agent: <https://www.anthropic.com/customers/intercom>
- [5-22] Intercom: Intercom Announces Fin AI: From "Chatbots" to AI Agents That Actually Solve Customer Problems: <https://www.intercom.com/blog/announcing-fin-ai/>
- [5-23] SEC EDGAR (Klarna Filing): Klarna AI Assistant: Customer Service Moves from Answering Questions to Executing Tasks: <https://www.sec.gov/Archives/edgar/data/2003292/000200329226000007/klar-20251231.htm>
- [5-24] GetPerspective: Klarna AI Customer Service: Replacing 700 Agents — A 2026 Case Study: <https://getperspective.ai/blog/klarna-ai-customer-service-replacing-700-agents-conversational-ai-case-study>
- [5-25] Klarna (Press): Klarna AI Assistant Handles Two-Thirds of Customer Service Chats: <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats/>

- [5-26] Klarna (Press): Klarna AI Assistant Handles Two-Thirds of Customer-Service Chats, Delivering a \$40M Improvement to Annual Profits: <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-for-a-40m-improvement-to-annual-profits/>
- [5-27] Klarna (Press): Klarna AI Assistant Case Study: Handling Two-Thirds of Customer-Service Conversations in Its First Month: <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-first-month/>
- [5-28] Klarna (Press): Klarna AI Assistant: 2.3 Million Conversations in One Month — the Equivalent of 700 Full-Time Agents, Projected to Add \$40M to Profits in 2024: <https://www.klarna.com/international/press/klarna-ai-assistant-handles-two-thirds-of-customer-service-chats-in-its-first-month/>
- [5-29] Forbes (Bernard Marr): Klarna's AI Customer-Service Strategy Revisited (2026): Over-Aggressive Layoffs Backfired, and the Value of "Roughly 100 Highly Skilled Humans" Was Reassessed: <https://www.forbes.com/sites/bernardmarr/2026/07/16/how-klarnas-ai-agent-strategy-backfired-but-became-a-useful-lesson/>
- [5-30] Business Insider: Klarna Is Reassigning Engineers and Marketers to Customer Support After AI Quality Concerns: <https://www.insider.com/klarna-reassigns-workers-to-customer-support-after-ai-quality-concerns-2025-9>
- [5-31] Baidu Baijiahao: The Klarna Lesson: AI Customer Service Is No Silver Bullet — from the "Equivalent of 700 Agents" to Recalling Humans: <https://baijiahao.baidu.com/s?id=1865577917080864958>
- [5-32] Qualtrics: AI-Powered Customer Service Fails at Nearly Four Times the Rate of Other AI Uses — "Customers Can Tell When AI Is Used to Save Money Rather Than Solve Problems": <https://www.qualtrics.com/news/ai-powered-customer-service-fails-at-four-times-the-rate-of-other-tasks/>
- [5-33] GetTectonic: Salesforce Unveils 'Agentforce' to Empower a Billion AI Agents by 2025: <https://gettectonic.com/agentforce-unveiled/>
- [5-34] Atlan: What Is Sierra AI? Bret Taylor and Clay Bavor's Enterprise AI-Agent Platform for Customer Service (Architecture, Pricing, and Use Cases): <https://atlan.com/know/ai-agent/ai-agent-applications/what-is-sierra-ai/>

- [5-35] CMSWire: Zendesk Expands AI Capabilities: AI Agents Near 5 Billion Resolved Tickets; AI ARR Projected at \$200 Million: <https://www.cmswire.com/ai-news/zendesk-expands-ai-capabilities-for-service-teams/>
- [5-36] Forbes (Gil Press): One Negative Chatbot Experience Drives Away 30% of Customers: Churn Risk and the Data on "Broken Handoffs to Humans": <https://www.forbes.com/sites/gilpress/2023/02/01/one-negative-chatbot-experience-drives-away-30-of-customers/>
- [5-37] Tuck School of Business (Dartmouth): Even with Humans in the Loop, Agentic AI Systems Struggle with Angry Customers — Field Evidence from Alibaba/Taobao: <https://tuck.dartmouth.edu/news/articles/even-with-humans-in-the-loop-agentic-ai-systems-struggle>
- [5-38] CMSWire: Stop Calling It "Support": Your Customer-Service Team Is a Growth Engine (The ROI Logic from Service to Retention, Repeat Purchase, and Word of Mouth): <https://www.cmswire.com/customer-experience/stop-calling-it-support-your-service-team-is-your-growth-strategy/>
- [5-39] Financial Times: Air Canada Ordered to Compensate a Passenger over an AI Chatbot Hallucination: A Landmark Case for AI Customer-Service Liability: <https://www.ft.com/content/air-canada-chatbot-ruling>
- [5-40] Zendesk: Great Service = Growth: Bad Experiences Churn Customers, Good Experiences Command a Premium, and "Customer-Obsessed Organizations" Reap Financial Returns (92 Curated Statistics): <https://www.zendesk.com/blog/customer-service/satisfaction/customer-service-statistics/>
- [5-41] Cloopen: Cloopen's Financial Large-Model Copilot & Agent: Real-Time Agent Assistance Lifts Marketing Conversion Rates by 30%: <https://www.yuntongxun.com/hangyeshendu/54>
- [5-42] Mark Read Fintech: When Bank Customer Service Is No Longer "Human": Agentic AI's Sprint and Brakes in Financial Services (Global Banking Cases + Red Lines for Human Handoffs): <https://www.markreadfintech.com/p/ai-677>
- [5-43] Fast Company: I Was Interviewed by AI: A Complete Account of Applying for a Customer-Service Role: <https://www.fastcompany.com/91422082/i-was-interviewed-by-ai>

- [5-44] ifanr: Burger King Fits Employees with AI Headsets ("Patty"): Every "Thank You" Is Being Scored by AI: https://www.ifanr.com/1656044?utm_source=rss&utm_medium=rss&utm_campaign=
- [5-45] LessWrong: The Economics of Replacing Call Center Workers with AIs (An In-Depth Breakdown): <https://www.lesswrong.com/posts/rJatmEDcYrDQcwstT/the-economics-of-replacing-call-center-workers-with-ais>
- [5-46] California Management Review (Berkeley): Chatbot Frustration Is Real: 53–77% of Users Have Hit Bad Experiences — A Full List of the Hidden Costs: <https://cmr.berkeley.edu/2026/04/chatbot-frustration-is-real-hidden-costs-and-best-practices/>
- [5-47] China Daily: OneConnect's "Xiaoyi" Intelligent Customer-Service Robot: Tiered Large- and Small-Model Architecture with Human-Machine Collaboration Pushes the Human-Agent Replacement Rate Above 60%: <https://cn.chinadaily.com.cn/a/202507/01/WS6863c522a3106af2b3c71f6a.html>
- [5-48] Xinhuanet (citing the China Consumers Association's H1 2026 Complaint Analysis): AI Customer Service Becomes a New Hotspot in Consumer Complaints: <https://www.news.cn/tech/20260817/f05985ebfc3b4a17bfbb774c62ada5ab/c.html>
- [5-49] Sohu Tech: AI Customer-Service Complaints Are Rising — More Than 4,800 on the Black Cat Platform, and "The AI's Answer Doesn't Represent the Company" Stops Being a Liability Shield From September: <https://timeline.sohu.com/news/jyultl60Zy>
- [5-50] VentureBeat: Intercom's New Post-Trained Fin Apex 1.0 Beats GPT-5.4 and Claude Sonnet 4.6: <https://venturebeat.com/technology/intercoms-new-post-trained-fin-apex-1-0-beats-gpt-5-4-and-claude-sonnet-4-6>
- [5-51] CX Dive: Klarna Pursues an "Uber-Style" Customer Service Model: <https://www.customerexperiencedive.com/news/klarna-pursues-uber-style-customer-service-model/812763/>
- [5-52] Decagon (customer case study, Notion Head of Global Customer Experience Emma Auscher): How Notion Scaled Global Customer Experience with Decagon: <https://decagon.ai/case-studies/notion>
- [5-53] CNBC (citing Qualtrics): "I Hate Customer-Service Chatbots" — AI Support's Rocky Consumer Debut: <https://www.cnbc.com/2026/04/01/ai-chatbot-customer-service-complaints-refunds.html>

Sources for Chapter 6

- [6-01] ChurnZero: 2025 Customer Revenue Leadership Study: NRR Stabilizes — How Team Composition and Tech Stack Correlate with NRR: <https://churnzero.com/customer-success-leadership-study/>
- [6-02] Greylock: The Rise of AI-Native User Research: From Scheduled Interviews to On-Demand Insights: <https://greylock.com/blog/ai-user-research/>
- [6-03] Planhat: AI for Finding Expansion Opportunities: The Revenue CSMs Miss Is Hidden in Six Types of Signals: <https://www.planhat.com/editorial/ai-for-expansion-opportunities>
- [6-04] Gainsight: Leveraging AI as a Customer Success Manager: A 2025–2026 Practical Guide to AI in Customer Success: <https://www.gainsight.com/essential-guide/leveraging-ai-as-a-customer-success-manager/>
- [6-05] Velaris: A Practical Guide to AI Customer Health Scoring (from a CSM Perspective): <https://www.velaris.io/articles/ai-for-customer-health-scoring>
- [6-06] ChurnZero: Customer Health Scores in the Age of AI: <https://churnzero.com/blog/customer-health-scores-in-the-age-of-ai/>
- [6-07] VirtualAssistantVA: AI Win-Back Campaign Benchmarks: 22% Reactivation Rates and up to 500% ROI Within Six Months — 2026 Data Roundup: <https://virtualassistantva.com/news/ai-customer-win-back-campaigns-22-percent-reactivation-rate-500-percent-roi-2026>
- [6-08] G2: AI Reduces Churn: Findings from G2's 2026 Expert Research: <https://learn.g2.com/ai-in-churn-reduction>
- [6-09] Gainsight: AI-Powered Customer Success: The Future of Revenue Retention and Expansion (Gainsight's View): <https://www.gainsight.com/blog/ai-powered-customer-success-the-future-of-revenue-retention/>
- [6-10] Hypers: The Core Value of AI Marketing in Customer Lifecycle Management: <https://www.hypers.com/content/archives/10220>
- [6-11] FunnelStory: Predicting Churn and Retention in B2B SaaS 101: Data, Models, and Evaluation: <https://funnelstory.ai/blog/predicting-churn-retention-in-b2b-saas-101>
- [6-12] GetPerspective: Customer Health Score Software in 2026: 8 Tools Compared: <https://getperspective.ai/blog/customer-health-score-software-2026-8-tools-compared>

- [6-13] CX Today: Gartner's 2026 Magic Quadrant for VoC Platforms: AI Reshapes the Voice-of-the-Customer Market Across 12 Vendors: <https://www.cxtoday.com/customer-analytics-intelligence/gartner-magic-quadrant-voc-platforms-2026/>
- [6-14] Lucid: How SaaS Startups Use AI to Predict Customer Churn: <https://www.lucid.now/blog/saas-startups-use-ai-predict-churn/>
- [6-15] DEMG: The SaaS Win-Back Engine: Four AI-Personalized Sequences That Recover Revenue from Churned Customers: <https://demg.ai/blog/saas-win-back-engine-ai-sequences/>
- [6-16] SaaStr: SaaStr Builds Its Own AI "VP of Customer Success" (Qbee): A Success Engine Managing 100+ Customers: <https://www.saastr.com/how-we-built-our-ai-vp-of-customer-success-qbee-and-how-you-can-build-yours>
- [6-17] ProsperStack: Five "Pattern-Interrupt" Win-Back Campaign Cases: Getir, Sky, Urbani Truffles, Dollar Shave Club, and E-Commerce: <https://prosperstack.com/blog/winback-campaign-examples/>
- [6-18] NIM: Between Questionnaires and In-Depth Interviews: How AI-Moderated Interviews Are Reshaping Qualitative Market Research (An Interview with xelper): <https://www.nim.org/en/publications/detail/how-ai-is-rewiring-qualitative-research>
- [6-19] CancelFlow: Why Most Health Scores Fail: Telemetry and Behavioral Data $\neq$ Live Intent (Corroborated by ChurnZero's 2025 Research: 73%): <https://cancelflow.dev/blog/saas-churn-health-score>
- [6-20] Statify: Net Revenue Retention (NRR) Benchmarks and Five Levers to Lift It (2025–2026): <https://www.statify.com/blog/net-revenue-retention-benchmarks-2025>
- [6-21] MDPI (Machine Learning and Knowledge Extraction): Customer Churn Prediction: A Systematic Review of ML/DL Progress, Trends, and Challenges (2020–2024): <https://www.mdpi.com/2504-4990/7/3/105>
- [6-22] Fanruan (FinePedia): How to Optimize the Customer Lifecycle? AI-Driven Automated Reporting Lifts Operational Efficiency: <https://www.fanruan.com/finepedia/article/68baa7110bd240a239aa903d>
- [6-23] Iliyana Stareva: Building an AI Customer Health Score (and How to Avoid Bias): <https://www.iliyanastareva.com/blog/ai-customer-health-scoring-how-to-build-it-and-how-to-avoid-bias>

- [6-24] Gainsight: Predicting and Preventing Churn with AI: <https://www.gainsight.com/blog/predicting-and-preventing-churn-with-ai/>
- [6-25] Emerald (Innovation & Management Review): Churn Prediction for a SaaS Company with Machine Learning: The TecnoSpeed, Brazil, Case: <https://www.emerald.com/inmr/article/22/2/130/1251238/Churn-prediction-for-SaaS-company-with-machine>
- [6-26] US Tech Automations: E-Commerce Win-Back Case Study: Segmented Automated Sequences Reactivate 15.3% of Dormant Customers and Recover \$187K: <https://ustechautomations.com/resources/blog/ecommerce-win-back-campaign-case-study-15-percent-reactivation>
- [6-27] ChurnSolution: The 2026 Playbook for Subscription Win-Back Campaigns: Nearly 1 in 4 New Subscriptions Come from Formerly Churned Customers — Win-Back Probability of 20–40% vs. 5–20% for New Customers: <https://churnsolution.com/2026/06/18/the-playbook-for-subscription-win-back-campaigns-2026-revenue-recovery-guide/>
- [6-28] AIBP (AI Best Practices): Win-Back Campaign Automation Case Studies: AI Reactivation Results for Cashback Apps, Fintech, and Financial Customers: <https://ai-best-practices.com/case-studies/customer-win-back-campaign-automation>
- [6-29] Gruv: The Economics of Winning Back Churned Subscribers: When It Pays Off More Than Acquiring New Customers: <https://gruv.ai/blog/customer-winback-economics-churned-subscriber>
- [6-30] Gainsight: What Customer Success Teams Are Prioritizing in 2026 — and What the Data Shows Is Working: <https://www.gainsight.com/blog/what-customer-success-teams-are-prioritizing-in-2026-and-what-the-data-shows-is-working>
- [6-31] 1337 Sales (vendor case study): AI Customer Success Expansion Agent — Churn Down 40%, \$1.2M in New Expansion Revenue in Six Months, NRR Above 115%: <https://1337sales.com/case-studies/customer-success-expansion-agent>
- [6-32] Bain & Company: Why SaaS Stocks Have Dropped and What It Signals for Software's Next Chapter: <https://www.bain.com/insights/why-saas-stocks-have-dropped-and-what-it-signals-for-softwares-next-chapter/>

- [6-33] ChurnZero: ChurnZero Launches Retrospective AI Churn Analysis Agent: <https://churnzero.com/press-release/churnzero-launches-retrospective-ai-churn-analysis-agent/>
- [6-34] ChurnZero (customer case study): Lunit International — CSMs as Business Owners: <https://churnzero.com/case-studies/lunit-international/>
- [6-35] AdExchanger: Sam's Club Debuts Ad Targeting Suite That Confidently Asserts Future Buyers: <https://www.adexchanger.com/commerce/sams-club-debuts-ad-targeting-suite-that-confidently-assert-future-buyers/>
- [6-36] SaaStr (author Jason Lemkin): Should You Discount to Save a Renewal? The NRR Trap Says No: <https://www.saastr.com/should-you-discount-to-save-a-renewal-the-nrr-trap-says-no-you-probably-should-anyway/>
- [6-37] Huxiu: The Real Life-or-Death Line for Children's AI Hardware Is 180-Day Retention — An Interview with Tinglixiong Founder Yuan Lin: <https://www.huxiu.com/article/4892015.html>

Sources for Chapter 7

- [7-01] Tealium: 2025 State of the CDP Report: 92% of CDP Users Meet Their Goals and 84% Say the CDP Simplified AI Projects (Official Survey): <https://tealium.com/resource/whitepaper/2025-state-of-the-cdp/>
- [7-02] Osano: 2025 Marketing Data Privacy Compliance Guide: GDPR/CCPA Consent, Auditing, and Privacy-First Marketing (Practical Guide): <https://www.osano.com/articles/marketing-data-privacy-guide>
- [7-03] SuperOffice: 50+ CRM Statistics for 2026: The Biggest Bottleneck for AI-CRM Is Data, Not Technology (Trend Report): <https://www.superoffice.com/blog/50-crm-statistics/>
- [7-04] Altudo: AI-Powered Customer Identity Resolution: The Foundation of a Unified Customer Experience — and Why It Should Precede Personalization and AI Initiatives (Consulting Perspective): <https://www.altudo.co/insights/blogs/ai-powered-customer-identity-resolution-for-cx-transformation>
- [7-05] California Management Review (Berkeley): Balancing Personalized Marketing and Data Privacy in the Age of AI: 60% of Large Enterprises Automated GDPR Compliance with AI in 2025 (Academic Perspective): <https://cmr.berkeley.edu/2025/02/balancing-personalized-marketing-and-data-privacy-in-the-era-of-ai/>

- [7-06] HubSpot: Agent Hub — All of Your AI Agents, Managed in One Place: <https://www.hubspot.com/products/artificial-intelligence>
- [7-07] 11x: Agentforce Review 2026: The Strengths, Shortcomings, and Real Adoption Rates of Salesforce's Autonomous AI Platform (Third-Party Review): <https://www.11x.ai/guides/agentforce-review>
- [7-08] Salesforce: Agentforce: The AI Agent Platform: <https://www.salesforce.com/agentforce/>
- [7-09] The Starr Conspiracy: B2B Intent Data Benchmarks 2025: 21.3% vs. 8.4% Conversion Rates and a 62% Signal-Validation Failure Rate (Benchmark Library): <https://www.thestarrconspiracy.com/insights/benchmarks/b2b-intent-data-benchmarks-2025>
- [7-10] Salesloft: Blue Yonder's ABM with 6sense Intent Data and a Drift Chatbot: Accounts in the Buying Stage up 40%, CPL Down to \$145: <https://www.salesloft.com/resources/case-studies/blue-yonder>
- [7-11] MarketBetter: Bombora Company Surge 2026 Independent Review: \$30K+ Annual Fee and a Claimed 81% Intent Accuracy — but Company-Level Signals Only, with Execution Tools Sold Separately: <https://marketbetter.ai/blog/bombora-review-2026/>
- [7-12] CMSWire: Treasure Data Unveils AI Marketing Cloud and Its Super-Agent Vision at CDP World 2025 (Third-Party Report): <https://www.cmswire.com/customer-data-platforms/treasure-data-unveils-ai-marketing-cloud-and-super-agent-vision-at-cdp-world-2025/>
- [7-13] Big Technology: Google Cloud CEO Kurian on Gemini Enterprise: Enterprise Data + an Agent Builder + Customer-Service Tools: <https://www.bigtechnology.com/p/google-cloud-ceo-thomas-kurian-on-ecd>
- [7-14] HubSpot: Breeze: HubSpot's AI Suite for Customer Operations, Built for SMBs: <https://www.hubspot.com/products/artificial-intelligence/breeze>
- [7-15] CMSWire: HubSpot Shifts Breeze AI Agents to Pay-Per-Result Pricing: \$0.50 per Resolved Conversation (Third-Party Report): <https://www.cmswire.com/customer-experience/hubspot-shifts-breeze-ai-agents-to-pay-per-result-pricing/>
- [7-16] SaaStr: Marc Benioff: A Separate "AI Budget" Is the Wrong Framing — Every Budget Is an AI Budget: <https://www.saastr.com/marc-benioff-ceo-salesforce-on>

- [7-17] Simply Wall St: Salesforce Agentforce in the Enterprise: Smarsh Achieves 72% Self-Service Deflection (2026-09): <https://simplywall.st/stocks/us/software/nyse-crm/salesforce/news/salesforce-crm-agentforce-shows-72-self-service-deflection-i>
- [7-18] Salesforce (Investor Relations): Salesforce Delivers Record Third Quarter Fiscal 2026 Results, Driven by Agentforce & Data 360: <https://investor.salesforce.com/news/news-details/2025/Salesforce-Delivers-Record-Third-Quarter-Fiscal-2026-Results-Driven-by-Agentforce--Data-360/default.aspx>
- [7-19] Salesforce: Salesforce Introduces Agentforce, the Digital Labor Platform for Enterprises: <https://www.salesforce.com/news/press-releases/2024/09/09/agentforce-announcement/>
- [7-20] Salesforce: Salesforce Announces Agentforce: Autonomous AI Agents Enter the Full Enterprise Sales and Service Workflow: <https://www.salesforce.com/news/press-releases/2024/09/12/agentforce-announcement/>
- [7-21] Salesforce: Salesforce Rolls Out Agentforce: From Copilot Toward an Autonomous Agentic CRM: <https://www.salesforce.com/news/press-releases/2024/09/17/agentforce-announcement/>
- [7-22] Salesforce: Salesforce's Agentforce Is Here: Trusted, Autonomous AI Agents to Scale Your Workforce: <https://www.salesforce.com/news/press-releases/2024/10/29/agentforce-general-availability-announcement/>
- [7-23] East Money (dfcfw.com PDF archive): Artificial Intelligence Customer Relationship Management (AI CRM) Systems Research Report: https://pdf.dfcfw.com/pdf/H3_AP202509191746485186_1.pdf?1758277576000.pdf
- [7-24] CMSWire: What Is a Customer Data Platform (CDP)? A 2026 Market Guide — from Optional Tool to the Core of Operations (Market Overview): <https://www.cmswire.com/customer-data-platforms/what-is-a-customer-data-platform-cdp/>
- [7-25] Digiday: What Is Federated Learning? Training Models on Local Data for Interest Targeting in a Cookie-Free World (Explainer): <https://digiday.com/media/what-is-federated-learning/>
- [7-26] Zhihu Column (Neocrm Official): From Engage2025 to 2026: One Year of the Agent Explosion — Where Neocrm's AI CRM (NeoAgent 2.0) Stands Now (Official): <https://zhuanlan.zhihu.com/p/2076619354880358366>

- [7-27] SAMexpert: Microsoft's Sales AI Portfolio, Explained: Bundled Copilot Offerings vs. Consumption-Billed Sales Agents (Product Deep Dive): <https://samexpert.com/microsoft-sales-ai-products-vs-agents/>
- [7-28] Autobound: Intent Data Providers Compared (2026): A \$4.49B Market, and 94% of Buying Groups Settle on a Preferred Vendor Before Ever Talking to Sales (Market Comparison): <https://www.autobound.ai/blog/top-15-intent-data-providers-compared-2026>
- [7-29] Salesmotion: Intent Data Platforms Deep Review and Real Prices: 6sense at a \$62K Median Annual Fee; Bombora's Co-Op Network Covers 17,000+ Topics (Review): <https://salesmotion.io/blog/intent-data-providers>
- [7-30] PPC Land: The Intent-Data Channel War: Intentsify's 1.1 Trillion Signals Join the Clay Data Marketplace as Intent Data Sinks into Agentic Prospecting Workflows: <https://ppc.land/clay-gains-intentsify-1-1-trillion-intent-signals-via-new-partnership/>
- [7-31] Fenxiangxiaoke: Fenxiangxiaoke ShareAI: 15,000 Words on the Platform Architecture and On-the-Ground Results of AI-Native CRM (Official): <https://www.fxiaokey.com/crm/blog-80128.html>
- [7-32] OmniFunnelMarketing: Zero-Party Data Strategy: Collection Methods in the Age of Customer Trust — and How AI Amplifies Them (Strategy Overview): <https://www.omnifunnelmarketing.com/blog/building-zero-party-data-collection-strategy-customers-trust>
- [7-33] King & Wood Mallesons: Mastering Uni Marketing: Data-Compliance Paths for Retail Marketing in the Digital Era (Part 2) (PIPL in Practice): <https://www.kingandwood.com/cn/zh/insights/latest-thinking/uni-marketing-data-compliance-in-retail-marketing-in-era-of-digitalization.html>
- [7-34] Everyone Is a Product Manager (author Congcong): Private-Domain Reach Rate From 4.04% to Above 10% — Replacing "Three Behavior Buckets + Four Fixed RFM Labels" with AI-Driven Three-Dimensional Dynamic Segmentation Lifted Repurchase Rate 5% Year Over Year: <https://www.woshipm.com/ai/6464019.html>
- [7-35] OpenInstall: Meituan's AI Agent Xiaomei Connects Into Tencent Yuanbao — The Attribution Blind Spot of Cross-Ecosystem Agents: <https://www.openinstall.com/article/news/meituan-xiaomei-tencent-yuanbao-agent>

- [7-36] Forbes Technology Council (author Celebrus CEO Bill Bruno): Why the Next Evolution of the CDP Starts Before Customer Data: <https://www.forbes.com/councils/forbestechcouncil/2026/09/02/why-the-next-evolution-of-the-cdp-starts-before-customer-data/>
- [7-37] 6sense: 6sense MCP Server — Expanded Capabilities and New Sales Intelligence Features: <https://revcity.6sense.com/home/discussion/2584/6sense-mcp-server-expanded-capabilities-new-sales-intelligence-features>
- [7-38] Salesforce (Dreamforce 2026): Introducing Alforce — CRM Capabilities Everywhere Agents Work: <https://www.salesforce.com/news/stories/aiforce-announcement/>
- [7-39] Model Context Protocol (official registry): The Official MCP Registry — The Distribution and Discovery Entry Point for Agent Tools: <https://registry.modelcontextprotocol.io>
- [7-40] Yilantop: Alibaba Cloud Releases Fanhua Program 2.0, Upgrading Its Cloud Marketplace into an AI Application Marketplace: <https://www.yilantop.com/article/27146>
- [7-41] Cloudflare Blog: Building the Agent Cloud — Cloudflare's Agents Week 2026 Product Announcements: <https://blog.cloudflare.com/zh-cn/agents-week-in-review>

Chapter 8 Sources

- [8-01] From 0 to 20 AI Agents in 10 Months: SaaStr's Production-Deployment Playbook: <https://www.saastr.com/from-zero-to-20-ai-agents-in-10-months-the-saastr-playbook-for-actually-deploying-ai-agents-that-work>
- [8-02] The State of AI 2025 Global Survey: Value from Pilots to Scale Still Not Realized (McKinsey Global Survey): <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai-2025>
- [8-03] The Future State of Sales 2025 (Mercuri International Research Report, PDF): <https://mercuri.nl/wp-content/uploads/sites/15/2025/02/future-state-sales-2025.pdf>
- [8-04] 2026 Guide to Building Customer-Service System Use Cases: From Pain-Point Analysis to On-the-Ground Practice (Tencent Cloud Developer Community, 2026-04-21): <https://cloud.tencent.com/developer/article/2657797>

- [8-05] Best AI Sales Roleplay Tools of 2026: 8 Platforms Reviewed (Knowlee, Updated April 2026): <https://www.knowlee.ai/blog/best-ai-sales-roleplay-tools-2026>
- [8-06] The 60% AI Team: More Efficient, More Profitable, but Surprisingly Lonely (SaaStr): <https://www.saastr.com/building-a-60-ai-team-efficient-profitable-and-surprisingly-lonely>
- [8-07] AI Index Report 2024 - Chapter 7: Economy: <https://hai.stanford.edu/research/hai-ai-index/ai-index-2024>
- [8-08] AI Role-Play: Accelerating Sales Ramp and Improving New-Hire Success Rates (Hyperbound): <https://www.hyperbound.ai/blog/ai-roleplay-sales-ramp>
- [8-09] The State of AI SDR 2026: Adoption, Cost, and Performance Benchmarks (Laxis Research): <https://www.laxis.com/blog/state-of-ai-sdr-2026/>
- [8-10] AI SDRs: What Works and What Doesn't (A First-Hand Review of an 800-Account Controlled Experiment): <https://greenway.ai/post/ai-sdr-what-works-what-doesnt/>
- [8-11] How AI Will Redefine Sales Productivity: From Activity Metrics to EQ Metrics (Vivun x Gartner 2026): <https://www.vivun.com/blog/how-ai-will-redefine-sales-productivity-in-2026>
- [8-12] Sales Operations Planning in the AI Era (Gartner): <https://www.gartner.com/en/articles/sales-operations>
- [8-13] AI Is Transforming Productivity, but Sales Remains the New Frontier (Bain 2025 Technology Report): <https://www.bain.com/insights/ai-transforming-productivity-sales-remains-new-frontier-technology-report-2025/>
- [8-14] The State of AI Sales Agents 2026: Adoption, ROI, and Market Benchmarks (Laxis Research Industry Report): <https://www.laxis.com/blog/state-of-ai-sales-agent-2026/>
- [8-15] 30 Statistics on AI in Sales Enablement for 2026 (Cubeo AI Executive Briefing): <https://www.cubeo.ai/30-statistics-of-ai-in-sales-enablement-in-2026/>
- [8-16] AI Sales Role-Play Fully Explained: 2026 Guide to Selecting and Practicing Intelligent Sales-Training Systems (Chinese Industry Overview): <https://blogs.findpowerai.com/?p=1526>
- [8-17] The Complete Guide to AI Sales Role-Play and Deal-Coaching: 2025 Hands-On Playbook for Enterprise Sales Enablement: https://www.exexm.com/article_detail.php?id=591

- [8-18] 6 Skills Sales Teams Need in the Age of Agentic AI (Seismic): <https://www.seismic.com/blog/agentic-ai-era-sales-skills/>
- [8-19] How Asymbi Uses Agentforce to Grow Like a Company One-Tenth Its Size (Salesforce Customer Story): <https://www.salesforce.com/customer-stories/asymbi/>
- [8-20] The Future of B2B Sales: How Growth Champions Rewrite Their Playbooks with AI (McKinsey): <https://www.mckinsey.com/capabilities/growth-marketing-and-sales/our-insights/the-future-of-b2b-sales-how-growth-champions-rewire-their-playbooks-with-ai>
- [8-21] BCG: 74% of Companies Can't Scale AI Value — Roughly 70% of Deployment Barriers Come from "People," Not Technology or Algorithms: <https://www.bcg.com/press/24october2024-ai-adoption-in-2024-74-of-companies-struggle-to-achieve-and-scale-value>
- [8-22] Duolingo CEO's Company-Wide Letter: Five AI-First Practices (Including AI Deciding Performance and Layoffs): <https://www.businessinsider.com/duolingo-ceo-how-ai-will-be-used-performance-reviews-headcount-2025-4>
- [8-23] A Fortune 500 Data-Consulting Firm Rebuilds Outbound with AiSDR: 8,000+ Emails in 4.5 Months (Official Case Study): <https://aisdr.com/case-study-fortune-500-company/>
- [8-24] Gartner: Only 32% of Leaders Achieve "Healthy Change Adoption"; 79% of Employees Have Low Trust in Change: <https://www.gartner.com/en/newsroom/press-releases/2025-07-08-gartner-hr-research-finds-just-32-percent-of-business-leaders-report-achieving-healthy-change-adoption-by-employees>
- [8-25] Microsoft/LinkedIn 2024 Work Trend Index: 78% of Employees Are "Bringing Their Own AI" to Work; Half Won't Admit Using AI for Key Work: <https://www.microsoft.com/en-us/worklab/work-trend-index/ai-at-work-is-here-now-comes-the-hard-part>
- [8-26] One Year Since Gartner's "80% AI Resolution" Prediction — 2026 Reality Check (Maven AGI): <https://www.mavenagi.com/resources/one-year-since-gartners-ai-resolution-prediction>
- [8-27] Only One-Quarter of AI Customer Service Use Cases Produce ROI (CX Dive): <https://www.customerexperiencedive.com/news/only-one-quarter-of-ai-customer-service-use-cases-produce-roi/827951/>

- [8-28] How OpenAI Built a Global Customer Success Organization in 18 Months: From the First CS Hire to Teams in 8 Locations: <https://www.saastr.com/how-openai-built-a-global-cs-organization-in-18-months-vanessa-gutiehis-playbook-for-ai-powered-customer-success>
- [8-29] Pew: More Than Half of U.S. Workers Worry About AI in the Workplace — “Fear of Being Replaced” Outweighs Hope: <https://www.pewresearch.org/social-trends/2025/02/25/workers-views-of-ai-use-in-the-workplace/>
- [8-30] ROI of AI Customer Service: 2026 Benchmarks & Data (Intercom Fin, Official): <https://fin.ai/learn/roi-ai-customer-service-agents-benchmarks>
- [8-31] Salesforce 2026 State of Sales: AI and Agents Become the No. 1 Growth Tactic (Survey of 4,050 People): <https://www.salesforce.com/news/stories/state-of-sales-report-announcement-2026/>
- [8-32] Second Nature vs. Hyperbound: Two Routes for AI Sales Role-Play Platforms (Official Second Nature Comparison): <https://secondnature.ai/second-nature-vs-hyperbound-which-ai-sales-roleplay-platform-is-right-for-your-team/>
- [8-33] Second Nature Publishes Oracle NetSuite Case Study: AI Role-Play Training Drives +32% More Sales Opportunities (Third-Party Coverage): <https://completeaitraining.com/news/second-nature-links-ai-sales-role-play-platform-to-32/>
- [8-34] The ServiceTitan AI Playbook: How the CRO Uses AI to Rebuild Sales and Customer Operations (\$30M → \$860M ARR): <https://www.saastr.com/the-servicetitan-ai-playbook-how-cro-ross-use-ai-to-build-a-merit-based-sales-machine>
- [8-35] The Essential Customer Success Trends of 2026 — Expert Predictions (ChurnZero Interview Series, 2026): <https://churnzero.com/blog/2026-customer-success-trends-experts/>
- [8-36] Sinostrong AI Practice-Coach System: A 7x24 Hands-On Sales-Training Platform (Vendor Product Page): <https://www.sinostrong.com/products/ai.html>
- [8-37] Bringing Deep Research into GTM: Kyle Poyar’s Hands-On Workflow: <https://www.growthunhinged.com/p/deep-research-for-gtm>
- [8-38] Who Manages the AI Agent Fleet: RevOps as the Management Layer (RevSure Playbook): <https://www.revsure.ai/research-playbooks/who-owns-the-agents>

- [8-39] 10 Mistakes Almost Everyone Makes After Deploying 20+ AI Agents (SaaStr): <https://www.saastr.com/weve-deployed-20-ai-agents-here-are-the-10-mistakes-almost-everyone-makes>
- [8-40] Generative AI in Sales: 7 Use Cases Ranked by Revenue Impact (SetSmart, Based on 828,000 Conversations): <https://setsmart.io/blog/generative-ai-for-sales>
- [8-41] Datawhale's "FDE Case 100": 24 Frontline FDE Interviews on Enterprise AI Deployment: <https://assets.datawhale.cn/Datawhale%20FDE%E6%A1%88%E4%BE%8B100.pdf>
- [8-42] Applied AI Doesn't Work (Varick): <https://x.com/vasuman/status/2095999742031675738>
- [8-43] AI-Native Leaders: The Organizational Playbook for Engineering Transformation at Scale (ByteByteGo): <https://blog.bytebytego.com/p/ai-native-leaders-the-organizational>
- [8-44] AI-Native Engineering Teams: 10 Practices That Separate the Best (2026) (Coderio): <https://www.coderio.com/blog/biz-tech/ai-native-engineering-teams-10-practices/>
- [8-45] Organizational Transformation in the Age of AI: How Companies Maximize AI's Potential (WEF x Accenture White Paper): https://www.sohu.com/a/1013292532_121838863
- [8-46] Bain & Company: AI Won't Transform Your Go-to-Market Until You Transform the Work (2026 B2B Growth Agenda; 1,125 Leaders Surveyed): <https://www.bain.com/insights/ai-wont-transform-your-go-to-market-until-you-transform-the-work-b2b-growth-agenda-2026/>
- [8-47] Atlassian: The State of Teams 2026 — CMO Brief (94% Use AI; 3% of CMOs Are Confident Returns Are Quantifiable): <https://www.atlassian.com/blog/guides-research/the-state-of-teams-2026-cmo-brief>
- [8-48] Salesforce: What Leaders Should Do About American Workers' AI Skepticism (1,500+ Office Workers Surveyed): <https://www.salesforce.com/news/stories/what-leaders-should-do-about-american-workers-ai-skepticism/>
- [8-49] Salesforce: Agentic AI Leaders Survey on ROI (2,025 Enterprise Decision-Makers Surveyed): <https://www.salesforce.com/news/stories/agentic-ai-leaders-survey-on-roi/>

- [8-50] TMTPost: Zhipu and MiniMax Double Revenue While Cutting Sales Spending — The Sales Role Shifts From Pre-Sales to Mid- and Post-Sales: <https://www.tmtpost.com/8141022.html>
- [8-51] Salesforce: Agentic Enterprise Index Insights 2026: <https://www.salesforce.com/uk/news/stories/agentic-enterprise-index-insights-2026/>
- [8-52] ContentGrip (reading Gartner's CMO Spend Survey): CMO AI Spending Rises, but the Readiness Gap Remains (2026): <https://www.contentgrip.com/cmo-ai-readiness-gap>
- [8-53] The Ad Spend (based on Gartner's 2026 CMO Spend Survey): Marketing Org Design in the AI Era — A CMO Guide (2026): <https://theadspend.com/blog/marketing-org-design-ai-era>

Chapter 9 Sources

- [9-01] 2025 Year-End Recommendations for Intelligent AI Customer-Service Vendors: Focusing on Government/Enterprise and E-Commerce Scenarios: <https://www.cnblogs.com/pinpaituijan/p/19351238>
- [9-02] Six Legal-Tech Trends for 2026: Lawyer AI Adoption Rises from 11% to 30% — Firms That Use Technology to Capture Cases Generate 51% More Leads and 52% Higher Revenue: <https://getperspective.ai/blog/legal-tech-trends-2026-6-data-backed-shifts-law-firm-ai-adoption>
- [9-03] Fxiaoke: Deep Dive into the 2026 AI CRM System: <https://www.fxiaokey.com/crm/blog-85286.html>
- [9-04] AI Agent Market Research & Agentic Commerce Protocol: <https://github.com/lemonadesocial/atlas-protocol/blob/main/whitepaper/MARKET-RESEARCH.md>
- [9-05] AI Writing SaaS QuillBot Runs PLG Customer Operations on Amplitude: Self-Service Analytics for Everyone Drives Retention and the Growth Flywheel: <https://amplitude.com/case-studies/quillbot>
- [9-06] AI Becomes the Banker: 21 Digital-Banking Customer-Experience Case Studies (2024–2025): <https://www.finextra.com/blogposting/28841/ai-becomes-the-banker-21-case-studies-transforming-digital-banking-cx>
- [9-07] AI Breaks Through GTM: A Practitioner's Guide: <https://www.woshipm.com/ai/6192618.html>

- [9-08] How AI Is Reinventing the Wealth Advisor: The Continuum Story (Digital Native): <https://www.digitalnative.tech/p/how-ai-is-reinventing-wealth-advisory>
- [9-09] AI + Healthcare: Promising Prospects, but the Boundaries Still Need to Be Clarified: <http://www.news.cn/20260123/5448ceb1e3c7437b97c9b24f19f624e2/c.html>
- [9-10] How AI Rebuilds Traditional Patient-Management Models: <https://www.woshipm.com/?p=5799425>
- [9-11] Five Application Scenarios for AI Agents in B2B Marketing: Practice Paths and Real-World Cases: https://blog.csdn.net/JINGdigital_/article/details/148333942
- [9-12] AI-Empowered B2B Marketing: Five Stages of Full-Process Intelligence Practice and Value Creation: <https://www.woshipm.com/marketing/6250810.html>
- [9-13] AI Reshapes B2B Marketing: From Linear Funnels to Intelligent Growth Ecosystems: <https://blog.csdn.net/q997815917/article/details/149907431>
- [9-14] Air Canada’s chatbot illustrates persistent agency and responsibility: <https://link.springer.com/article/10.1007/s00146-024-02096-7>
- [9-15] Amazon announces Rufus, a new generative AI-powered conversational shopping assistant: <https://www.aboutamazon.com/news/retail/amazon-rufus>
- [9-16] Avangrid Utility Launches AI Energy Assistant “Ava”: First Covering Billing/Outage/Start-Stop Service Q&A to Divert Volume from the Call Center: <https://www.avangrid.com/w/avangrid-launches-ai-energy-assistant-to-enhance-customer-satisfaction-and-experience>
- [9-17] BCG Report: AI-Driven Customer Retention and Churn Warning in Financial Wealth Management: <https://www.bcg.com/publications/2024/ai-in-financial-services-wealth-management-customer-retention>
- [9-18] BofA’s Erica Surpasses 2 Billion Interactions, Helping 42 Million Clients: <https://newsroom.bankofamerica.com/content/newsroom/press-releases/2024/04/bofa-s-erica-surpasses-2-billion-interactions--helping-42-millio.html>
- [9-19] CSX Chessie: Logistics Customers Move from Asking Questions to Directly Calling Business Data: <https://www.microsoft.com/en/customers/story/24556-csx-microsoft-copilot-studio>
- [9-20] Caixin Weekly | Examining the State of “AI Teachers”: <https://www.caixinglobal.com/2025-07-18/caixin-weekly-examining-the-state-of-ai-teachers-102342780.html>

- [9-21] Chime Scales Member Support with AI-Powered Automation — Delivering Efficiency and a Stronger Customer Experience: <https://www.decagon.ai/case-studies/chime>
- [9-22] Clio Releases Clio Duo, a Built-In Generative AI Assistant: Embedding AI Across Legal Practice Management and Client Communication: <https://www.lawnext.com/2024/10/clio-releases-clio-duo-generative-ai-built-directly-into-its-law-practice-management-platform.html>
- [9-23] Davana Law Firm Grows with Clio Management + AI: Clients Scale from 30 to 3,000+ (100×), Recovering Over \$52 Million for Clients: <https://www.casestudies.com/company/clio/case-study/how-davana-law-firm-grew-exponentially-with-clio-solutions>
- [9-24] Decagon review: Is it the right AI agent for you in 2026?: <https://www.eesel.ai/blog/decagon-review>
- [9-25] Duolingo Max: Using GPT-4 for AI-Powered Personalized Conversation Learning to Boost Retention and Paid Subscriptions: <https://blog.duolingo.com/duolingo-max/>
- [9-26] Duolingo’s AI-First Strategy: DAU Up 51% YoY, MAU Reaches 130 Million, 2025 Revenue Outlook Tops \$1 Billion: <https://chiefaiofficer.com/duolingos-ai-strategy-fuels-51-user-growth-and-1b-revenue/>
- [9-27] Evri: Automating Customer Service Ultimately Improves the “Real-Person Phone Channel”: <https://www.microsoft.com/en/customers/story/25358-evri-microsoft-power-platform>
- [9-28] Fidal IA: Internal Knowledge Assets Are the Key to Customer Operations in Professional Services: <https://www.microsoft.com/en/customers/story/25200-fidal-azure-openai>
- [9-29] Forrester TEI Study: Amplitude Customers See 217% ROI in Three Years — Self-Service Analytics Drives Acquisition and Retention Gains: <https://investors.amplitude.com/news-releases/news-release-details/amplitude-customers-saw-217-return-investment-according>
- [9-30] The New GTM Challenge: Your Next Customer May Be an AI: <https://www.huxiu.com/article/4858286.html>
- [9-31] Google Cloud’s 101 Real-World Generative AI Use Cases (Including Customer Service and Lifecycle): <https://cloud.google.com/transform/101-real-world-generative-ai-use-cases-from-industry-leaders>

- [9-32] Greenfield Solar Uses AI Voice/SMS Agents for Lead Qualification: Close Rate Jumps from 11% to 29%, First Response Under 30 Seconds: <https://www.plura.ai/case-studies/solar-company-ai-intelligence-18-percent-conversion>
- [9-33] Hims & Hers Launches Labs AI: The First AI Care Agent, Turning 130 Biomarker Test Results into Personalized Insights: <https://news.hims.com/newsroom/meet-labs-ai-the-first-ai-care-agent-from-hims-hers>
- [9-34] How ClassPass Embraces AI Across the Entire Customer Stack: <https://www.decagon.ai/case-studies/classpass>
- [9-35] How Duolingo Partners with Decagon (Duolingo English Test Customer-Support Transformation): <https://www.decagon.ai/case-studies/duolingo>
- [9-36] Indeed Will Speed Up Hiring with Agentforce's Digital Workforce: <https://www.salesforce.com/customer-stories/indeed/>
- [9-37] Intercom's Claude-Powered AI Support Agent Fin: Resolving Up to 86% of Support Volume, Response Time Cut from 30 Minutes to Seconds: <https://claude.com/customers/intercom>
- [9-38] Intercom's AI Support Product Fin Races Toward \$100M ARR: 7,000+ Customers, 51% Out-of-the-Box Resolution Rate, 99.9% Accuracy: <https://www.saasultra.com/intercoms-fin-is-the-ai-agent-racing/>
- [9-39] Klarna's AI Assistant Automates Two-Thirds of Customer Service (2024-2025 Update): <https://www.usefini.com/blog/klarna-automates-two-thirds-of-customer-service-with-ai-assistant>
- [9-40] Klarna says its AI agent is doing the work of 853 employees: <https://finance.yahoo.com/news/klarna-says-ai-agent-doing-070000869.html>
- [9-41] Kraken Official Case Study: Octopus Energy Cuts Service Costs 40% with a Unified Support Platform + AI Assistant; AI Drafts 60% of Emails: <https://www.kraken.tech/case-studies/octopus-energy-customer>
- [9-42] Large language models in healthcare: A survey of applications and challenges: <https://www.nature.com/articles/s41591-024-02802-4>
- [9-43] Netflix Uses AI Personalization to Turn Content Recommendations into a Retention Engine: 80% of Viewing Comes from Recommendations; the Company Says It Saves Over \$1 Billion a Year in Churn Costs: <https://marketingino.com/the-netflix-recommendation-algorithm-how-personalization-drives-80-of-viewer-engagement/>

- [9-44] Octopus Energy Unlocks Support Capacity with AI Email Tool Magic Ink: AI Drafts 35% of Emails; Satisfaction Beats Human Agents by 15 Percentage Points: <https://transformationplaybook.ai/blog/octopus-energy-ai-customer-service>
- [9-45] Quadrivia: Healthcare AI Extends from Q&A to the Full Patient Journey: <https://cloud.google.com/customers/quadrivia>
- [9-46] SIAG: Combining Citizen Relationship Management with AI Agents: <https://www.microsoft.com/en/customers/story/26118-siag-dynamics-365-customer-service>
- [9-47] Salesforce Agentforce Customer Success Story Collection (Including 1-800Accountant, Asymbi, and More): <https://www.salesforce.com/news/stories/agentforce-customer-success-stories/>
- [9-48] Sandvik Manufacturing Copilot: Turning Industrial Knowledge into Sales/Service Capability: <https://www.microsoft.com/en/customers/story/22965-sandvik-azure-open-ai-service/>
- [9-49] Spotify Turns AI Personalization into a Growth Engine: First Full Year of Profit in 2024 (~€1.1 Billion), with 281 Million Paying Subscribers: <https://www.g-co.agency/insights/spotify-case-study-personalization-strategy-as-a-growth-engine>
- [9-50] State of AI in Fintech Customer Acquisition 2024: <https://www.cbinsights.com/research/report/ai-fintech-customer-acquisition-2024/>
- [9-51] The AI Playbook for Financial Services (WEF + Accenture, June 2026 Report): https://reports.weforum.org/docs/WEF_The_AI_Playbook_for_Financial_Services_2026.pdf
- [9-52] Virgin Money Redi: From Reactive Customer Service to Proactive Customer Journeys: <https://www.microsoft.com/en/customers/story/24471-virgin-money-microsoft-365-copilot>
- [9-53] Postmortem of a Failed Hospital AI Follow-Up Project (The Affiliated Hospital of Guilin Medical University): <https://www.woshipm.com/pd/6086013.html>
- [9-54] Sany Heavy Industry's CRM Localization: AI Sales-Forecast Model + Equipment-Customer-Service Digital Twin Lifts Lead Conversion to 28%: <https://www.fxiaokey.com/crm/information-69990.html>

- [9-55] Listed Insurers Fire the Opening Shot in the AI Escalation War: Ping An Underwrites 93% of Policies in Seconds; New China Life's AI Support Equals 600 Human Agents: <https://stcn.com/article/detail/1682984.html>
- [9-56] Ping An: AI-Empowered Full-Lifecycle Customer Operations and Intelligent Service Practice: https://www.pingan.com.cn/news/detail/20240415_cs_ai.html
- [9-57] A Mid-Size Law Firm's AI Intake and Client-Lifecycle System: \$420,000 in New Revenue in the First 90 Days; Lead Response Time Cut from 36-48 Hours to Under 90 Minutes: <https://cappedoutlabs.com/case-studies/law-firm-ai-intake>
- [9-58] AI + Cross-Border E-Commerce: Top 10 Cross-Border E-Commerce Cases — Walmart and Shopify's Agentic Commerce Positioning Battle: <https://weibo.com/ttarticle/p/show?id=2309405292065516028187>
- [9-59] Personal-Injury Law Firm Website Visitor Identification: 1,500 Anonymous Visitors Identified Per Month, New Engagements +0.86% in a Single Quarter: <https://www.fullthrottle.ai/success-story/personal-injury-law-firm/>
- [9-60] From "Click Reach" to "Long-Term Retention": Mastering B2B Customer Acquisition, Retention, and Loyalty Cultivation: <https://www.sap.cn/blogs/top-three-b2b-customer-acquisition-strategies>
- [9-61] Zuoyebang Learning Machine Tops August 2025 Online Sales with a 38.6% Share: Driven by the Self-Developed Galaxy LLM + AI Super Teacher: <https://finance.sina.com.cn/jjxw/2025-09-20/doc-infrcarp2204317.shtml>
- [9-62] AI Customer Service for Outage Scenarios: Florida Electric Cooperative SECO Energy Cuts Per-Call Cost 66% with Intelligent Virtual Agents; 32% of Calls Fully Automated: <https://capacity.com/blog/outage-management/>
- [9-63] ABM Media Pilot for a Global Manufacturer: 50% Target-Account Reach, 100% of Leads Are MQLs, 80 Key Accounts Activated: <https://www.icrossing.com/case-studies/abm-media-strategy-manufacturing-parts-services>
- [9-64] Beijing E-Town Launches the City's First Government-LLM Platform "Yizhi": Smart Support Agent "Xiaoyi" Covers 2 Million Knowledge Items and 97 High-Frequency Services: http://banshi.beijing.gov.cn/tzgg/202407/t20240702_428169.html
- [9-65] Xiamen Cardiovascular Hospital (Xiaxin): Rebuilding Diagnosis and Treatment with "Business Agents" — Outpatient Precise Triage, Medical-Record QC, and a Closed Risk-Warning Loop: <https://www.chima.org.cn/Html/News/Articles/17465.html>

- [9-66] Key-Account Management, Digital Assets, and AI Applications: <https://www.woshipm.com/share/6050158.html>
- [9-67] LLMs Empower Lalamove's Customer Service: Intelligent Practice in Intent Recognition, Emotional Comfort, and Vehicle-Type Recommendations: <https://juejin.cn/post/7486089532283879424>
- [9-68] Perfect Diary's Private-Domain Community Operations Deconstructed: Parcel Cards + Personal Accounts + Community Funnels, with SCRM Automated Tagging and Segmentation: <http://www.ikscrm.com/school/syll/1493.html>
- [9-69] Now That Customers Ask AI First, How Should B2B Companies Change Content-Led Acquisition?: <https://segmentfault.com/a/1190000048010907>
- [9-70] Xiaoman Technology's OKKI: Full-Scenario AI for Foreign-Trade CRM — 75% AI Adoption, No. 1 with 14.7% Share in the Tool Category: <https://mp.m.ofweek.com/ai/a456714443567>
- [9-71] Watsons Private Domain Fully Deconstructed: 65 Million Members, 4,100+ Stores, 5,000+ Communities, Online Digital Business Revenue Exceeds RMB 500 Million: <https://www.woshipm.com/operate/5452718.html>
- [9-72] Guangdong Government Service Online Launches an AI Smart Service Model for Government Affairs: 400,000 Sessions in Two Months, 95%+ Request-Recognition Rate: http://psc.zs.gov.cn/gkmlpt/content/2/2388/post_2388833.html
- [9-73] Guangdong Power Grid's "Five-in-One" Intelligent Customer-Service System: AI Resolved 13.72 Million Cases in the First 11 Months of 2023, Diverting 70% of Manual Work: <https://zhuanlan.zhihu.com/p/694784136>
- [9-74] WeBank's AI Computing Power Grows 3.5x, Serving More Than 440 Million Individual Customers: <https://www.webank.com/mediaCenter/mtbd/3141>
- [9-75] KuaiShang's "Legal LLM + Omnichannel Aggregation" AI Support Cracks Law-Firm Acquisition: AI Handles 90% of Nighttime Inquiries, Acquisition Cost Down 35%: <https://www.kuaishang.cn/jiqiren/23679.html>
- [9-76] China Merchants Bank × Tongyi Foundation Models: Intelligent Investment-Research Assistant "CMB Think-Tank AI Xiaoyan" Wins the 2024 AI Best-Practice Award: <https://finance.sina.com.cn/roll/2025-01-15/doc-ineezhcc8057421.shtml>

- [9-87] A U.S. Nonprofit Health System Uses AI + Salesforce Health Cloud to Transform Patient Engagement: <https://www.pwc.com/us/en/library/case-studies/ai-healthcare-engagement-transformation.html>
- [9-88] Tencent’s Healthcare LLM Deployed at ~1,300 Hospitals: AI Checkup Add-Ons Lift Average Order Value +28% Across Triage/Pre-Consultation/Support Scenarios: <https://developer.cloud.tencent.com/article/2677992>
- [9-89] Mango TV Members Surpass 75.6 Million: Hunan Broadcasting Group’s “Mango LLM” Incubates 80+ Agents Across 30+ Shows: <https://www.chinaz.com/ainews/27182.shtml>
- [9-90] Suzhou 12345 Hotline’s “Intelligent Call-Handling LLM”: Real-Time Transcription + Smart Form-Filling + Knowledge Recommendations + Automatic Dispatch: <https://www.suzhou.gov.cn/szsrnzf/szyw/202503/12c05d46f7fe4eddaa56ad0a1274b957.shtml>
- [9-91] Cainiao’s Intelligent AI Support Serves 430,000 Sessions a Day: Freight Outbound Calls, Rider Dispatch, and Cross-Border Multilingual Service: <https://www.10100.com/article/57965212>
- [9-92] BlueFocus’s “Self-Revolution”: AI-Driven Revenue of RMB 3.725 Billion in 2025, Token Consumption Past 1 Trillion, Pivoting to “Selling Systems” for Brand Clients: <https://www.36kr.com/p/3768876765496065>
- [9-93] Resolution Rate from 68% to 86%! Shenzhen Tong Uses AI to Field Tens of Millions of Residents’ “Soul-Searching” Questions: <https://grow.163.com/cms/vEpaeuYg.html>
- [9-94] Banks Build an Efficient “Human-Machine Collaboration” Service Ecosystem: <https://news.10jqka.com.cn/20251127/c672767667.shtml>
- [9-95] Retail AI Readiness: AI-Readiness Rankings of the Top 20 Retailers by Market Cap (CB Insights): <https://www.cbinsights.com/research/ai-readiness-index-for-retail>
- [9-96] ZKH Industrial’s Agent Matrix Named a Shanghai “AI+” Excellent Case: AI Material Manager Lifts Procurement-Management Efficiency by 30%: <https://finance.sina.com.cn/stock/bxjj/2025-07-15/doc-inffpwtm5528562.shtml>
- [9-97] The AI Trump Cards of Nine Companies — SF Express, JD Logistics, Cainiao, ZTO, YTO, Shentong and More — All in One Look: <https://news.qq.com/rain/a/20260720A09GET00>

- [9-98] How 16 Banks Are Building AI Agents: "Scenarios + Agents" Becomes the Next Battleground (Shaqiu Community, Chinese: 沙丘社区): <https://www.shaqiu.cn/article/ZBnaY03KVzp7>
- [9-99] 9 Agentic AI Use Cases Banks Are Scaling in 2026 (Backbase): <https://www.backbase.com/blog/agentic-ai-use-cases-in-banking>
- [9-100] 2026: The Year of Agentic AI, and a New Era for Finance (Lloyds Banking Group): <https://www.lloydsbankinggroup.com/insights/2026-the-year-of-agentic-ai-and-a-new-era-for-finance.html>
- [9-101] A Compliance Guide for AI in Banking and Insurance: Interpreting Yinfa [2026] No. 8: <https://zhuanlan.zhihu.com/p/2060758478633931195>
- [9-102] Manufacturing's 2026 Mandate: From AI Pilot to Agentic Profit (Dataiku): <https://www.dataiku.com/blog/manufacturing-ai-trends-2026>
- [9-103] 2026 Digital Government Citizen Survey: Citizens Want Governments to Accelerate AI Adoption (BCG): <https://www.bcg.com/publications/2026/why-citizens-want-governments-to-accelerate-ai-adoption>
- [9-104] 36Kr: Futong Tianxia Wires Agents Into the Entire Foreign-Trade Chain — 24/7 Multilingual Replies, Automated Background Checks, Lead Scoring, and Quote Generation; 60,000 Foreign-Trade Firms Onboarded: <https://www.36kr.com/p/3844402952391177>
- [9-105] Becker's Hospital Review: 7,000 Registrants, 350 Appointments — 3 Health Systems Prove Digital Marketing ROI: <https://www.beckershospitalreview.com/digital-marketing/7000-registrants-350-appointments-3-health-systems-prove-digital-marketing-roi/>
- [9-106] The Paper (interview): Ping An's Fu Xin — AI Applications Are Judged on Scale, Entry Point, Scenario, and Substitutability, and Every Project Must Cost the Business Volume It Will Generate at Approval: https://www.thepaper.cn/newsDetail_forward_33846502
- [9-107] Securities Times: The Agent Wave Sweeps the Securities Industry — Brokerages' AI Transformation Moves from Point Improvements to System Rebuilding: <https://wap.eastmoney.com/a/202609283884544867.html>
- [9-108] Freshfields (firm website): Tech, Data and AI — The Digital Frontier: <https://www.freshfields.com/en/our-thinking/campaigns/tech-data-and-ai-the-digital-frontier>

- [9-109] Warner Music Group (official news): Warner Music Group Opens a Technology and Engineering Hub in Bengaluru: <https://www.wmg.com/news/warner-music-group-signs-ai-licensing-deal-with-music-technology-company-klay>

Chapter 10 Sources

- [10-01] AI-Generated and Synthetic Content Must Carry Labels from September 1: China's New AI Content-Compliance Rules Take Effect: http://www.ce.cn/xwzx/gnsz/gdxw/202509/t20250901_2460291.shtml
- [10-02] AI Outbound Calls Dial a Thousand Nuisance Calls a Day — Where Is the “Compliance Boundary”? http://news.china.com.cn/2026-07/28/content_118621717.shtml
- [10-03] AI-Powered Deception: A Deeper Dimension of Dark Design Patterns in Conversational AI Tools and Platforms: <https://cdt.org/insights/ai-powered-deception-a-deeper-dimension-of-dark-design-patterns-in-conversational-ai-tools-and-platforms/>
- [10-04] Air Canada Held Responsible for Chatbot's Hallucinations: <https://www.aibusiness.com/nlp/air-canada-held-responsible-for-chatbot-s-hallucinations->
- [10-05] Can Companies Insure Against AI's Growing Risks? (Tufts Now, 2026-05-12): <https://now.tufts.edu/2026/05/12/can-companies-insure-against-ais-growing-risks>
- [10-06] The DPD System-Update Incident: An AI Support Chatbot Swears at a Customer and Writes a Poem Slamming Its Own Company (DPD error caused chatbot to swear at customer): <https://www.bbc.com/news/technology-68025677>
- [10-07] EU AI Act Article 50 Transparency Obligations Take Effect on 2026-08-02: Customer-Facing AI Interactions Must “Reveal AI Identity”; Human Review Can Serve as a Liability Exemption: <https://artificialintelligenceact.eu/transparency-rules-article-50/>
- [10-08] FCC Passes Declaratory Ruling Unanimously: AI-Generated Robocalls Are Illegal: <https://www.fcc.gov/document/fcc-makes-ai-generated-voices-robocalls-illegal>

- [10-09] FTC’s “Operation AI Comply”: An Enforcement Storm of Five Cases Against Exaggerated and Fraudulent AI Marketing Claims: <https://www.ftc.gov/news-events/news/press-releases/2024/09/ftc-announces-crackdown-deceptive-ai-claims-schemes>
- [10-10] FTC Impersonation Rule: Impersonating Government and Businesses Is Prohibited (Including AI-Generated Impersonation): <https://www.ftc.gov/legal-library/browse/rules/impersonation-government-businesses-rule>
- [10-11] FTC Penalizes “Robot Lawyer” DoNotPay: Unsubstantiated AI Customer-Acquisition Promises Are Unlawful: <https://www.ftc.gov/legal-library/browse/cases-proceedings/donotpay>
- [10-12] HIPAA Enforcement Turns to the “Browser”: In 2024, the OCR Issued \$9.9 Million in Fines over Website Tracking Tools on Medical Sites — AI Support Chatbots Become a New PHI-Leak Vector: <https://feroot.com/blog/hipaa-violation-penalties-website-tracking/>
- [10-13] Moffatt v. Air Canada, 2024 BCCRT 149 (Judgment in the Moffatt v. Air Canada Case): <https://decisions.civilresolutionbc.ca/crt/crtd/en/item/525448/index.do>
- [10-14] PwC: With the UK FCA Adding No New AI Rules, Consumer Duty Will Be the Main Compliance Lens for Scrutinizing “Customer-Facing AI”: <https://www.pwc.co.uk/industries/financial-services/understanding-regulatory-developments/scaling-customer-facing-ai-unlocking-better-outcomes-and-consumer-duty-compliance.html>
- [10-15] SEC’s First “AI Washing” Enforcement: Two Investment Advisers Fined \$400,000 for False AI Marketing Claims: <https://www.sec.gov/newsroom/press-releases/2024-36>
- [10-16] The Florida Bar Issues the Nation’s First AI-Use Guidelines: One-Third of Members Have Never Used AI; Client-Data Confidentiality and Disclosure Statements Are the Red Lines: <https://www.floridabar.org/the-florida-bar-news/new-florida-bar-guide-helps-lawyers-incorporate-ai-responsibly-in-their-work/>
- [10-17] First Generative-AI “Hallucination” Tort Case Ruled in China: AI’s “Compensation Promise” Is Not the Company’s Expression of Intent; the Service Provider Is Exempted by Prominent Disclaimers (Hangzhou Internet Court): <http://society.people.com.cn/n1/2026/0126/c1008-40652440.html>

- [10-18] Air Canada ordered to pay customer who was misled by airline's chatbot: <https://www.theguardian.com/world/2024/feb/16/air-canada-chatbot-lawsuit>
- [10-19] Italy's Garante Fines AI Emotional-Companion Chatbot Replika €5 Million: No Legal Basis for Processing + No Age Verification: <https://www.reuters.com/sustainability/boards-policy-regulation/italys-data-watchdog-fines-ai-company-replikas-developer-56-million-2025-05-19/>
- [10-20] Munich Re's HSB Launches AI Liability Insurance for Small and Medium Businesses: "AI Lawsuits" Become an Insurable Risk: <https://www.munichre.com/hsb/en/press-and-publications/press-releases/2026/2026-03-18-introducing-ai-liability-insurance-for-small-businesses.html>
- [10-21] 12 U.S. States Have Enacted Laws Regulating Consumer-Facing AI Chatbots: AI Identity Disclosure, Minor Protection, and Accountability for "Machines Masquerading as Humans" Form the Compliance Baseline: <https://www.koleyjessen.com/insights/publications/ai-chatbot-compliance-what-your-business-needs-to-know>
- [10-22] 7 Frontier AI Models Ran Real Businesses for 72 Hours: \$12,431 in Fake Invoices, 2,797 Spam Emails, \$0 Revenue (Bottleneck Labs): <https://www.bottlenecklabs.com/blog/benchmarking-7-autonomous-businesses>
- [10-23] OpenAI's Rogue Agents Reached 10+ More Websites: University Wikis, a Chemistry Wiki, and Text-Sharing Sites as Message Boards (Fortune): <https://fortune.com/2026/09/09/openai-rogue-ai-agents-reached-12-more-websites/>
- [10-24] An OpenAI Agent Swarm Attacked RubyGems Without Disclosure: 2,000+ Malicious Packages, a Novel Flaw Exploited to Steal API Keys: <https://www.rubyhack.ai/>
- [10-25] The Register (reporting on a RAND report and on W. R. Berkley and Verisk/ISO policy language): AI Risks Make Some Insurers Wary of Corporate Liability: <https://www.theregister.com/ai-and-ml/2026/09/18/ai-risks-make-some-insurers-wary-of-corporate-liability/5297347>
- [10-26] Reuters: Lawyer for State Farm Fined Over AI Hallucination in Los Angeles Lawsuit: <https://www.reuters.com/legal/litigation/lawyer-state-farm-fined-over-ai-hallucination-los-angeles-lawsuit-2026-09-14/>
- [10-27] Xinhuanet Zhejiang (covering two Hangzhou rulings): AI Causes Trouble — Who Pays, and How to Govern It: <http://www.zj.xinhuanet.com/20260414/f07218aaacdc46a1b2f8d0984a264601/c.html>

- [10-28] People's Daily (Satire and Humor supplement): Beware of "AI Hallucination": http://paper.people.com.cn/fcyym/pc/content/202606/05/content_30164935.html
- [10-29] Sina Finance: Why "AI Outbound Calls" Harassment Cannot Be Curbed — One Valid Customer Record Sells for RMB 100, With a Return Above 1:5: <https://finance.sina.com.cn/jjxw/2025-05-13/doc-inewmaka2042042.shtml>
- [10-30] Beijing Daily (carrying China Media Group's Voice of China): AI Outbound Calling Dials 800–1,500 Numbers a Day at RMB 0.12–0.16 per Minute — the Compliance Boundary Is Whether the User Gave Clear Prior Consent: <https://xinwen.bjd.com.cn/content/s6a68080fe4b0e45f3fd54adc.html>
- [10-31] MIIT and Other Agencies: Work Plan for the Special Campaign to Comprehensively Rectify Nuisance Calls (full text republished by Security Inside Reference): <https://www.secrss.com/articles/6129>
- [10-32] Xinhuanet (carrying the State Administration for Market Regulation): SAMR Publishes Ten Representative Illegal-Advertising Cases — 44,521 Advertising Violations Handled Nationwide in 2025: <https://www.news.cn/fortune/20260131/b7a71709bc5f4ec2bf6dff05ce08486d/c.html>
- [10-33] 21st Century Business Herald: The March 15 Program Exposes the AI "Poisoning" Gray Supply Chain, Revealing High-Risk Algorithmic Gaps Behind Large Models: <https://www.21jingji.com/article/20260316/herald/8cf9afdb3bc8ba06b10b2f89aef3bc17.html>
- [10-34] Xinhua: The Hangzhou Internet Court Rules That "Inaccurate AI Information" Is Not Necessarily Infringing — Liability Turns on Whether the Developer Met Its Obligations: <https://www.news.cn/legal/20260330/670d3456cc1a48be8dd9df42ad13e472/c.html>
- [10-35] BBC Media Centre (2025-02-11): BBC Research Shows Issues With Answers From Artificial Intelligence Assistants: <https://www.bbc.com/mediacentre/2025/bbc-research-shows-issues-with-answers-from-artificial-intelligence-assistants>
- [10-36] Jiefang Daily / Shangguan News (4,308 valid responses): Shanghai Consumer Council's 2026 "618" Online Shopping Experience Survey — AI Matching Accuracy at Only 16.06% and 38.79% Reporting That Higher-Priced Items Are Pushed First: <https://www.jfdaily.com/wx/detail.do?id=1134996>

- [10-37] Volcano Engine Developer Community: Three Numbers on Brand Hallucination — 23% of AI Answers Contain Factual Errors and Brand-Related Content Errs at 31%: <https://developer.volcengine.com/articles/7621130200160682038>
- [10-38] Growth Hacking AI Weekly EP.71: "Mass-Automated Heartbreak Farming" — Dating Apps Put Large Models Into the Matching Pool to Impersonate Real People: https://mp.weixin.qq.com/s?__biz=Mzg5ODU4OTU3MQ==&mid=2247502366&idx=1&sn=07d845ffa276a6eb68edd0a82cf0
- [10-39] Owlsh (covering Moffatt v. Air Canada and OLG Hamm): Who's Liable When Your AI Support Agent Gives a Wrong Answer?: <https://owlsh.bot/blog/ai-chatbot-liability>
- [10-40] ePrivacy Blog (analysis of the OLG Hamm ruling and EU AI Act Article 50): Landmark Higher Regional Court Ruling on the Attribution of AI Statements: <https://blog.eprivacy.eu?p=2805>
- [10-41] Xinhua (via Science and Technology Daily): China Consumers Association Advisory — Don't Let AI Services Mislead Consumption Decisions: <https://www.news.cn/tech/20260902/77bfdb038b5040a8bec54802df6e3577/c.html>
- [10-42] China Consumers Association (official site): Compilation of Complaint Cases (including AI-customer-service cases): <https://cca.org.cn/single/11186/69866.html>
- [10-43] Xinhua (via Legal Daily): Some Users Misled by AI "Hallucinations" — Experts Urge Clarifying the Multiple Liability Boundaries: <https://app.xinhuanet.com/news/article.html?articleId=20260629ff39cf20d4b4482e8b1e0e9da60d2bb3>
- [10-44] Sina Finance (via Southern Metropolis Daily): Yuanbao, Doubao, and Qwen Take Their Agent Features Offline as New AI-Anthropomorphism Rules Take Effect: <https://finance.sina.com.cn/stock/t/2026-07-04/doc-inifscxv2472777.shtml>
- [10-45] ByeRisk Blog: New Rules on Labeling AI-Generated Content — How Douyin, Xiaohongshu, and WeChat Channels Differ in Implementation: <https://www.byerisk.com/blog/aigc-content-labeling-3-platform-comparison>

Chapter 11 Sources

- [11-01] 2025 Edelman Trust Barometer Special Report on Brands: Brand Trust Shifts “From We to Me”: <https://www.edelman.com/trust/2025/trust-barometer/special-report-brands>
- [11-02] Analysis of 680 Million AI Citations: ChatGPT, Google AI Overviews, and Perplexity Cite in Completely Different Ways — “AI Search”: <https://www.leapd.ai/blog/ai-visibility/how-chatgpt-google-ai-overviews-and-perplexity-source-information-in-2026>
- [11-03] AI Brand Visibility Statistics (2026): How Often Brands Are Cited in ChatGPT/Perplexity/Claude/Gemini (Boring Marketing): <https://boringmarketing.com/ai-visibility-statistics>
- [11-04] AIGOS All-Domain Intelligent Growth Operating System: From SEO/SEM/ASO/DSO to AEO/GEO/GEM/PDE/Agent and Agentic C...: https://blog.csdn.net/weixin_41860745/article/details/162148098
- [11-05] Agentic Commerce: A Guide to AI Shopping Agents — Redefining Retail Economics (Deloitte Report): <https://www.deloitte.com/us/en/Industries/consumer/articles/agentic-commerce-ai-shopping-agents-guide.html>
- [11-06] The Agentic Commerce Opportunity: How AI Agents Are Ushering In a New Era for Consumers and Merchants (McKinsey Report): <https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-agentic-commerce-opportunity-how-ai-agents-are-ushering-in-a-new-era-for-consumers-and-merchants>
- [11-07] Ahrefs Studies 75,000 Brands: What Determines Whether a Brand Is Recommended in ChatGPT / AI Mode / AI Overviews: <https://ahrefs.com/blog/ai-brand-visibility-correlations/>
- [11-08] BCG Global Survey: More Than 60% of Consumers Highly Trust GenAI Results — Consumers Believe “AI Can Help Me Buy Better”: <https://www.bcg.com/publications/2026/consumers-trust-ai-to-buy-better-brands-must-adapt>
- [11-09] Buy it in ChatGPT: Instant Checkout and the Agentic Commerce Protocol: <https://openai.com/index/buy-it-in-chatgpt/>
- [11-10] Forrester Tracks B2C Agentic Payments: Protocol Wars Heat Up, OpenAI Instant Checkout Pulls Back, Agents...: <https://www.forrester.com/blogs/agentic-payments-in-b2c-commerce-where-we-are-now/>

- [11-11] GEO Case Study: Go Fish Digital Does Generative Engine Optimization with Four Levers — AI-Channel Inquiries Triple, Converting 25× Better Than Traditional Search: <https://gofishdigital.com/blog/generative-engine-optimization-geo-case-study-driving-leads/>
- [11-12] Gartner Predicts: Traditional Search-Engine Traffic Will Drop 25% by 2026 as AI Chatbots Take Share: <https://www.gartner.com/en/newsroom/press-releases/2024-02-19-gartner-predicts-search-engine-volume-will-drop-25-percent-by-2026-due-to-ai-chatbots-and-other-virtual-agents>
- [11-13] Gartner: AI Agents Will Intermediate More Than \$15 Trillion in B2B Purchases by 2028; 90% of B2B Buying Will Be Handled by Agents: <https://www.digitalcommerce360.com/2025/11/28/gartner-ai-agents-15-trillion-in-b2b-purchases-by-2028/>
- [11-14] IBM × NRF Global Study “Own the Agentic Commerce Experience”: 45% of Consumers Already Use AI in Their Shopping Journeys...: <https://newsroom.ibm.com/2026-01-07-ibm-nrf-study-brands-and-retailers-navigate-a-new-reality-as-ai-shapes-consumer-decisions-before-shopping-begins>
- [11-15] Metrigy Study: 85% of Consumers Prefer Interacting with Humans Rather Than AI Agents for Customer Service: <https://www.metrigy.com/press-release-metrigy-study-85-of-consumers-prefer-interacting-with-humans-vs-ai-agents-for-customer-service/>
- [11-16] OpenAI Operator Launches: A Signal That the Agentic Commerce Era Has Arrived (Crunchbase News Guest Commentary): <https://news.crunchbase.com/ai/openai-operator-agentic-era-ecommerce-friend-bain/>
- [11-17] Salesforce 2025 Holiday Shopping Data: AI and Agents Drive \$262 Billion in Sales — Evidence for the First Year of “Agentic Shopping”: <https://www.salesforce.com/news/stories/2025-holiday-shopping-data/>
- [11-18] Stripe Launches Agentic Commerce Suite: Low-Code, Single Integration Lets Merchants Put Their Products in Front of Every AI Agent: <https://stripe.com/newsroom/news/agentic-commerce-suite>
- [11-19] SurveyMonkey 2026 Customer-Service Statistics: 79% of Americans Strongly Prefer Human Agents; 42% Would Pay Extra for a Human Channel: <https://www.surveymonkey.com/curiosity/customer-service-statistics/>
- [11-20] What Is Agentic Commerce? How AI Agents Are Reshaping Online Shopping: <https://fin.ai/learn/what-is-agentic-commerce>

- [11-21] The “Digital Trust 2026” Report: 52% of Consumers Would Pay 7% More on Average to Brands That “Win Trust with AI” — the Trust Premium Is Quantified for the First Time: <https://usercentrics.com/resources/state-of-digital-trust-report-2026/>
- [11-22] The Machine-Customer Economy Has Arrived: The Structural Break in Customer Experience — “Make Your Business Readable by AI Agents”: <https://www.cmswire.com/customer-experience/machine-customers-the-structural-break-in-customer-experience/>
- [11-23] Generative Engine Optimization (GEO) Statistics 2026: 60+ Data Points on AI Citations: <https://www.omnibound.ai/blog/generative-engine-optimization-statistics>
- [11-24] Survey: In 10 Years “Human Service” Will Become a Luxury — the Human Touch Is Commanding a Premium (CX Dive / Medallia): <https://www.customerexperiencedive.com/news/human-driven-customer-service-luxury-premium/816501/>
- [11-25] How Retailers Stay Visible to AI Shopping Agents: Competition Shifts from Clicks to “Being Chosen” (ML6, 2026-04): <https://www.ml6.eu/en/blog/how-do-retailers-stay-visible-with-ai-shopping-agents-choosing-products>
- [11-26] AI and Agentic Commerce Statistics 2026 (commercetools): <https://commercetools.com/blog/agentic-commerce-stats-enterprise-guide>
- [11-27] 2026 AI Agent Commerce Industry Report (Checkout.com): <https://www.checkout.com/zh-cn/guides-and-reports/agentic-commerce-2026>
- [11-28] Own the Agentic Commerce Experience (IBM Institute for Business Value): <https://www.ibm.com/thought-leadership/institute-business-value/en-us/report/agentic-commerce>
- [11-29] Ken Huang: Google and Shopify's UCP Just Won Agentic Commerce: <https://kenhuangus.substack.com/p/googles-ucp-just-won-agentic-commerce>
- [11-30] Opascope: AI Shopping Assistant Guide 2026 — Agentic Commerce Protocols (ACP vs. UCP): <https://opascope.com/insights/ai-shopping-assistant-guide-2026-agentic-commerce-protocols>
- [11-31] Amazon Seller Central (official forum notice): Business Solutions Agreement Adds an Agent Policy: <https://sellercentral.amazon.com/seller-forums/discussions/t/84e3f6b1-42f7-4cf3-a189-a5cc8d78d838>

- [11-32] Profasee: Amazon AI Agent Policy — What Sellers Need to Know (SP-API registration, 12-month audit trail, human-authorization thresholds): <https://profasee.com/blog/amazon-ai-agent-policy-what-sellers-need-to-know>
- [11-33] STCN (Securities Times): Doubao Responds to Its "12% Hotel Commission" — the AI Entry Point Becomes an Independent Transaction Channel: <https://www.stcn.com/article/detail/4068900.html>
- [11-34] Xinhuanet: Alipay Announces AI Agent Payments Have Passed 300 Million Transactions and Launches a Full-Stack AI-Native Payment System: <https://www.news.cn/tech/20260526/5169cdb4192548528203168d3ee109fec.html>
- [11-35] Shanghai Observer (Jiefang Daily): China's First Multi-Agent Cross-Device Interconnection Protocol Released — Agent Commercialization Accelerates as Merchant Payment Fees Are Discounted 40%: <https://www.shobserver.com/wx/detail.do?id=1162176>
- [11-36] Ant International: Agentic Mobile Protocol (AMP) Onboards 17 Institutions at Launch, Covering 1.5 Billion Consumers: <https://www.ant-intl.com/cn/news/detail?id=jsx9m407&redirect=news>
- [11-37] Fortune: Amazon Says Its AI Shopping Assistant Rufus Is on Pace to Pull in an Extra \$10 Billion in Sales: <https://fortune.com/2025/11/02/amazon-rufus-ai-shopping-assistant-chatbot-10-billion-sales-monetization>
- [11-38] Digital Commerce 360 (citing Adobe Analytics): Adobe — AI Referral Traffic Keeps Doubling and Converts Better Than Other Channels: <https://www.digitalcommerce360.com/2026/08/19/adobe-ai-referral-traffic-data-july-2026>
- [11-39] Demand Local: 23 AI Agent and Autonomous Shopping Statistics for Retailers in 2026: <https://www.demandlocal.com/blog/ai-agent-autonomous-shopping-statistics>
- [11-40] Braze: Agentic Commerce — A Complete Guide for Brands (citing the Retail Customer Engagement Review and IBM Institute for Business Value): <https://www.braze.com/resources/articles/agentic-commerce>
- [11-41] Mastercard (press release, 2026-06-10): Mastercard Launches Agent Pay for Machines: <https://www.mastercard.com/us/en/news-and-trends/press/2026/june/mastercard-launches-agent-pay-for-machines.html>

- [11-42] Visa (official product page): Visa Intelligent Commerce — A Trust Layer for AI Agent Shopping: <https://www.visa.com/en-us/solutions/intelligent-commerce>
- [11-43] Gartner (2026-05-27): Gartner Survey Finds Consumers Want AI Shopping Help but Not AI Purchase Decisions: <https://www.gartner.com/en/newsroom/press-releases/2026-05-27-gartner-survey-finds-consumers-want-ai-shopping-help-but-not-ai-purchase-decisions>
- [11-44] Checkout.com ("Agentic Commerce 2026"): Consumer Demand for AI Shopping Is Forming Fast, but Trust for Agentic Commerce Is Still Catching Up: <https://www.checkout.com/newsroom/consumer-demand-for-ai-shopping-is-forming-fast-but-trust-for-agentic-commerce-is-still-catching-up>
- [11-45] Mastercard (press release): Mastercard Report Predicts That One in 10 People Will Routinely Use AI Agents to Shop and Pay by 2030: <https://www.mastercard.com/news/europe/en/newsroom/press-releases/en/2026/mastercard-report-predicts-that-one-in-10-people-will-routinely-use-ai-agents-to-shop-and-pay-by-2030/>
- [11-46] Deloitte (2026-06-26): Agentic Commerce — The Future of B2B Commerce: <https://www.deloitte.com/us/en/what-we-do/capabilities/applied-artificial-intelligence/articles/b2b-agentic-commerce.html>
- [11-47] The Keyword (2026-06-29): Perplexity Shopping GMV Run Rate Estimated Near \$2 Billion: <https://www.thekeyword.co/news/perplexity-commerce-gmv-run-rate>
- [11-48] Digital Applied: Agentic Commerce Standards — UCP vs ACP vs AP2 in 2026: <https://www.digitalapplied.com/blog/agentic-commerce-standards-ucp-acp-ap2-2026-merchant-guide>
- [11-49] Airwallex: Explaining Agentic Commerce Protocols — What Merchants Most Need to Know: <https://www.airwallex.com/cn/blog/understanding-agentic-commerce-protocols>
- [11-50] Hexagon: The Agentic Commerce Protocol Stack — UCP, ACP, AP2, MCP, A2A Explained: <https://joinhexagon.com/blogs/the-agentic-commerce-protocol-stack-ucp-acp-ap2-mcp-a2a-expl-mmi9c03u-hb1k>
- [11-51] eesel AI: ChatGPT Instant Checkout 2026 — What Changed for Stores: <https://www.eesel.ai/blog/chatgpt-instant-checkout-2026>

- [11-52] GoDataFeed: How AI Shopping Surfaces Read Your Product Feed — ChatGPT, Perplexity, Google AI Mode & Copilot: <https://www.godatafeed.com/blog/ai-shopping-product-feed-optimization>
- [11-53] Shopify (official blog): Perplexity Shopping — How to Optimize Your Store for AI (2026): <https://www.shopify.com/blog/perplexity-shopping>
- [11-54] ACI Worldwide (YouGov survey): Six in Ten UK Consumers Would Stop Using an AI Shopping Agent After One Mistake: <https://investor.aciworldwide.com/news-releases/news-release-details/six-ten-uk-consumers-would-stop-using-ai-shopping-agent-after>
- [11-55] Klaviyo (2026 AI Consumer Trends Report): Consumer Trust in AI — What Brands Need to Know in 2026: <https://www.klaviyo.com/solutions/ai/consumer-trust-in-ai>
- [11-56] Alchemer (2026 Retail Report): Retail AI Adoption Outpaces Consumer Trust: <https://www.alchemer.com/resources/benchmark-report/2026-retail-report-ai>
- [11-57] paperclipped: Amazon Rufus AI Shopping Agent — \$12B Sales, Buy for Me & Auto Buy: <https://www.paperclipped.de/en/blog/amazon-rufus-ai-shopping-agent>

Narrative Sources (Stories Used Across Chapters)

- [12-01] “AI Is a People Problem, Not a Technology Problem”: Zapier CEO Wade Foster on the Post-GPT-4 “Code Red,” All-Company Hackathons, and More: <https://betweentwocoos.com/wade-foster-why-ai-adoption-is-a-people-problem/>
- [12-02] “The Signals Were Always There — Most CS Teams Just Weren’t Looking”: How Seven Customer-Success Executives at Pulse Moved Their Biggest Rocks with AI — from Rockwell...: <https://unchurned.gainsight.com/p/inside-pulse-how-7-cs-leaders-are>
- [12-03] “I Tore Down the Whole Company”: IgniteTech CEO Eric Vaughan Replaced Nearly 80% of Employees in a Year, Facing Resistance and Even Sabotage — Two Years Later He Says, “I’d Do It Again”: <https://fortune.com/2025/08/17/ceo-laid-off-80-percent-workforce-ai-sabotage/>

- [12-04] “I Need Fewer Heads”: Salesforce CEO Marc Benioff Admits Cutting the Customer-Service Team Down from 9,000 Once AI Agents Went Live: <https://fortune.com/2025/09/02/salesforce-ceo-billionaire-marc-benioff-ai-agents-jobs-layoffs-customer-service-sales/>
- [12-05] Artisan’s 8 Months from 0 to \$1.3M ARR: How 23-Year-Old Founder Jaspar Carmichael-Jack Won His First Customers: <https://www.pmf.show/blog/artisan-jaspar-carmichael-jack-0-to-1-3m-arr-ai-bdr-ava>
- [12-06] Klarna’s U-Turn: From “AI Can Replace 700 Support Agents” to the CEO Admitting “We Went Too Far” — Announcing at Stockholm HQ to Bloomberg That It Was Re-Hiring Humans: <https://www.maginitive.com/article/klarna-dials-back-its-ai-customer-service-strategy-now-its-hiring-humans-again/>
- [12-07] Confessions of an AI SDR Skeptic: ROI DNA’s Head of Marketing Gina Inks on Her Decision Process — from Resistance to Putting Piper at the Core of Growth (Qualified): <https://www.qualified.com/plus/articles/confessions-of-an-ai-sdr-skeptic>
- [12-08] A Tech Company Admits It Was Once an AI Skeptic: Intercom Principal Engineer Brian Scanlan on the “Skeptic to Believer in One Week” Turnaround: <https://workos.com/blog/intercom-skeptics-to-believers-ai>
- [12-09] Rising from the Ashes: Eoghan McCabe Returns After Four Years Away from the CEO Role Due to Illness, Sees Stalled Growth, and Bets Intercom on AI in “Founder Mode”: <https://www.lennysnewsletter.com/p/how-intercom-rose-from-the-ashes-eoghan-mccabe>
- [12-10] The AI SDR Launch Done Backwards: Jaspar Demos Live at SaaStr AI Annual, Cites His Own Customer Numbers, and Even Tells the Story of How Terrible the First Two Months Were: <https://www.saastr.com/artisans-ava-2-0-what-a-fully-autonomous-ai-bdr-actually-looks-like-in-production-with-ceo-jaspar-carmichael-jack/>
- [12-11] Eating Their Own Dog Food for Three Years: How Intercom’s Support Team Got Fin to Resolve 81% of Support Volume, Absorbed 300% Demand Growth, and Turned the Team into One Doing Retention: <https://www.intercom.com/blog/automate-customer-service-while-improving-customer-experience/>

- [12-12] Customer-Service AI Blowout on Camera: DPD’s Chatbot Was Provoked by a Customer into Swearing and Wrote a Poem Insulting the Company — the Firm Scrambled to Take the AI Module Offline (January 2024): <https://www.independent.co.uk/tech/chatbot-swears-ai-dpd-poem-b2481967.html>
- [12-13] Customer Service Goes from Cost Center to Growth Engine: Leadership Circle Serves 3 Million Customers with an AI-First Support System — Transformation Lead Miranda Dunn: <https://www.atlassian.com/blog/customer-service-management/ai-first-customer-support-transformation-lessons>
- [12-14] We Built an AI Customer-Success Agent Nobody Used — Then Killed It Ourselves: Boomi Enterprise-AI Team Lead Luke Hagstrand on the “Deconstruction”: <https://boomi.com/blog/customer-success-ai-agent/>
- [12-15] The Map-Makers Are Still in the Room: Field Notes from Pulse 2026 — From Frank Auger Proving with Two Customer Cohorts in 2010 That CSMs Can Curb Churn, to Pa...: <https://unchurned.gainsight.com/p/pulse-2026-field-notes-post-sales>
- [12-16] Artisan, Which Shouted “Don’t Hire Humans,” Raised \$25M and Is Still Hiring: The 23-Year-Old CEO Admits the Early Product Had “Hallucinations So Bad I Could Barely Watch” and Customers Churned in Waves: <https://techcrunch.com/2025/04/09/artisan-the-stop-hiring-humans-ai-agent-startup-raises-25m-and-is-still-hiring-humans/>

Appendix B: Common Metrics

A quick-reference table of the metrics commonly used in AI-powered customer lifecycle management, organized by stage of the customer lifecycle.

Acquisition

Metric	Meaning	Reference benchmark
CAC (customer acquisition cost)	The total cost of acquiring one new customer	Varies by channel; commonly from a few thousand to over ten thousand dollars in B2B SaaS
CPL (cost per qualified lead)	The cost of acquiring one qualified sales lead	Blue Yonder case: fell to \$145 after ABM + intent data
Cold email reply rate	Share of outbound emails that receive a reply	Generally below 5% in 2026; AI templated outreach decayed 60%+ over 18 months
AI SDR reply-rate decay	How the effectiveness of AI outreach changes over time	11.2% (launch) → 4.4% (18 months)

Conversion

Metric	Meaning	Reference benchmark
Conversion rate	Rate at which visitors/ leads convert into customers	E-commerce industry bottleneck of 2-3%; AI personalization can lift it by up to 25%
Revenue per visitor (RPV)	Revenue contributed by a single visitor	AI personalization lifts it by 20%+ on average
AI dynamic-pricing gains	Revenue change from real-time pricing	Airline load factors 72% → 80%+ (with fares actually lower)

Service

Metric	Meaning	Reference benchmark
Deflection rate (auto-resolution rate)	Share of conversations AI resolves without handoff to a human	Intercom Fin: 50-70%; vendor-reported figures commonly 60-80%, 30-40 percentage points above independent data
Average resolution time	Time from a question being raised to it being resolved	Klarna: AI 11 minutes → 2 minutes (humans ~12 minutes)
CSAT (customer satisfaction)	Customers' satisfaction score for the service	Klarna: AI on par with human agents
AI support ROI	Return for every \$1 invested	Vendor-reported ~\$3.5; independent data significantly lower

Retention & Expansion

Metric	Meaning	Reference benchmark
Churn rate	Share of customers lost over a period	US businesses lose \$136.8 billion annually to avoidable churn
NRR (net revenue retention)	Net revenue change contributed by existing customers	Median 100-115%; top quartile 115-130%; best-in-class 130%+ (usage-based platforms)
Health score	Composite metric of a customer's likelihood to renew/churn	73% of companies admit their current health scores cannot reliably predict churn
Win-back rate	Share of dormant customers successfully reactivated	AI-driven win-back ~22% vs. 5-10% for bulk sends

Organization & Change

Metric	Meaning	Reference benchmark
Share of selling time	Time salespeople actually spend with customers	Only ~25% (Bain)
New-rep ramp time	Time for a rep to go from onboarding to reaching quota	AI role-play: 210 days → 72 days (Hyperbound × Vanta)
ARR per FTE (output per employee)	Annualized revenue contributed per employee	Fleet of AI agents: \$370,000 → \$640,000 (ICONIQ)

Note: The metrics in this table use the conventions adopted throughout this book. Most "benchmarks" vary widely by industry and methodology—when citing, defer to the original sources in Appendix A.

Appendix C: Version History and Update Notes

This book is a "living book"—material keeps accumulating and the version keeps growing. This appendix records what each version updated and why.

v1.0.3 (2026-09-20, weekly update)

- Chapter 3 adds "Entry Points: In the Chinese Market, the Search Box Is Turning into a Dialogue Box," "AI Visibility: A Second Report Card Beyond Rankings," and "Agentish: The Boundary of Autonomous Prospecting"
- Chapter 5 adds "The Cost Ledger of Human Handoff Is Becoming a Compliance Ledger" and "Resolution Rate Becomes an Industry Unit, and People Are Re-Tiered by Value"
- Chapter 6 adds "From Health Scores to Triggers"
- Chapter 7 adds "Segmentation and Reach: A Postmortem on 4.04%"
- Chapter 8 adds "AI Will Not Save a Bad Process"
- Chapter 9 adds "New Industry Signals (Part Two): Turning 'Someone Has to Watch It' into Capacity"
- Chapter 10 adds "The Liability of AI Output Is Being Priced," "The Governance Time Lag on Outbound Calling," and "AI Poisoning: The Supply Chain of GEO Gray Markets"
- Chapter 11 adds "Protocols and Admission Tickets: The Barrier to Entry in the Machine-Customer Era" and "Machine Customers' Money and Trust"
- Appendix A adds 44 sources (3-55 to 3-68 / 5-48 to 5-51 / 6-30 to 6-32 / 7-34 / 8-46 to 8-49 / 9-104 to 9-105 / 10-25 to 10-33 / 11-29 to 11-40), synced across both languages

v1.0.2 (2026-09-13, weekly update)

- Chapter 3 adds "The AI-ization of Advertising and the Truthfulness of Traffic": conversational ads as a new entry point, AI-answer traffic converting 4.4x better, and 60% of Google app-install ads going to a bot farm

- Chapter 8 adds "FDE: The Intermediate State of AI Deployment" and "An Organizational Prescription for AI-Native Teams": 24 first-person FDE cases from Datawhale, three-bucket process redesign, and pods/Agent Champions/STO with system-quality metrics
- Chapter 9 adds "New Industry Signals": bank agent-platformization, more than 40% of agentic AI projects to be canceled in manufacturing, and citizens fearing government AI adoption is too slow
- Chapter 10 adds "When Agents Go Find Customers Themselves": 7 models running businesses autonomously for 72 hours with \$0 revenue, colluding rogue agents and the RubyGems attack, and the financial-sector AI compliance backstop
- Chapter 11 adds "The Machine-Customer Ledger": institutional market estimates for agentic commerce and the risk of brands being intermediated
- Appendix A adds 20 sources (3-52 to 3-54 / 8-41 to 8-45 / 9-98 to 9-103 / 10-22 to 10-24 / 11-26 to 11-28), synced across both languages

v1.0.1 (2026-09-08, English edition renamed)

- The English edition title changed from *AI-Powered Customer Operations* to **AI Customer Lifecycle Management**, with the subtitle *Acquisition, Conversion, Service, Retention, and Expansion in the Age of AI*
- Rationale: "customer lifecycle management" is an established market search phrase and covers the book's full arc from acquisition to expansion, while "operations" reads as execution-layer scope and undersells the front stages
- Scope of the rename: English terminology throughout the book, the glossary, the PDF cover, and the site meta description; the Chinese title and Chinese content are unchanged

v1.0.0 (2026-09-06, initial public release)

- First public release: open for reading on the GitHub repository and the companion website
- Content status: preface + 11 main chapters + afterword + Appendices A/B/C; 247 key data points in the main text carry numbered citations, and Appendix A contains 399 sources (each traceable back to its original text)

v0.2 (2026-09-06, citation-annotation release)

- Added numbered superscripts to the 247 key data points across the 11 main chapters (placed after the data point and before the source annotation; clickable, jumping to the corresponding entry in Appendix A)
- Appendix A upgraded from plain-text numbering to anchor-based numbering: all 399 sources now carry `<a id>` anchors with "chapter-sequence" ids, with 0 dangling citations in the main text and 0 duplicate anchors
- Added cross-topic sources (e.g., the Glasp AEO case [3-51]) so that in-book citations remain self-consistent

v0.1 (2026-09-05, first draft)

- First draft complete: preface + 11 main chapters + Appendices A/B/C. Material came from the continuously growing knowledge base (on AI-powered customer lifecycle management: 398 bilingual Chinese-English items, each tagged [official claims] or [third-party verified])
- Chapter structure: organized around the five stages of the customer lifecycle (acquisition → conversion → service → retention → expansion), plus the foundation, organization, industries, the dark side, and the future
- Writing conventions: scene-based openings, success and failure cases side by side, traceable data, no authorial voice
- To-do: English translation (deferred); full source verification of key data points against the originals; revisions based on reader feedback

Afterword

This book has a completed first draft, but strictly speaking, it will never be finished.

The process of writing this book was itself an experiment in AI-powered customer lifecycle management. The material came from an automated collection pipeline: every day it scans the global news on AI and customer operations, then distills, translates, structures, and verifies sources before storing them in a knowledge base. From topic kickoff to 398 pieces of material took less than two weeks; from material to first draft took a few days. The Klarna case "grew" out of the material base on its own—the pipeline first collected the 2024 news that AI had done the work of 700 customer service agents, and months later collected the reversal, "we went too far." Put the two items side by side, and the opening of Chapter 1 wrote itself.

This is exactly what I want this book to convey, and the most essential thing about AI-powered customer lifecycle management: **AI does not think for you—it amplifies your thinking. AI does not run the customer lifecycle for you—it gives you the capacity to run it for many more customers.**

Throughout the writing, several cases kept reminding me:

Klarna reminded me that efficiency is not the goal—customers are. DPD reminded me that AI needs guardrails, just as cars need brakes. Artisan reminded me that a channel dividend is spent once, and restraint is a strategy. Octopus Energy reminded me that when AI support is done well, customers really do stay. Intercom's engineers reminded me that turning skeptics into believers takes only one genuinely good product.

If you have read this far, I hope you take three sentences away with you:

First: AI-powered customer lifecycle management has no silver bullet, but it has a map. This book is that map—engine, foundation, organization, industries, boundaries. Every company can find where it stands today.

Second: data before AI, organization before tools. 74% of companies are stuck in pilots, and seven in ten of the obstacles come from people. Before buying tools, think through whether your data is clean and whether your people are ready.

Third: AI handles scale; humans handle trust. When AI handles 95% of standardized interactions, the remaining 5% of human interactions are the entire battlefield on which brands differentiate. Human service will not disappear—it will become a luxury.

Every figure and case in this book has a source (Appendix A). If you spot any errors or omissions, please let me know: weokk2025@gmail.com.

In the next edition, this book will keep growing—just like the subject it studies.